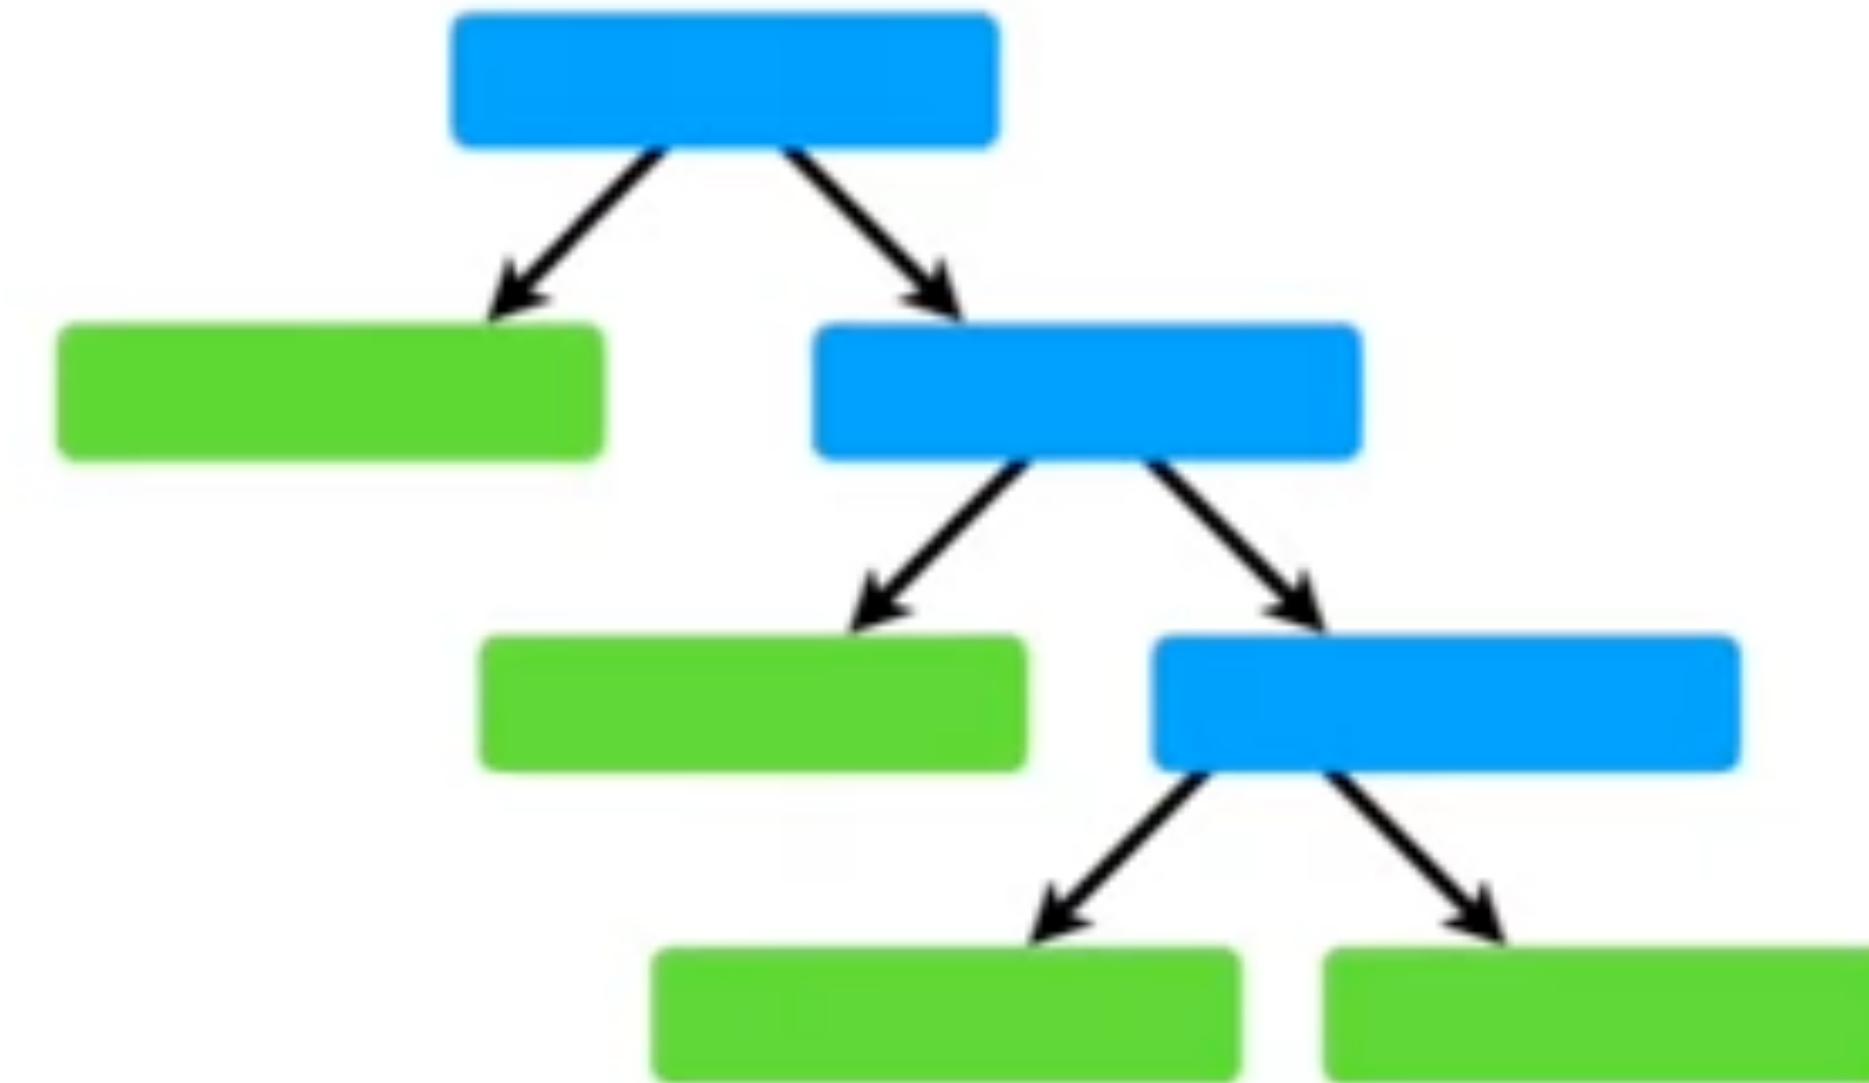


Lecture 24: Introduction to information theory

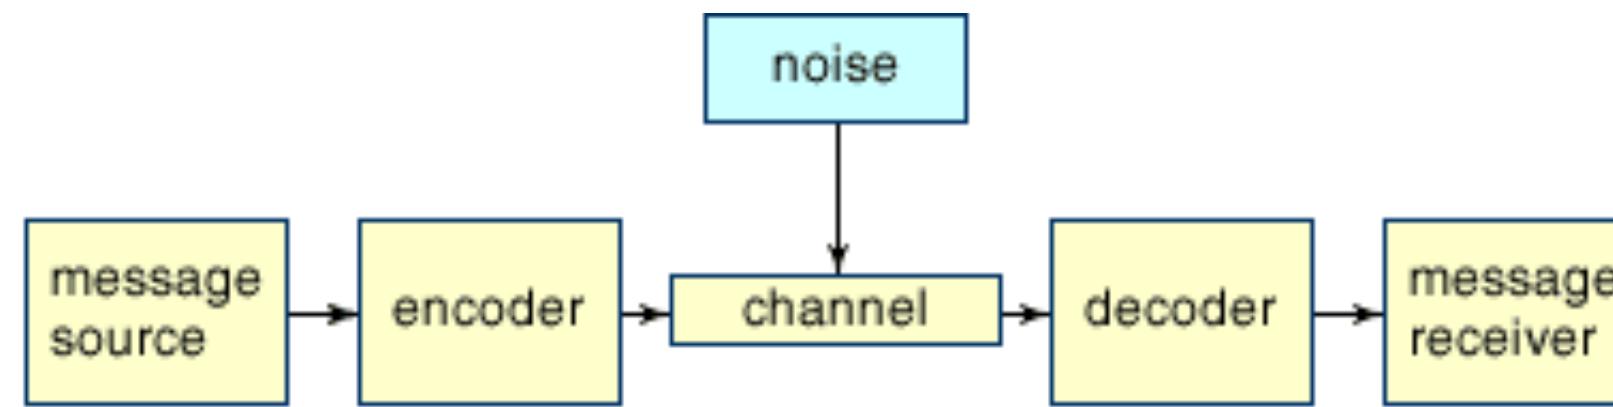
Instructor: Michael Szell

Nov 24, 2023



Today you will learn information theory basics for ML

Entropy

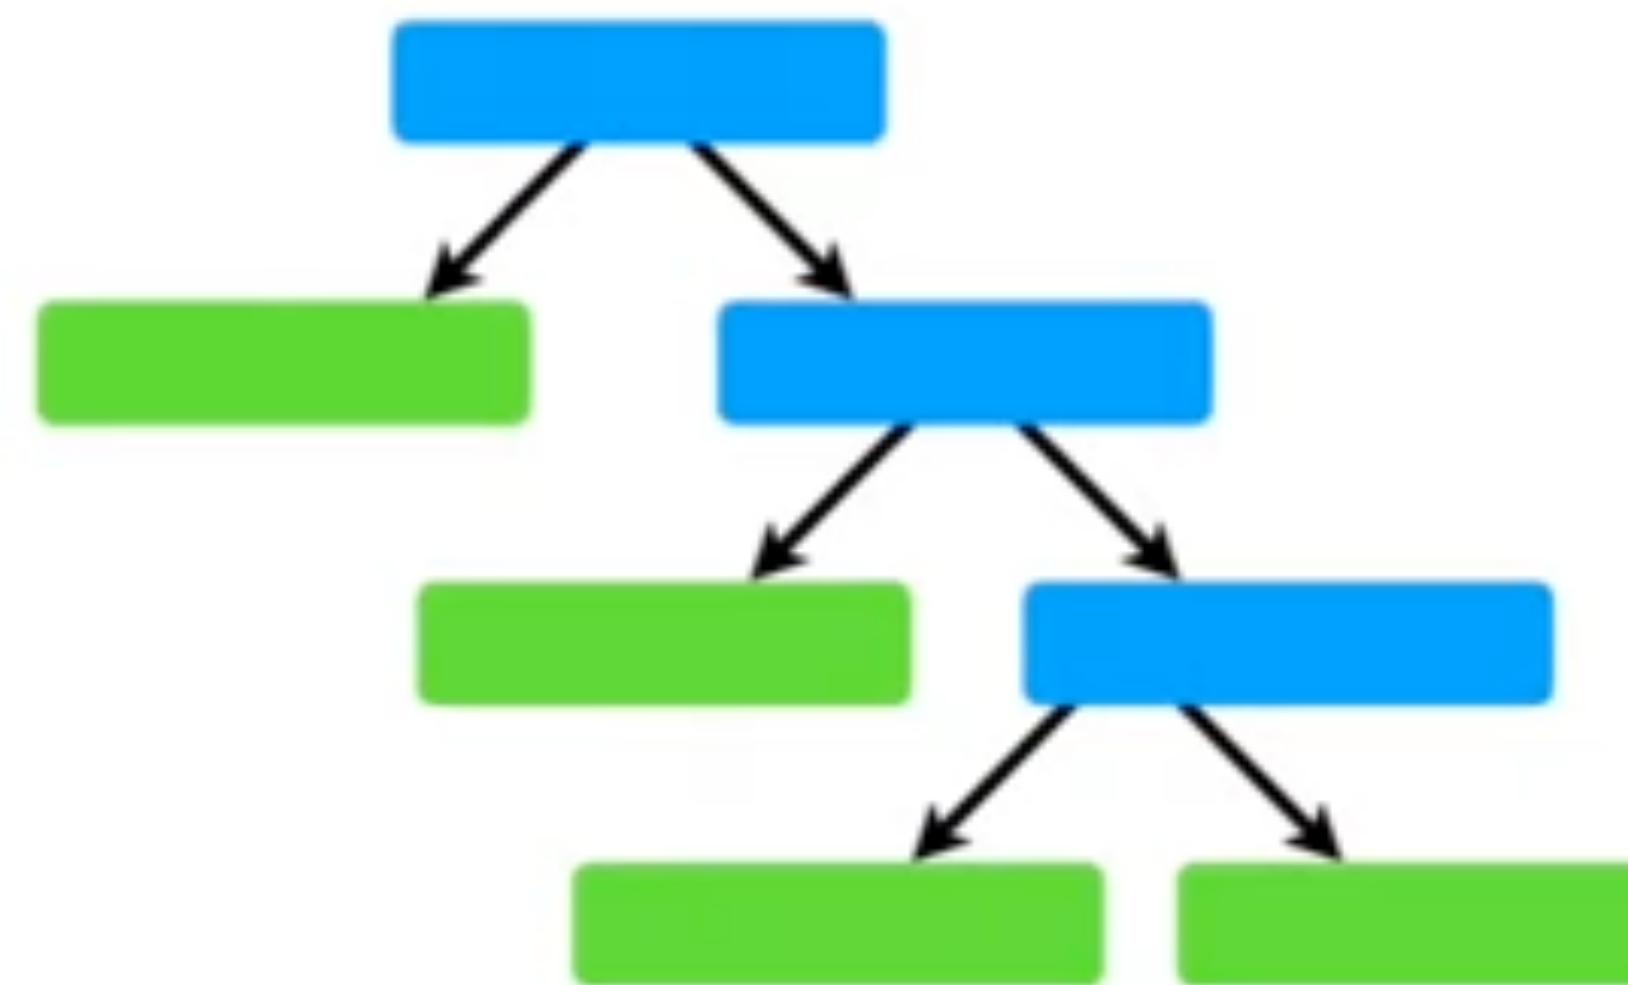


$$H(X) = \mathbb{E}(I(X))$$

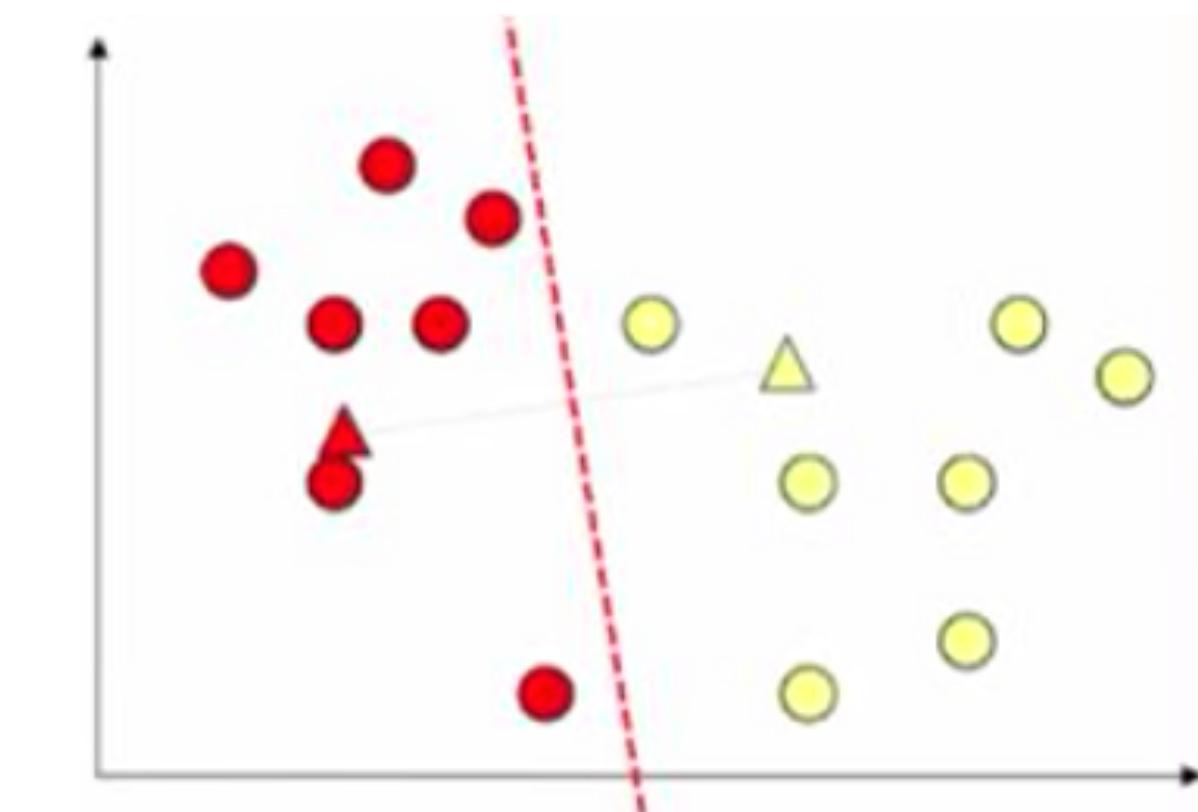
Data Science setup



Decision trees



Clustering



Let's play Bar Khokba

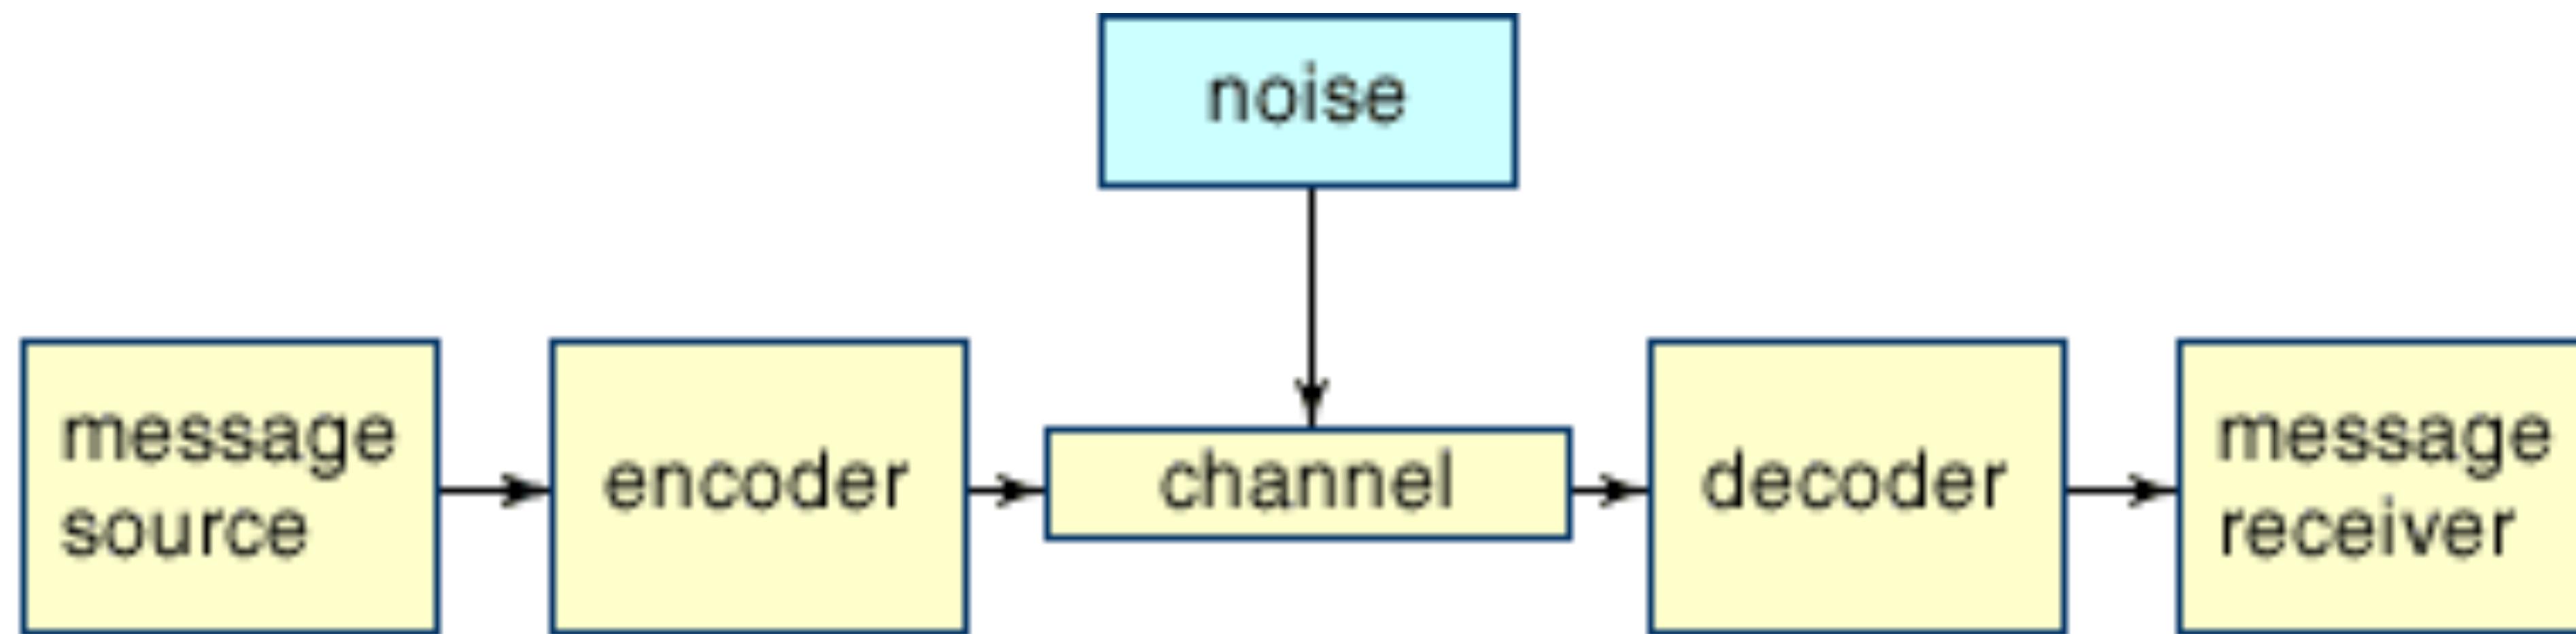


Let's play Bar Khokba

With 7 YES/NO
questions, find one
person in the room

Such problems can be analyzed with **information theory**

Information theory is the mathematical study of the quantification, storage, and communication of information.



Entropy is the basis for many concepts in data science

Building Decision/Classification trees

Mutual information

Kullback-Leibler Divergence

Entropy is the basis for many concepts in data science

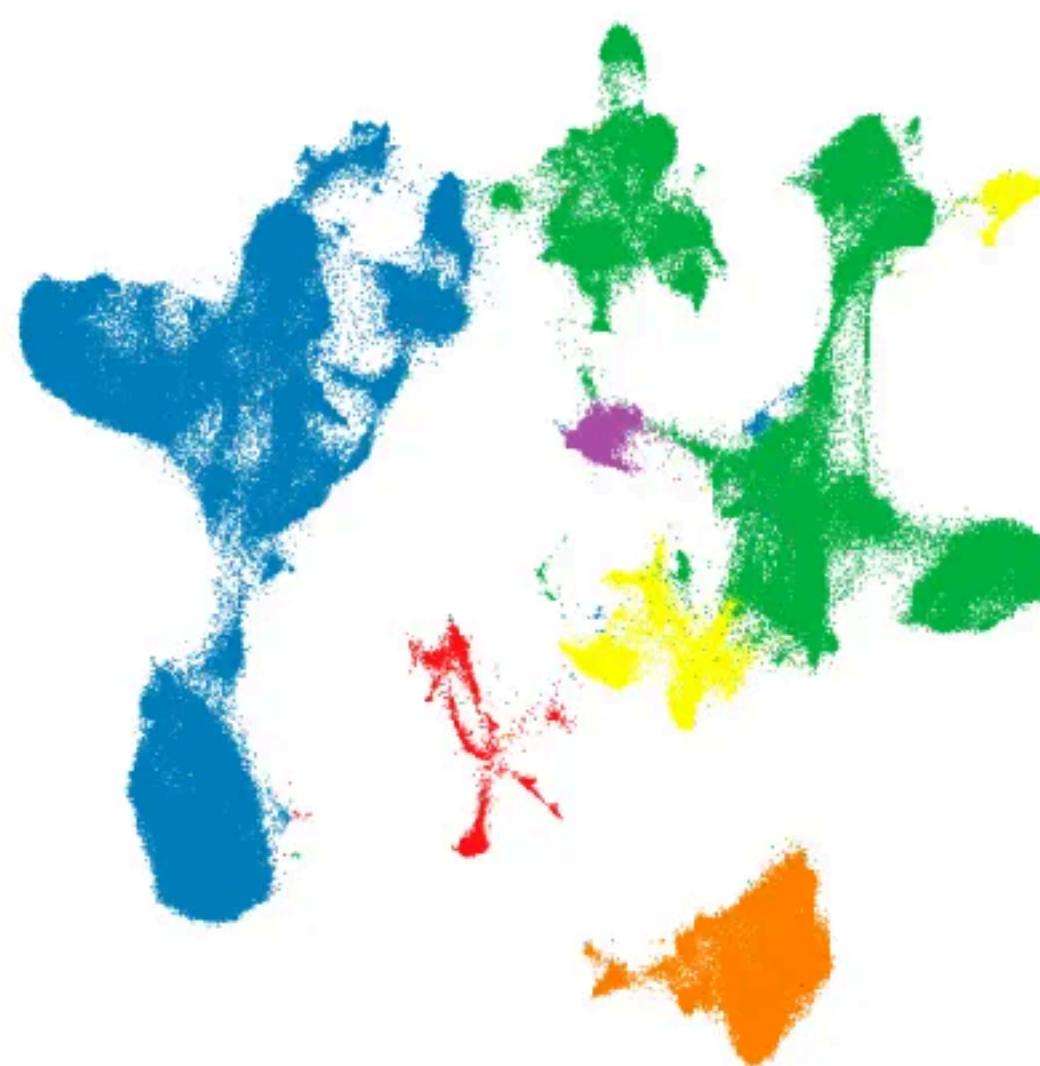
Building Decision/Classification trees

t-SNE

UMAP

a

UMAP

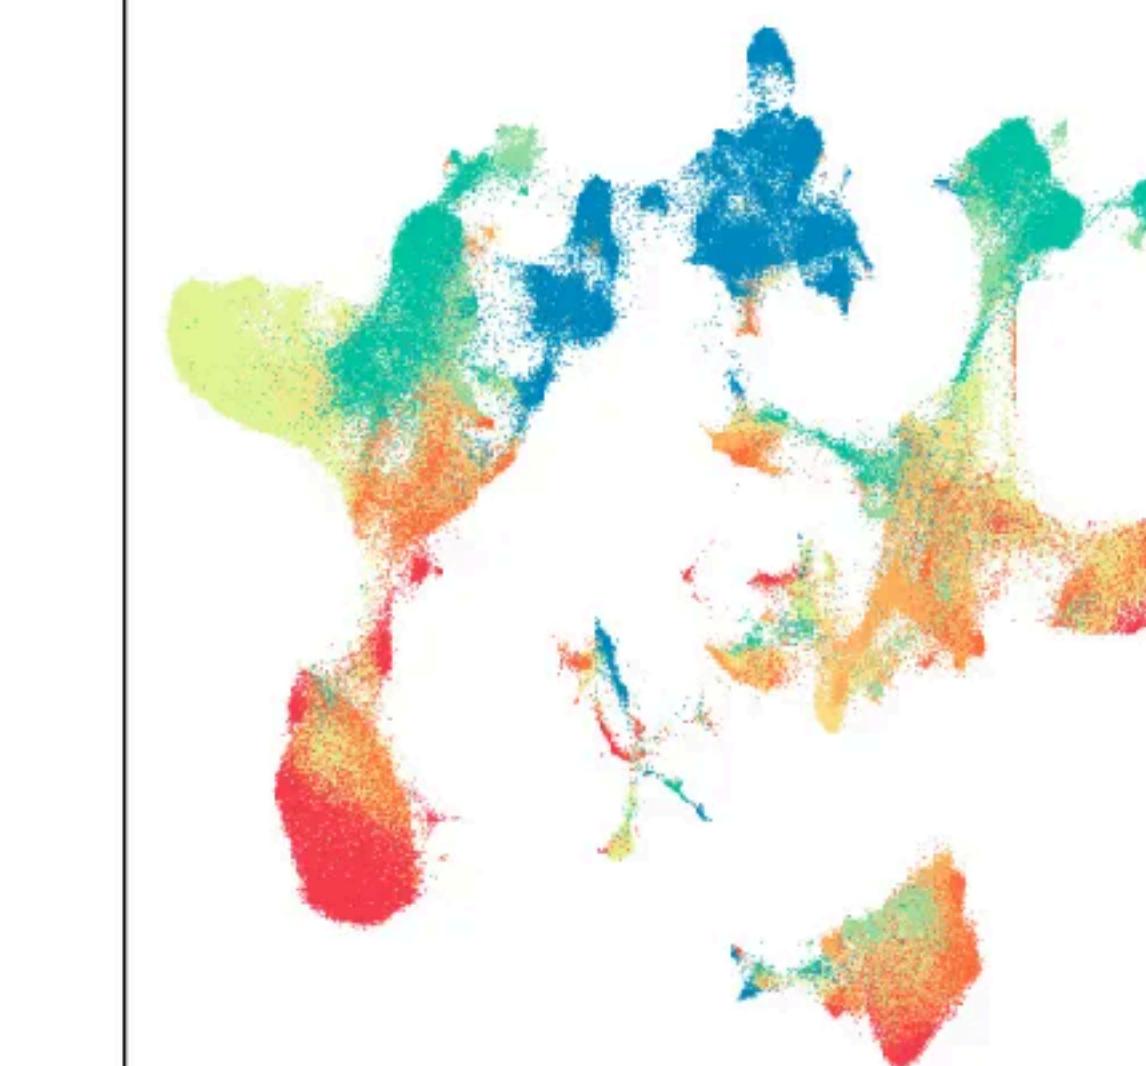


Cell types

- Contaminant (including B)
- CD4 T
- CD8 T
- MAIT
- NK/ILC
- $\gamma\delta$ T

b

UMAP



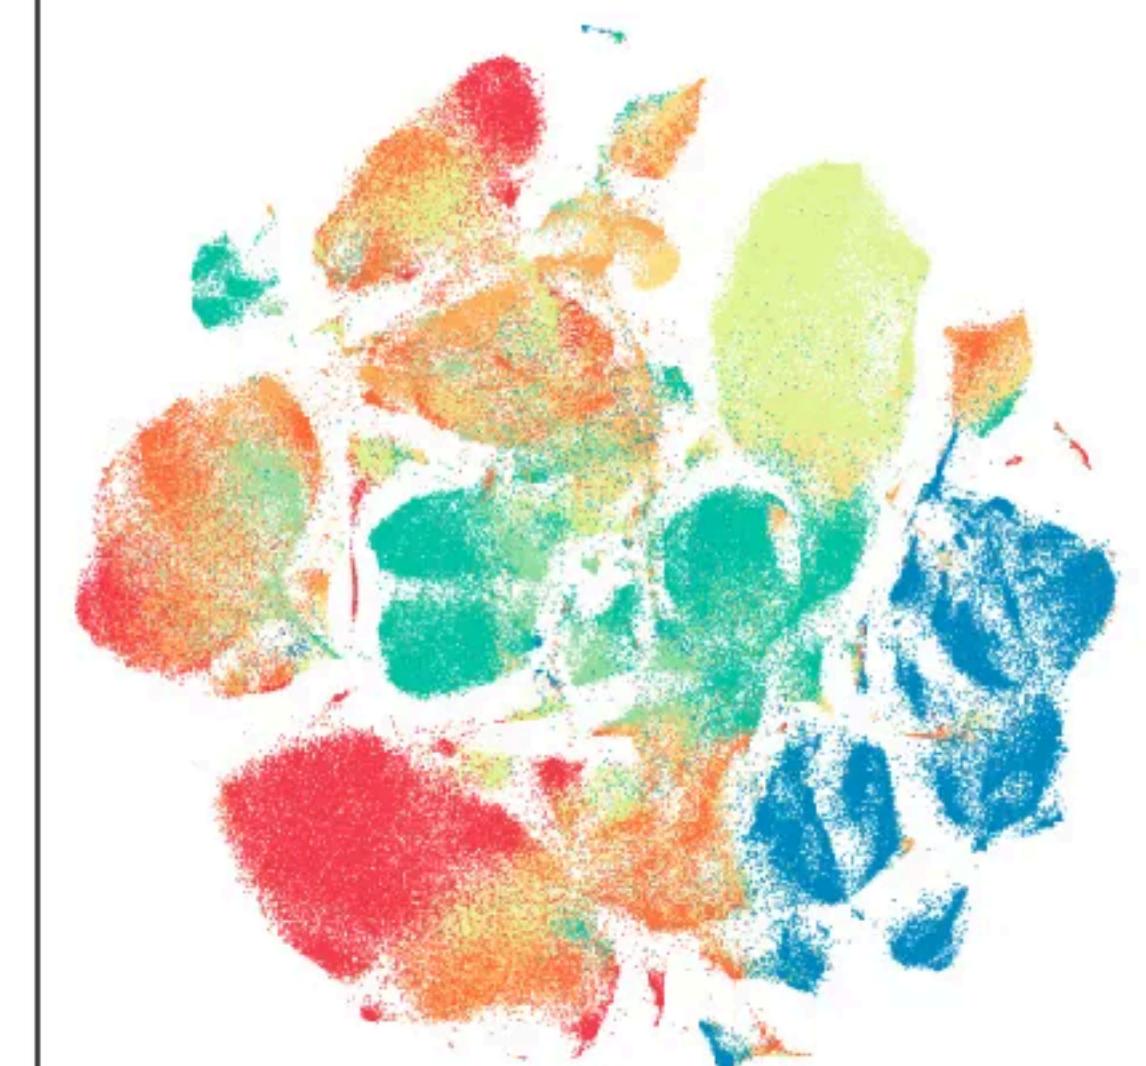
Sample types

- CB
- PBMC
- Liver
- Spleen
- Tonsil
- Lung
- Gut
- Skin

Mutual information

Kullback-Leibler Divergence

t-SNE



Self-information I of an event x is how "surprising" it is:



Self-information I of an event x is how "surprising" it is:

If $p(x) = 1$, it should be perfectly unsurprising, i.e. $I(x) = 0$

The less probable it is, the more surprising, and the higher I should be

Self-information I of an event x is how "surprising" it is:

If $p(x) = 1$, it should be perfectly unsurprising, i.e. $I(x) = 0$

The less probable it is, the more surprising, and the higher I should be

For two independent events, their total I should be the sum of their I s

Self-information I of an event x is how "surprising" it is:

If $p(x) = 1$, it should be perfectly unsurprising, i.e. $I(x) = 0$

The less probable it is, the more surprising, and the higher I should be

For two independent events, their total I should be the sum of their I s

$$I(x) = \log_2 \frac{1}{p(x)} = -\log_2 p(x)$$

Entropy is the expected surprisal

Blackboard

$$H(X) = \mathbb{E}(I(X))$$



Entropy is the expected surprisal

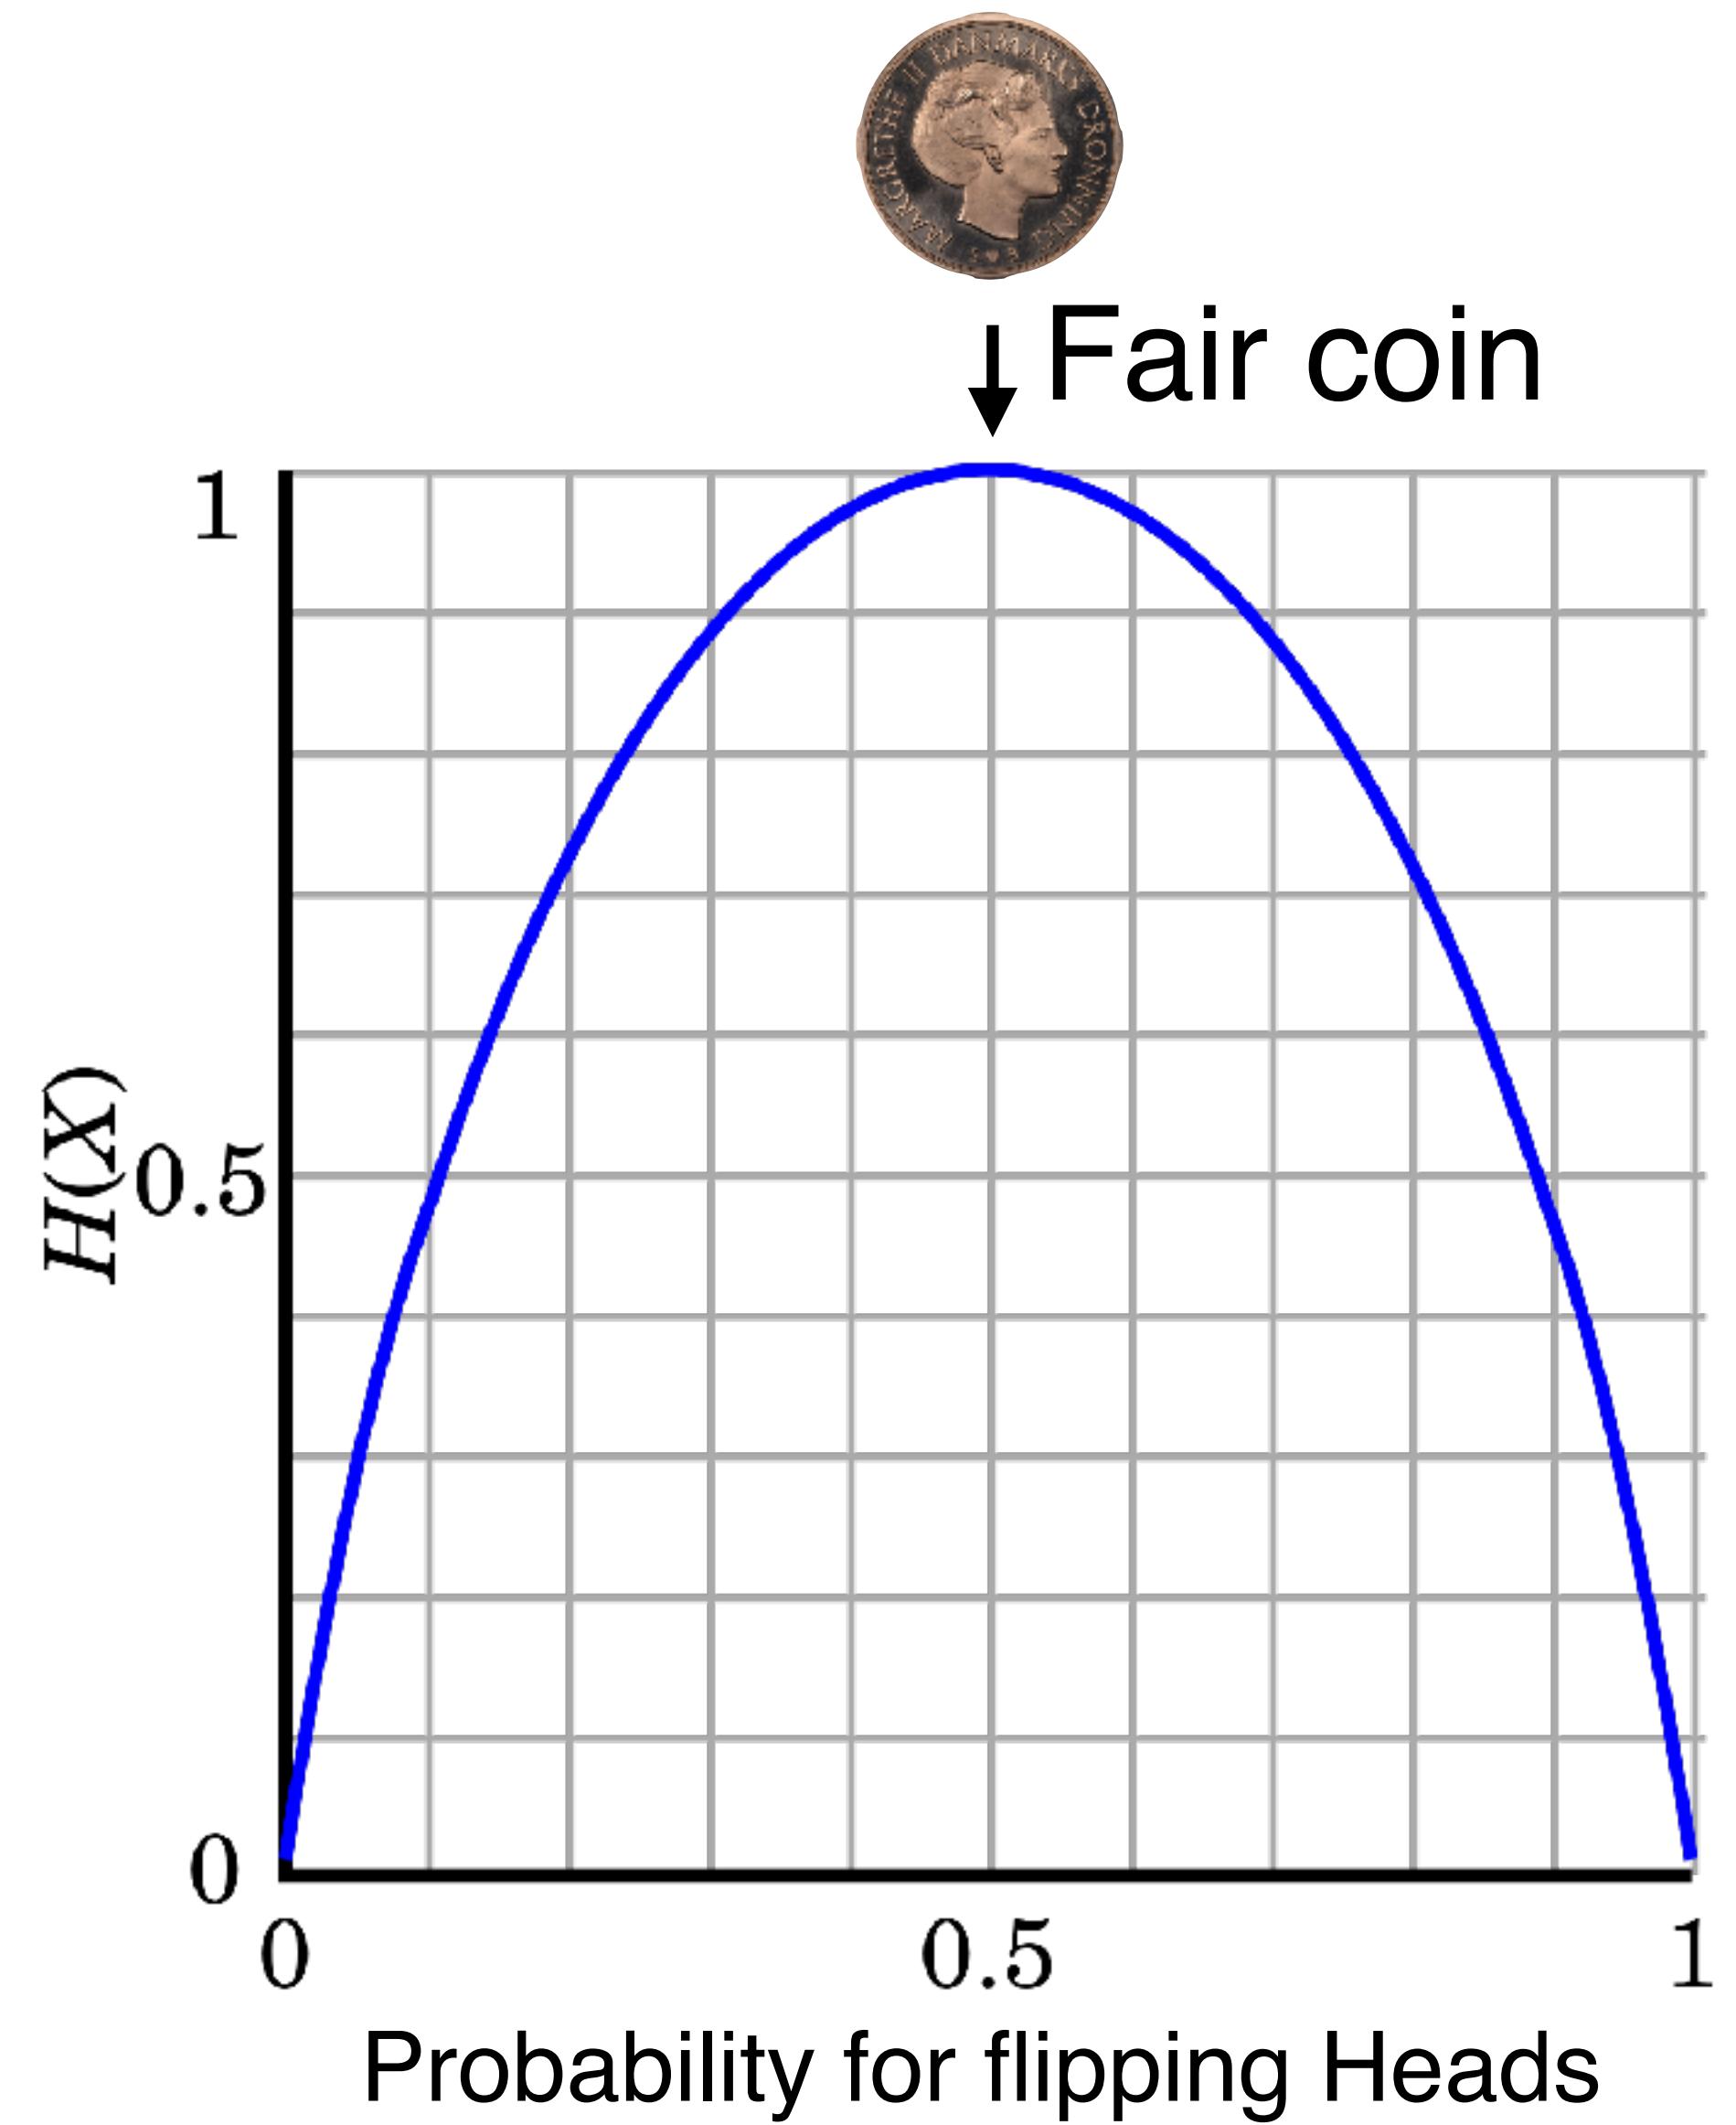
$$H(X) = \mathbb{E}(I(X))$$

$$= \sum_x p(x) I(x)$$

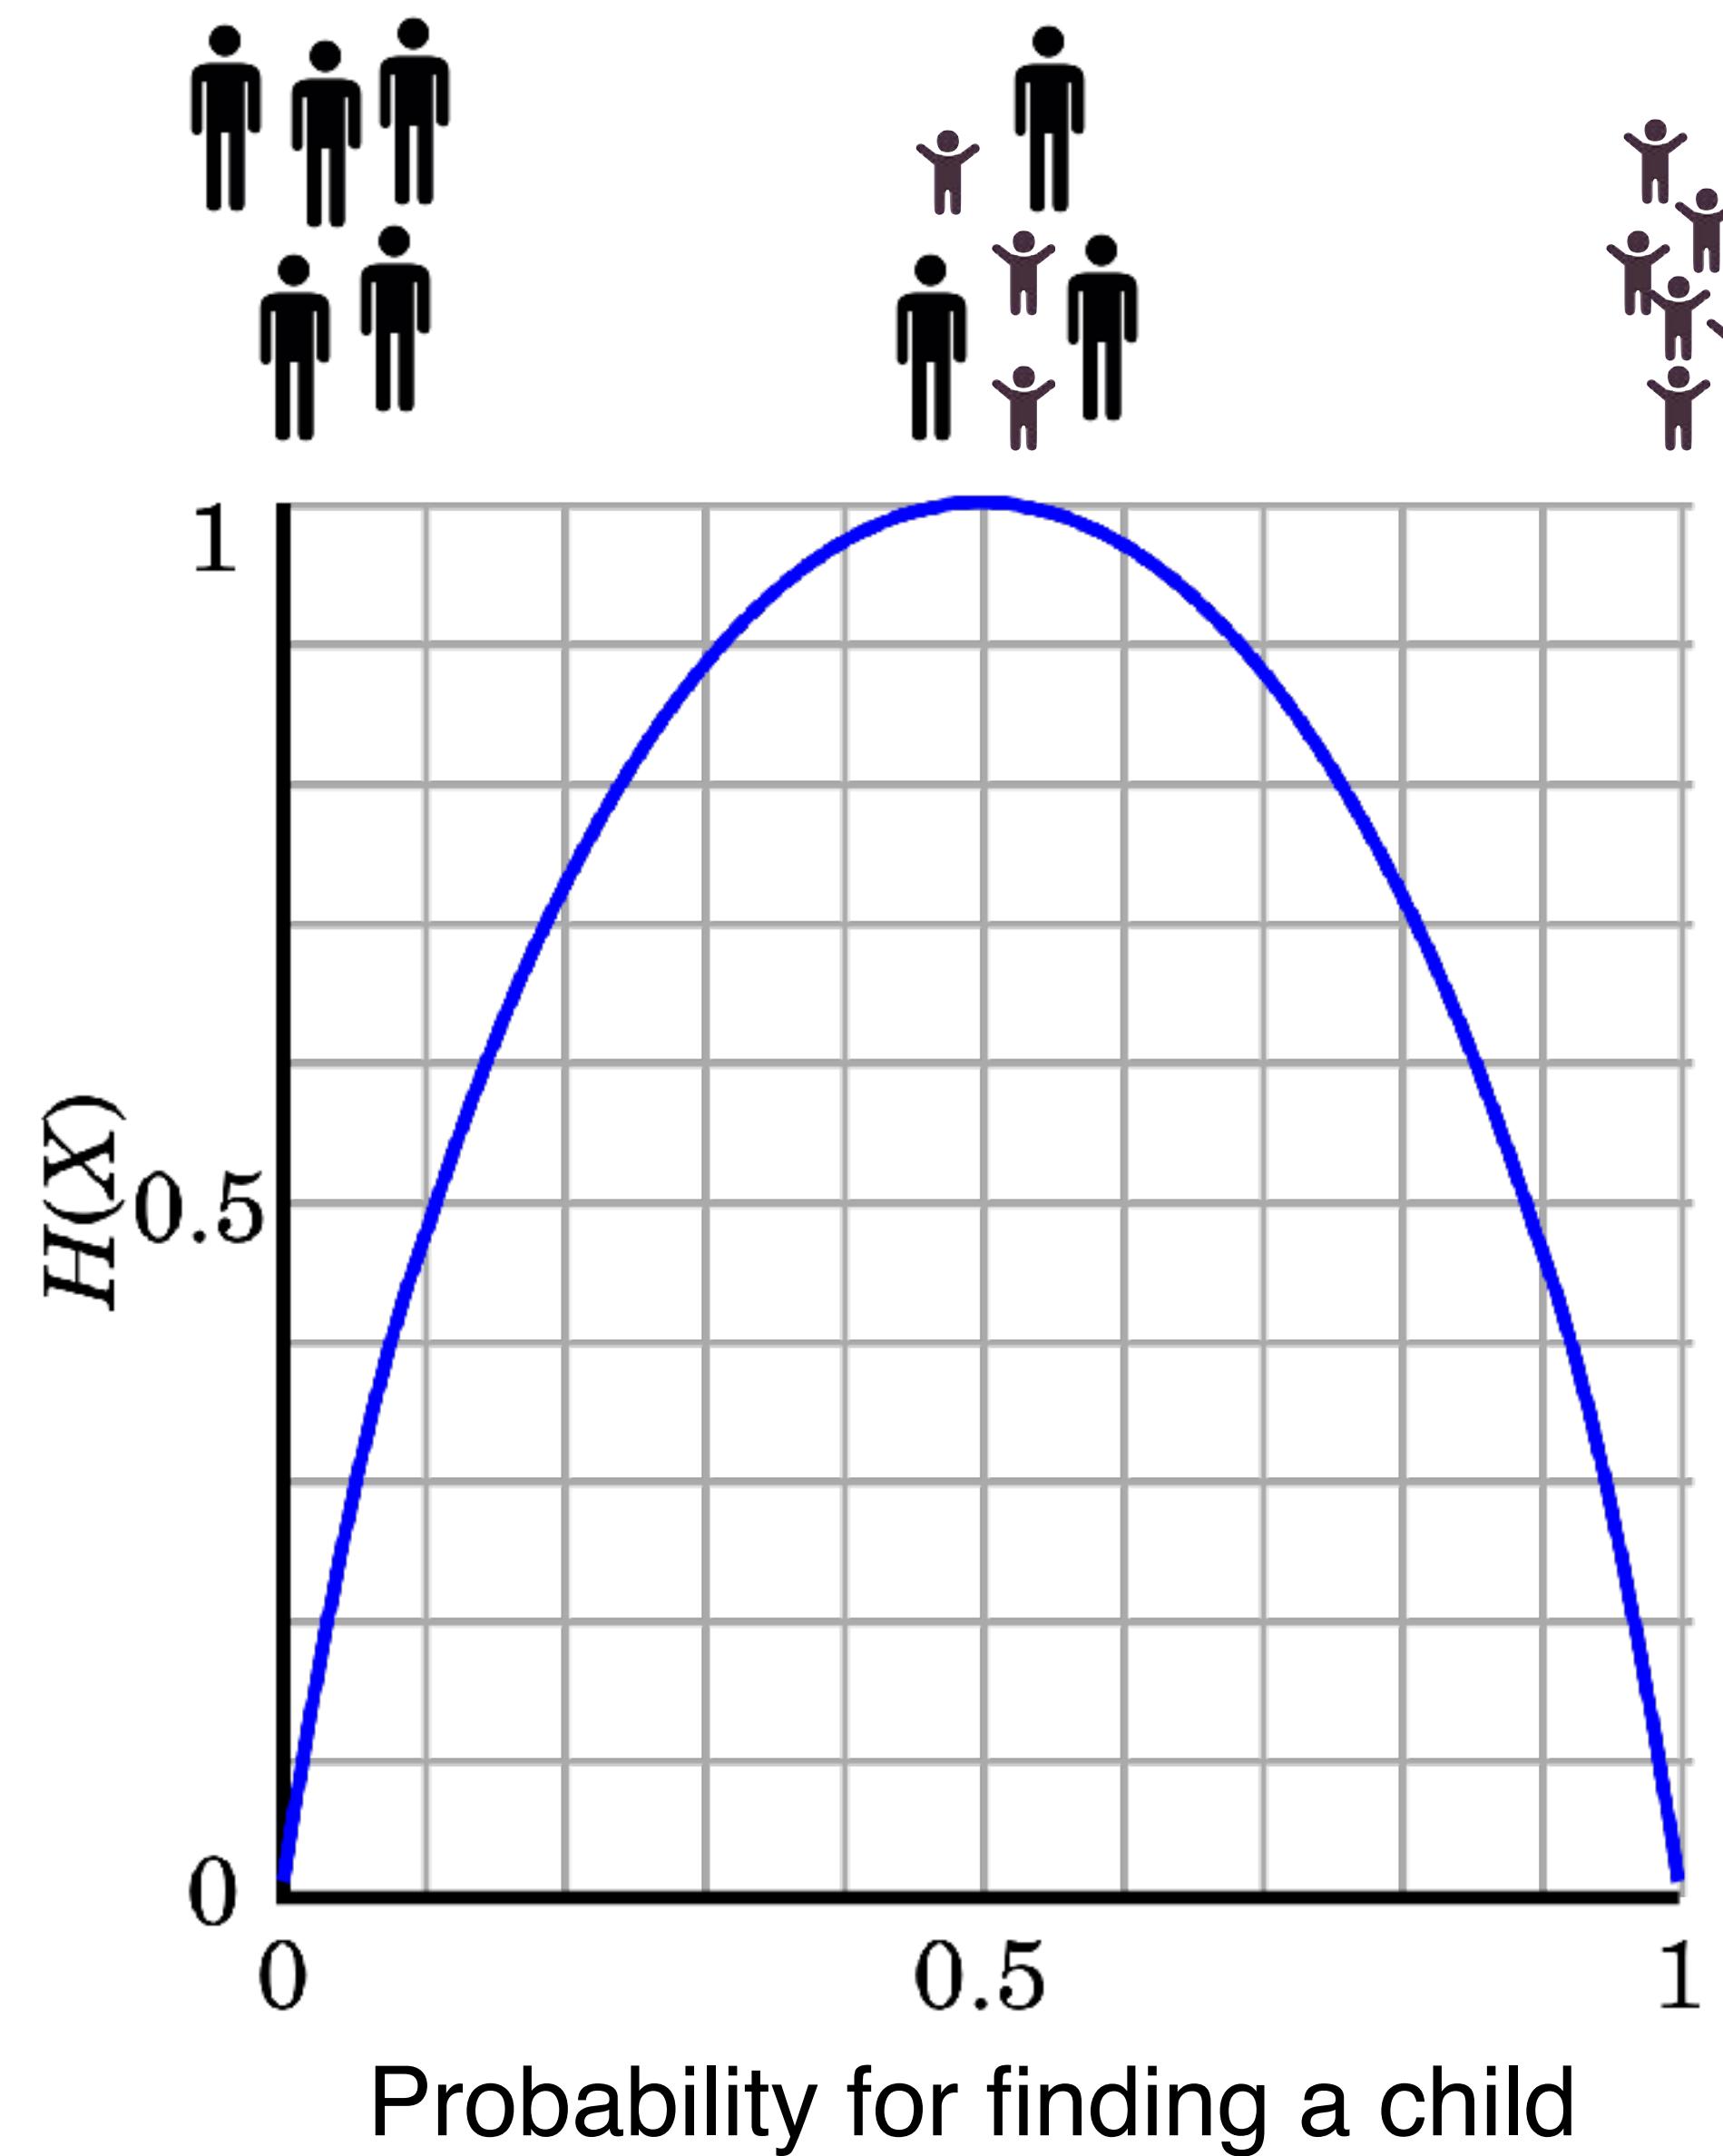
$$= - \sum_x p(x) \log_2 p(x)$$

See also: [https://www.youtube.com/
watch?v=YtebGVx-Fxw](https://www.youtube.com/watch?v=YtebGVx-Fxw)

Example: Entropy of a coin

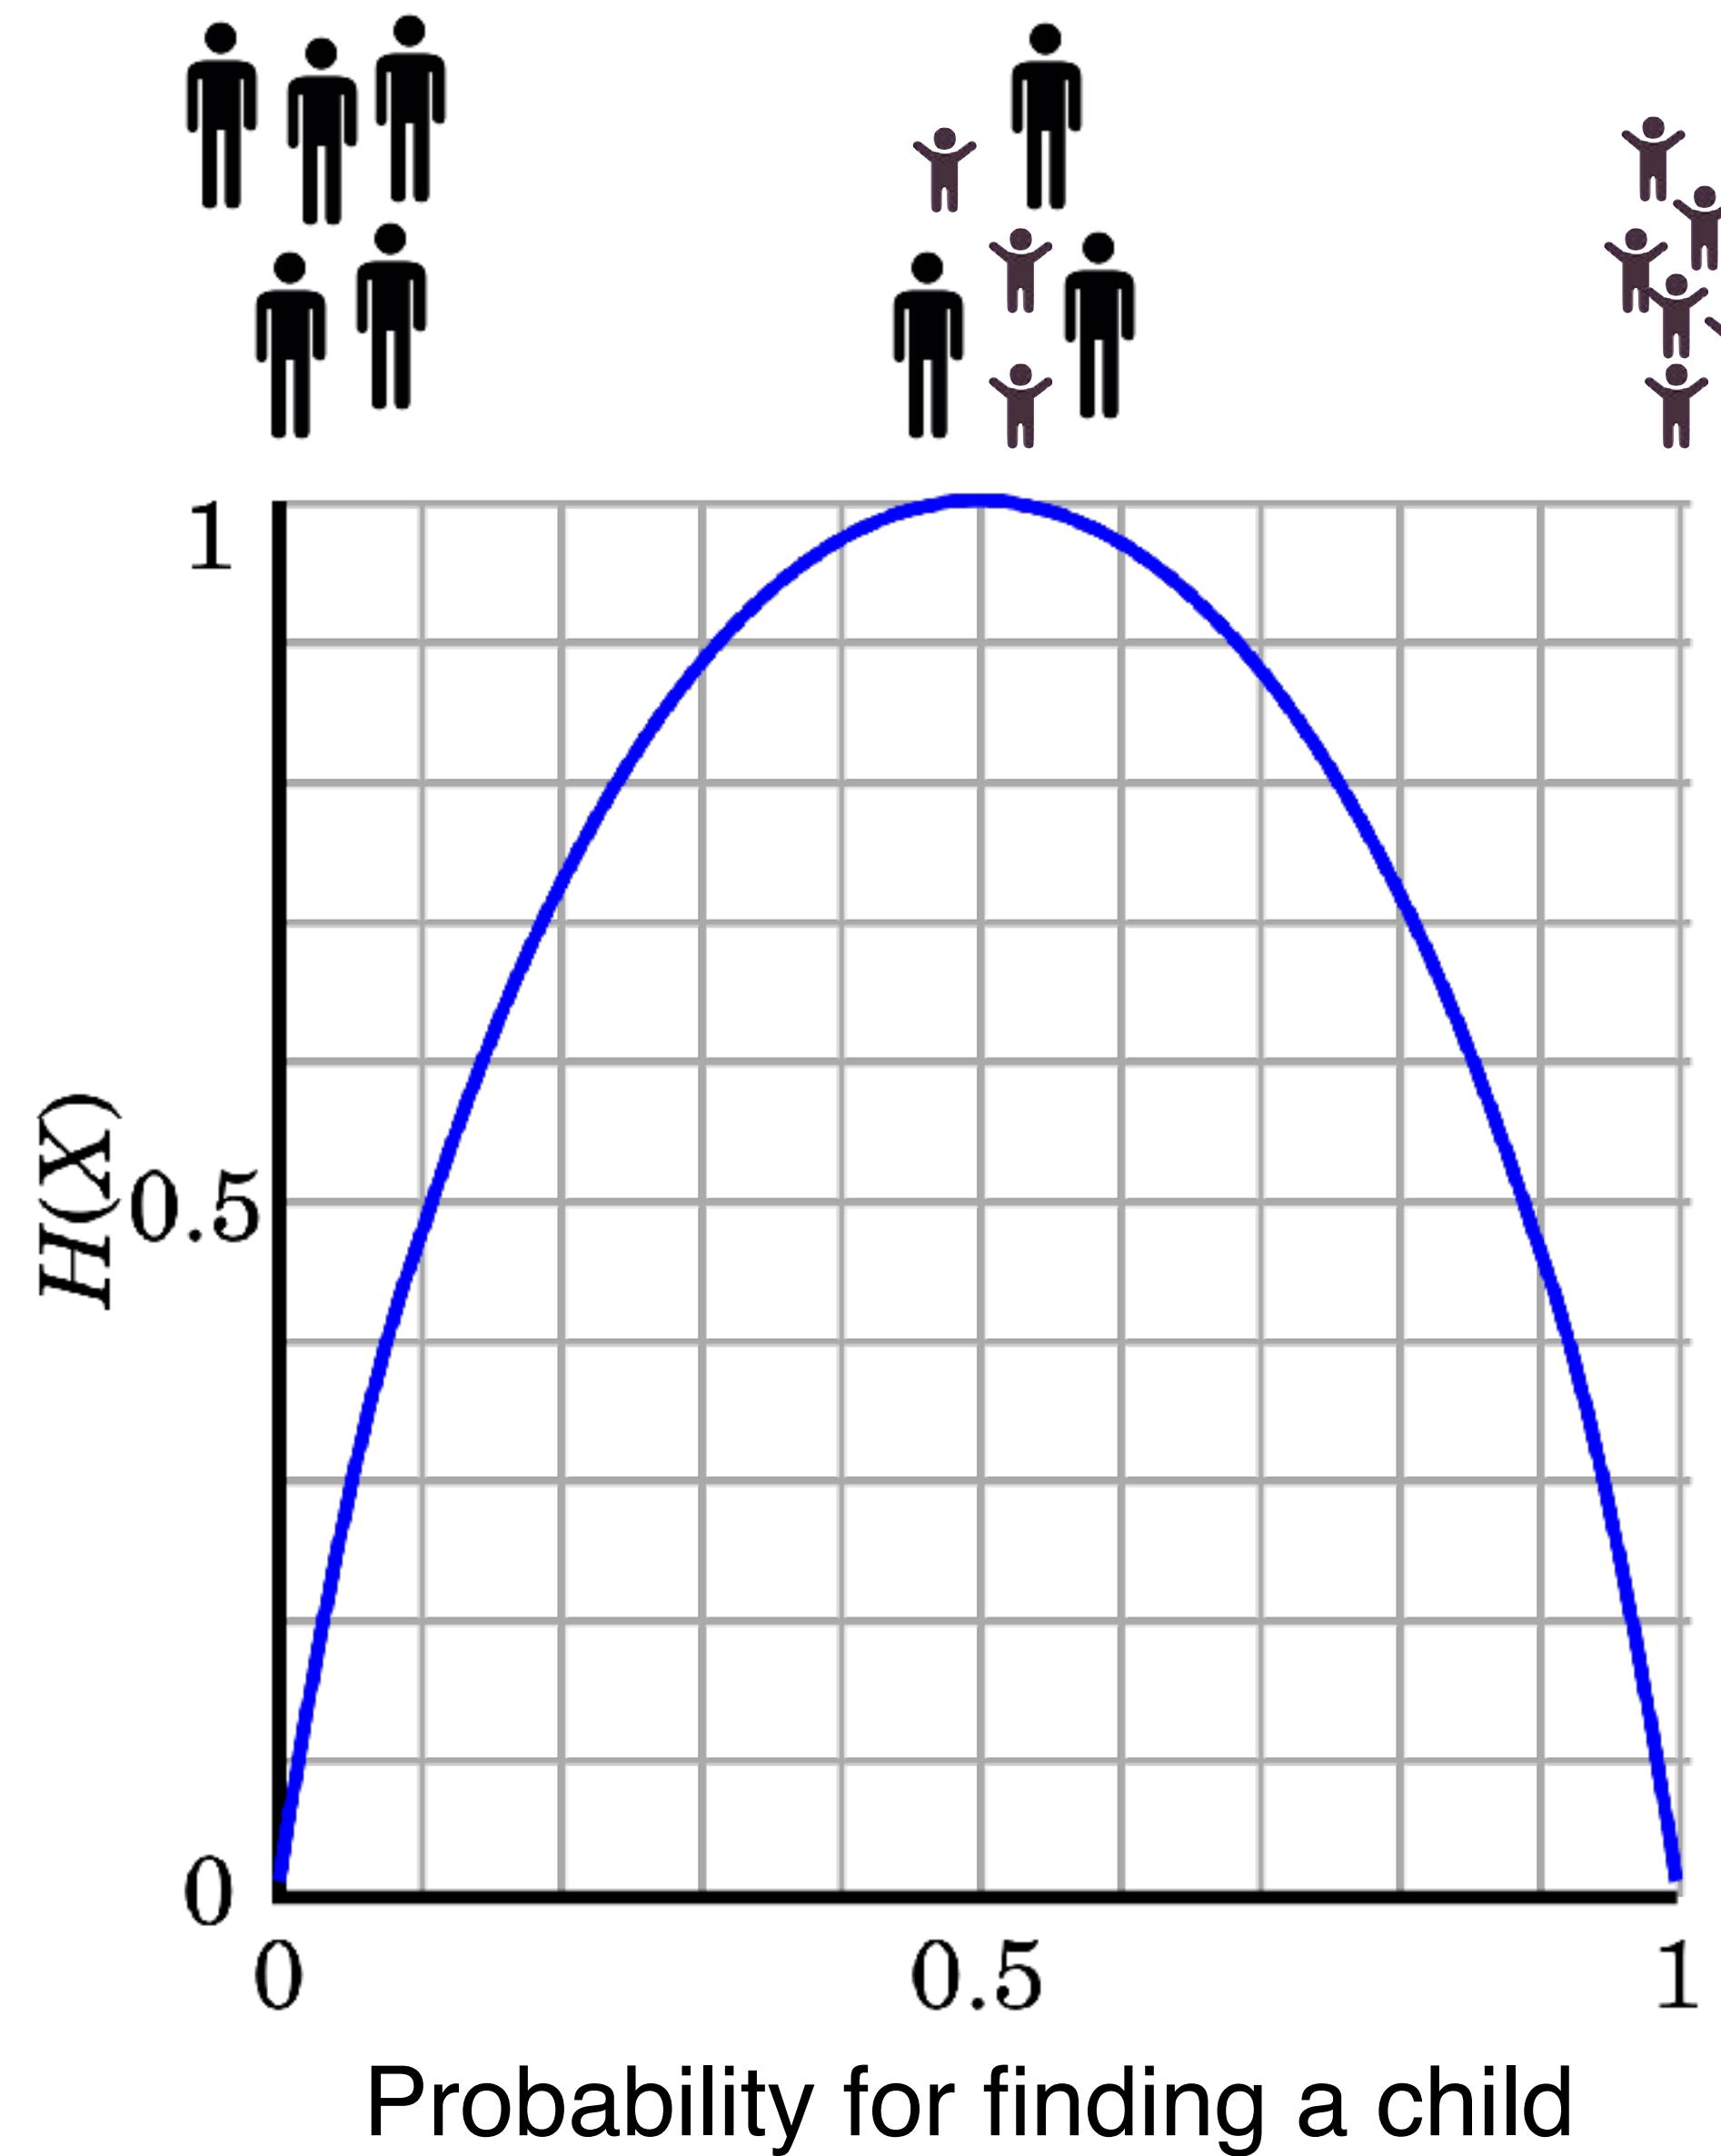


Example: Entropy of a binary data distribution



Here entropy is a measure of
balance or equality

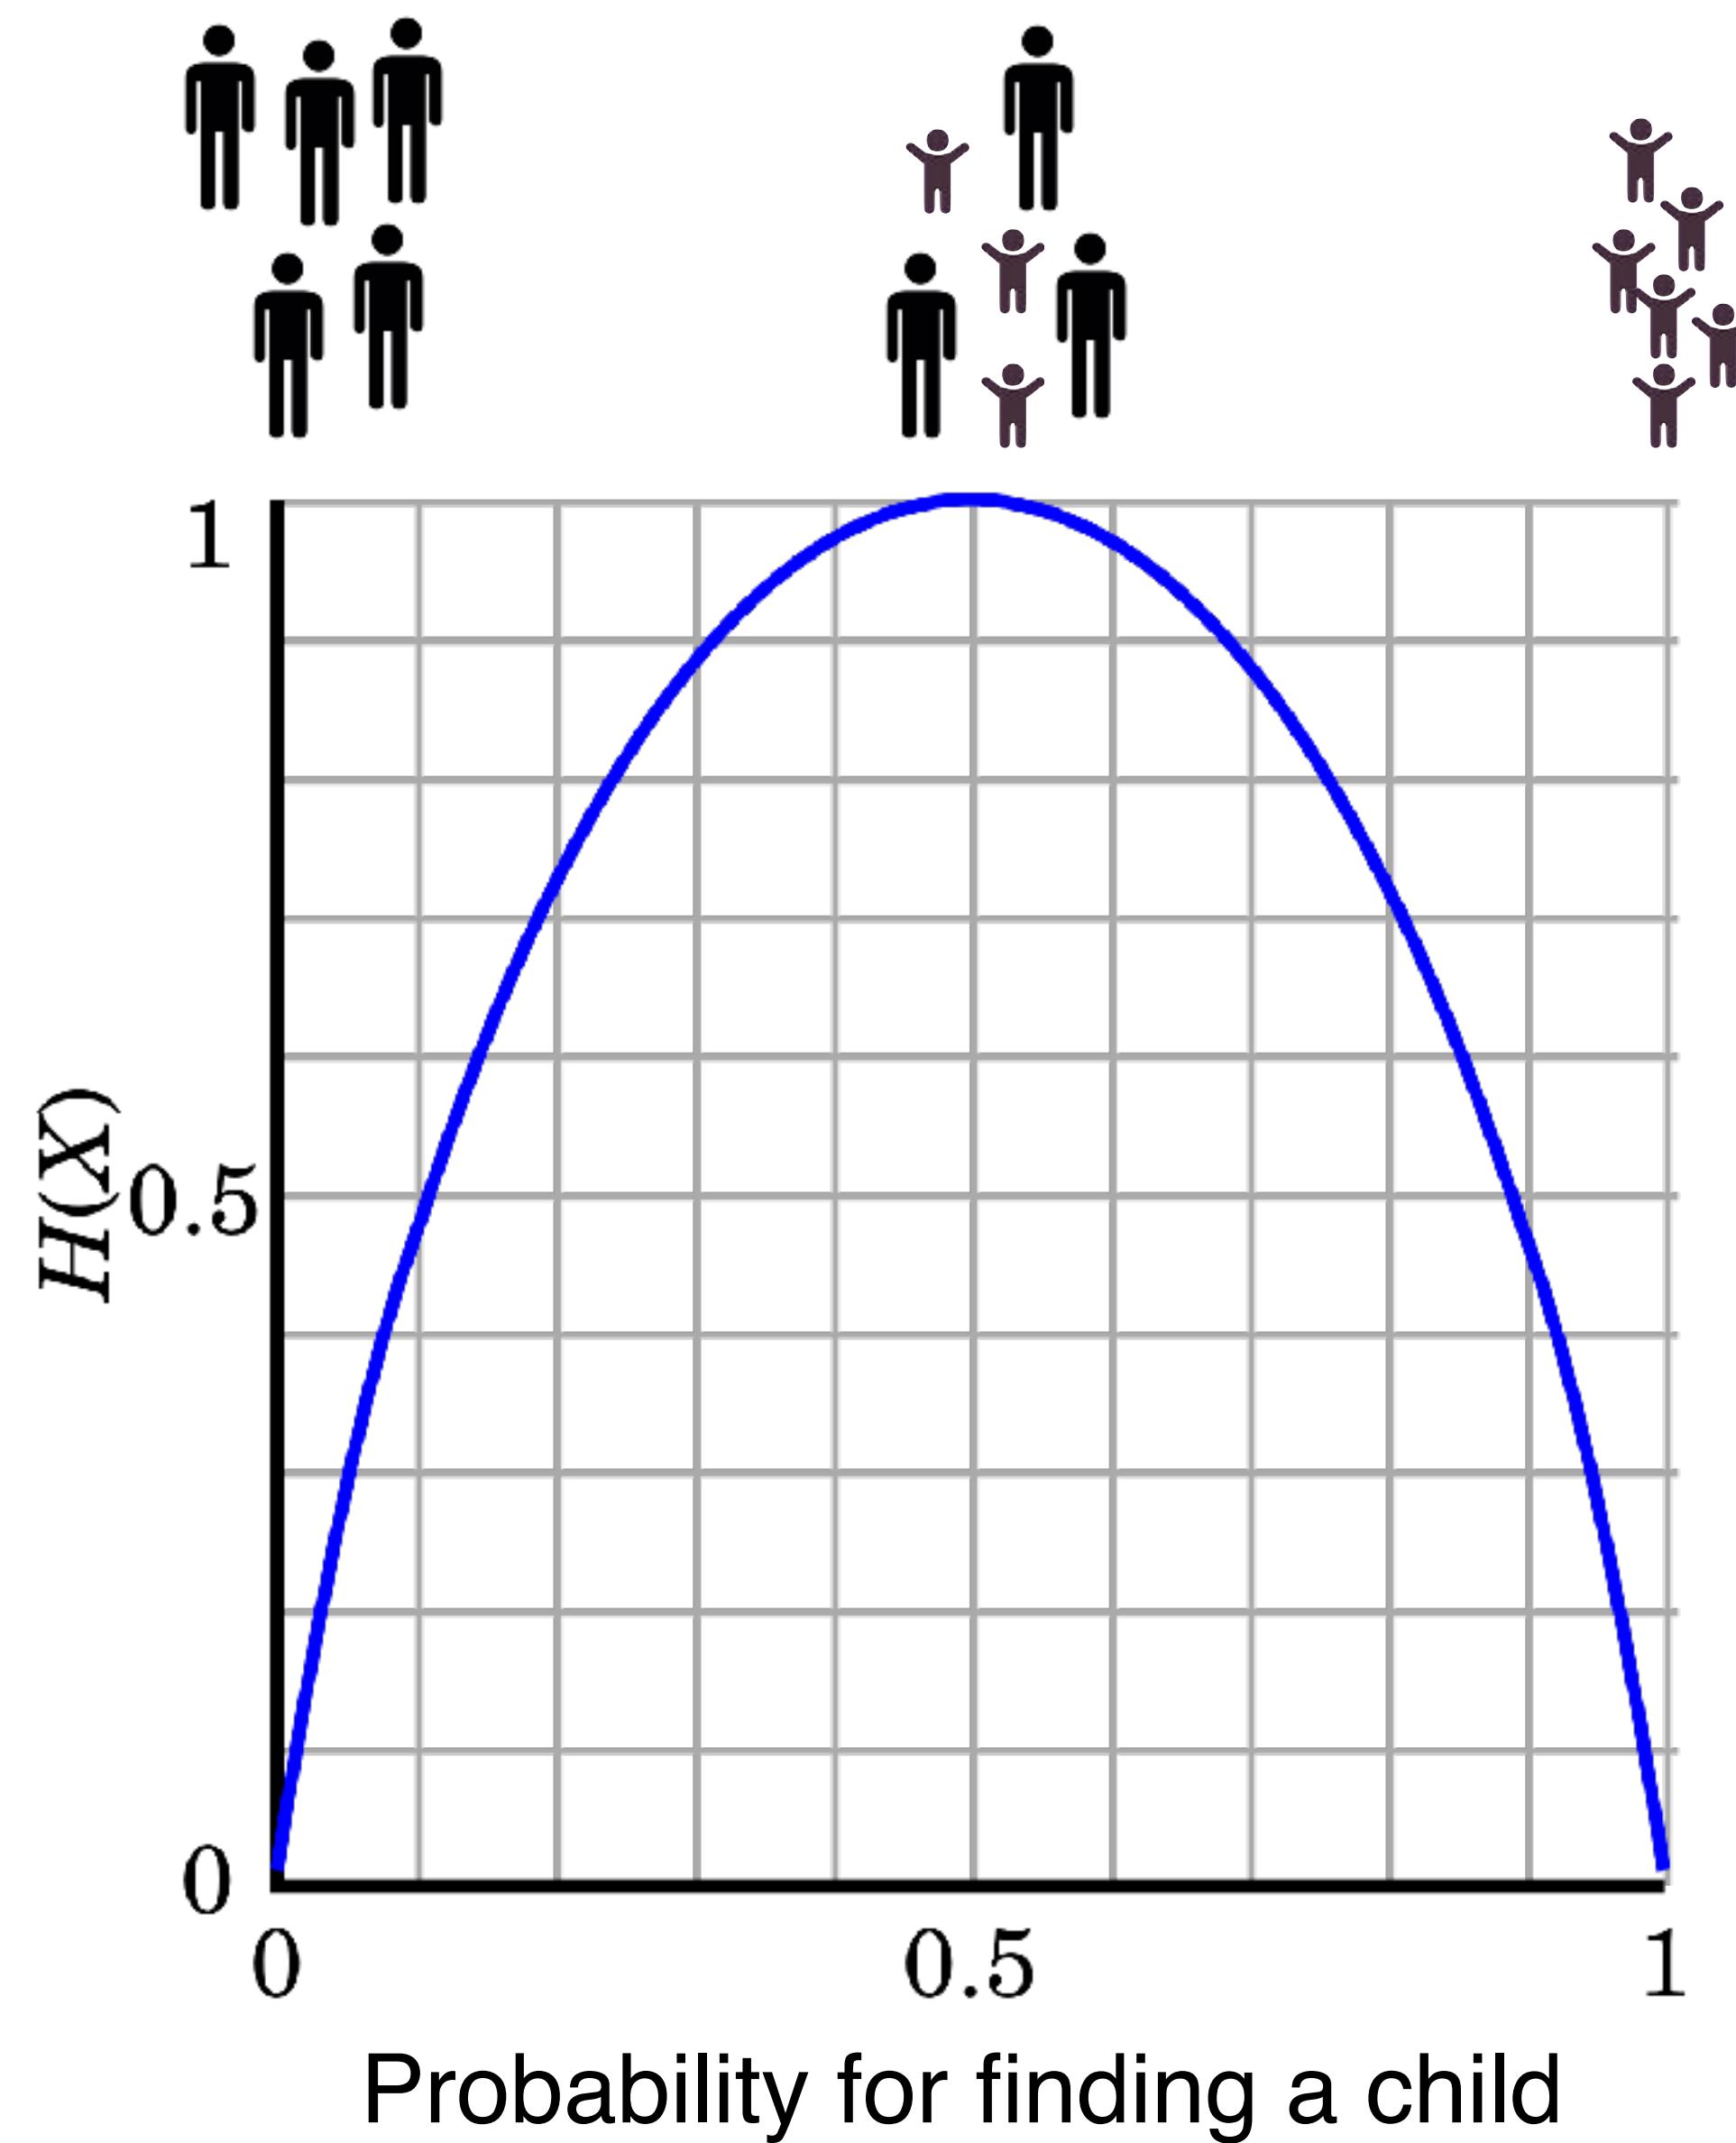
Example: Entropy of a binary data distribution



Here entropy is a measure of
balance or equality

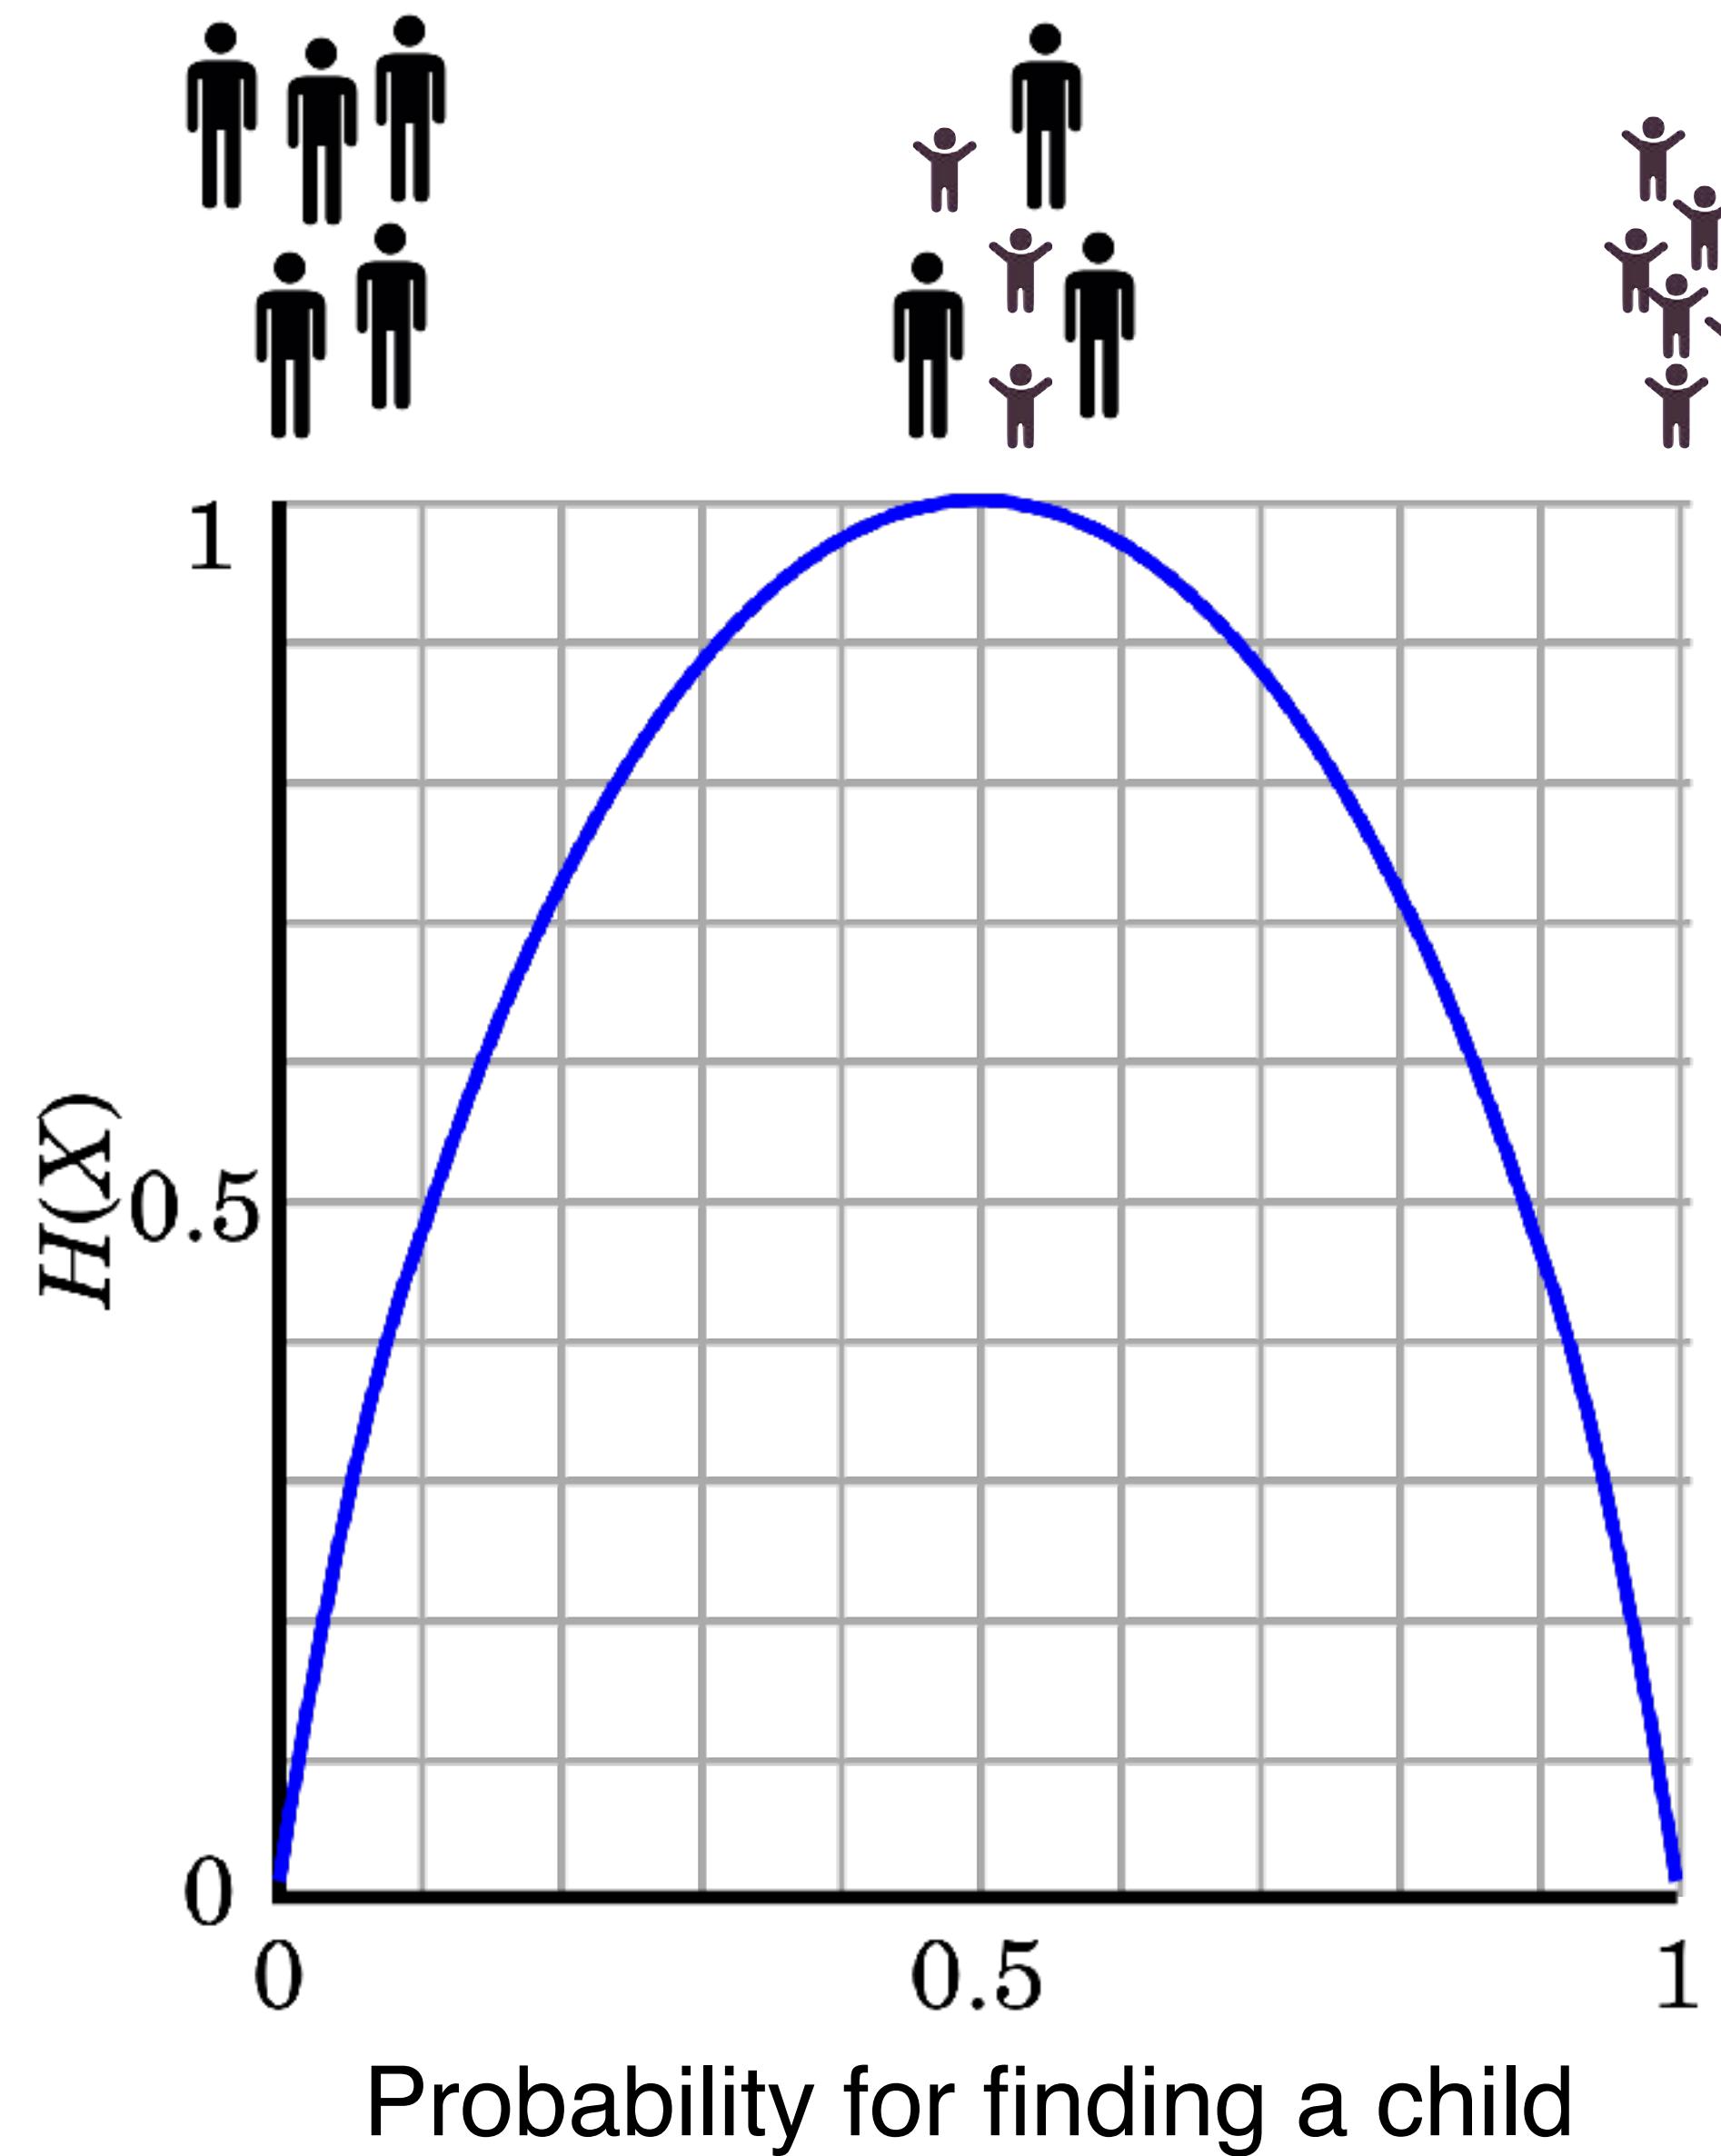
In machine learning: Impurity

Example: Entropy of a binary data distribution



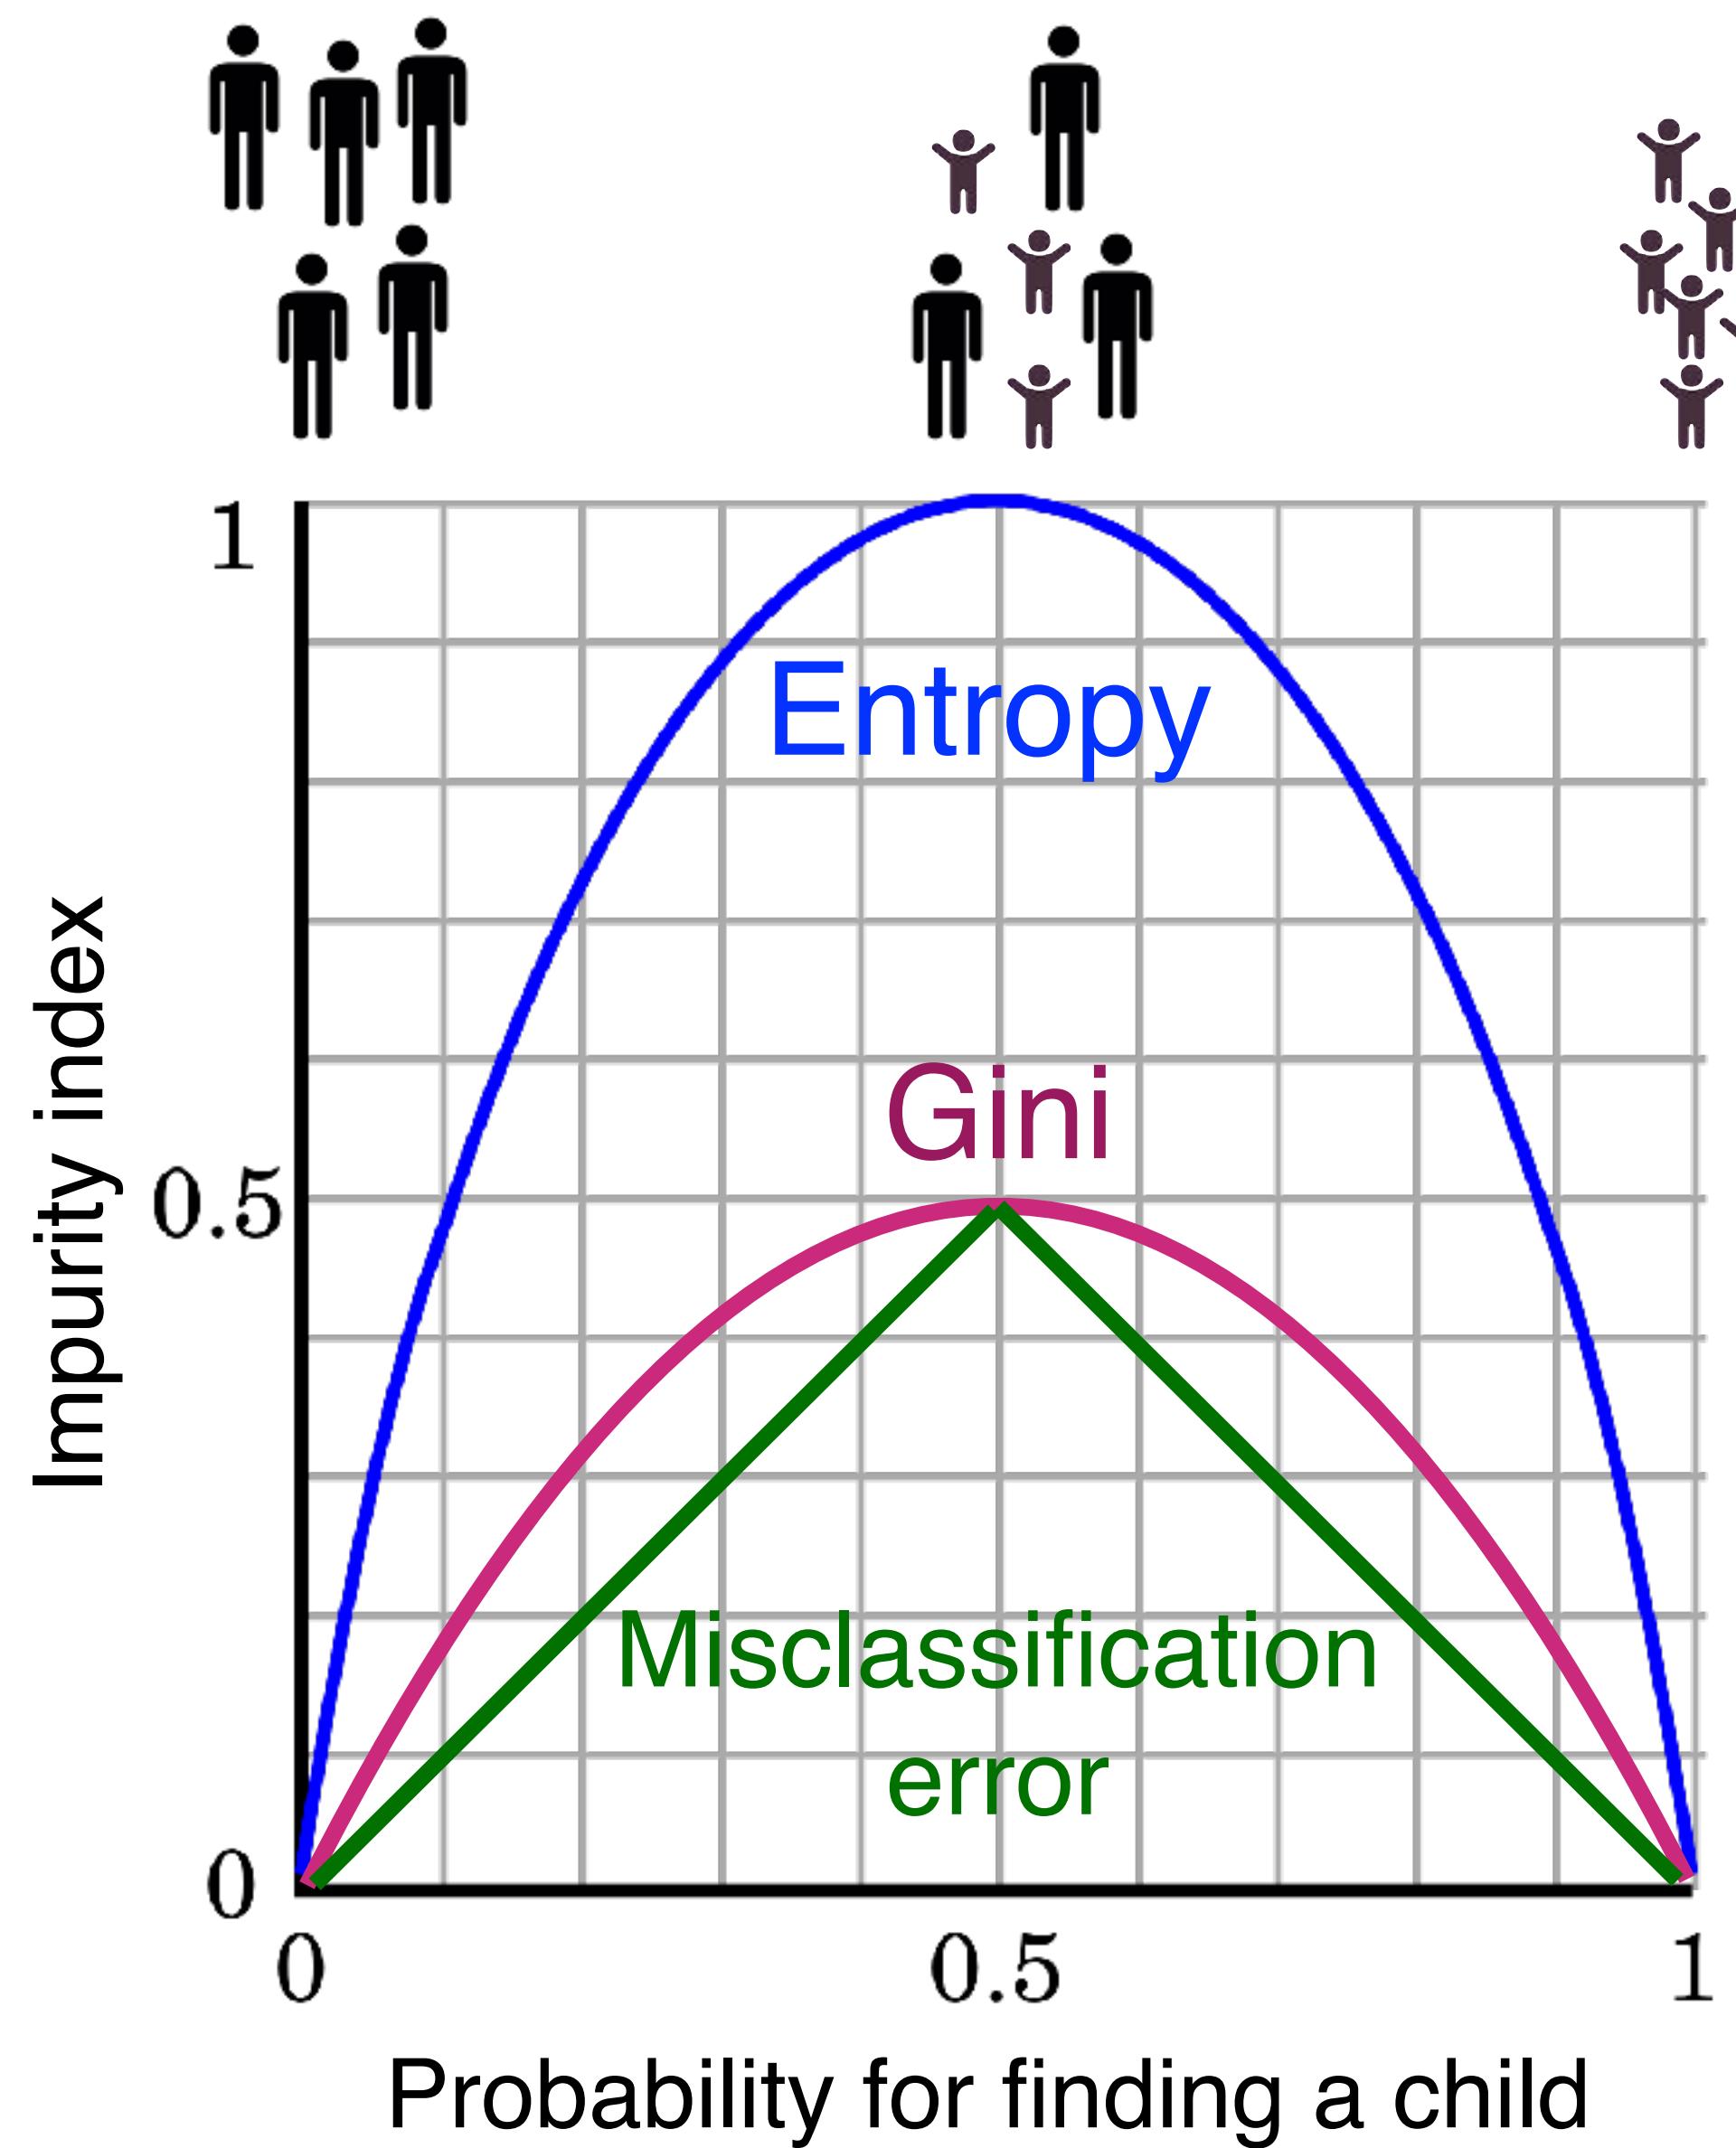
We could choose the age limit
for "child" to maximize impurity..
.. to give us the optimal strategy
to minimize yes/no questions

Example: Maximize Entropy for optimal Bar Khokba strategy



We get an optimal strategy by choosing questions that divide the yes/no outcomes equiprobably (=maximizing entropy)

There are other impurity measures



Impurity measures are used in decision tree classification

https://www.youtube.com/watch?v=_L39rN6gz7Y

unary

decimal

binary

hexadecimal

bits & bytes

Number systems, counting in binary

Positional systems:

unary

We write

decimal

$$a_n a_{n-1} a_{n-2} \dots a_0$$

binary

to represent

hexadecimal

$$a_n b^n + a_{n-1} b^{n-1} + a_{n-2} b^{n-2} + \dots + a_0 b^0$$

in base b with the digits

bits & bytes

$$a \in \{0, \dots, b-1\}$$

Memory to store events

Blackboard

coin flip

die roll

Data Science Setup

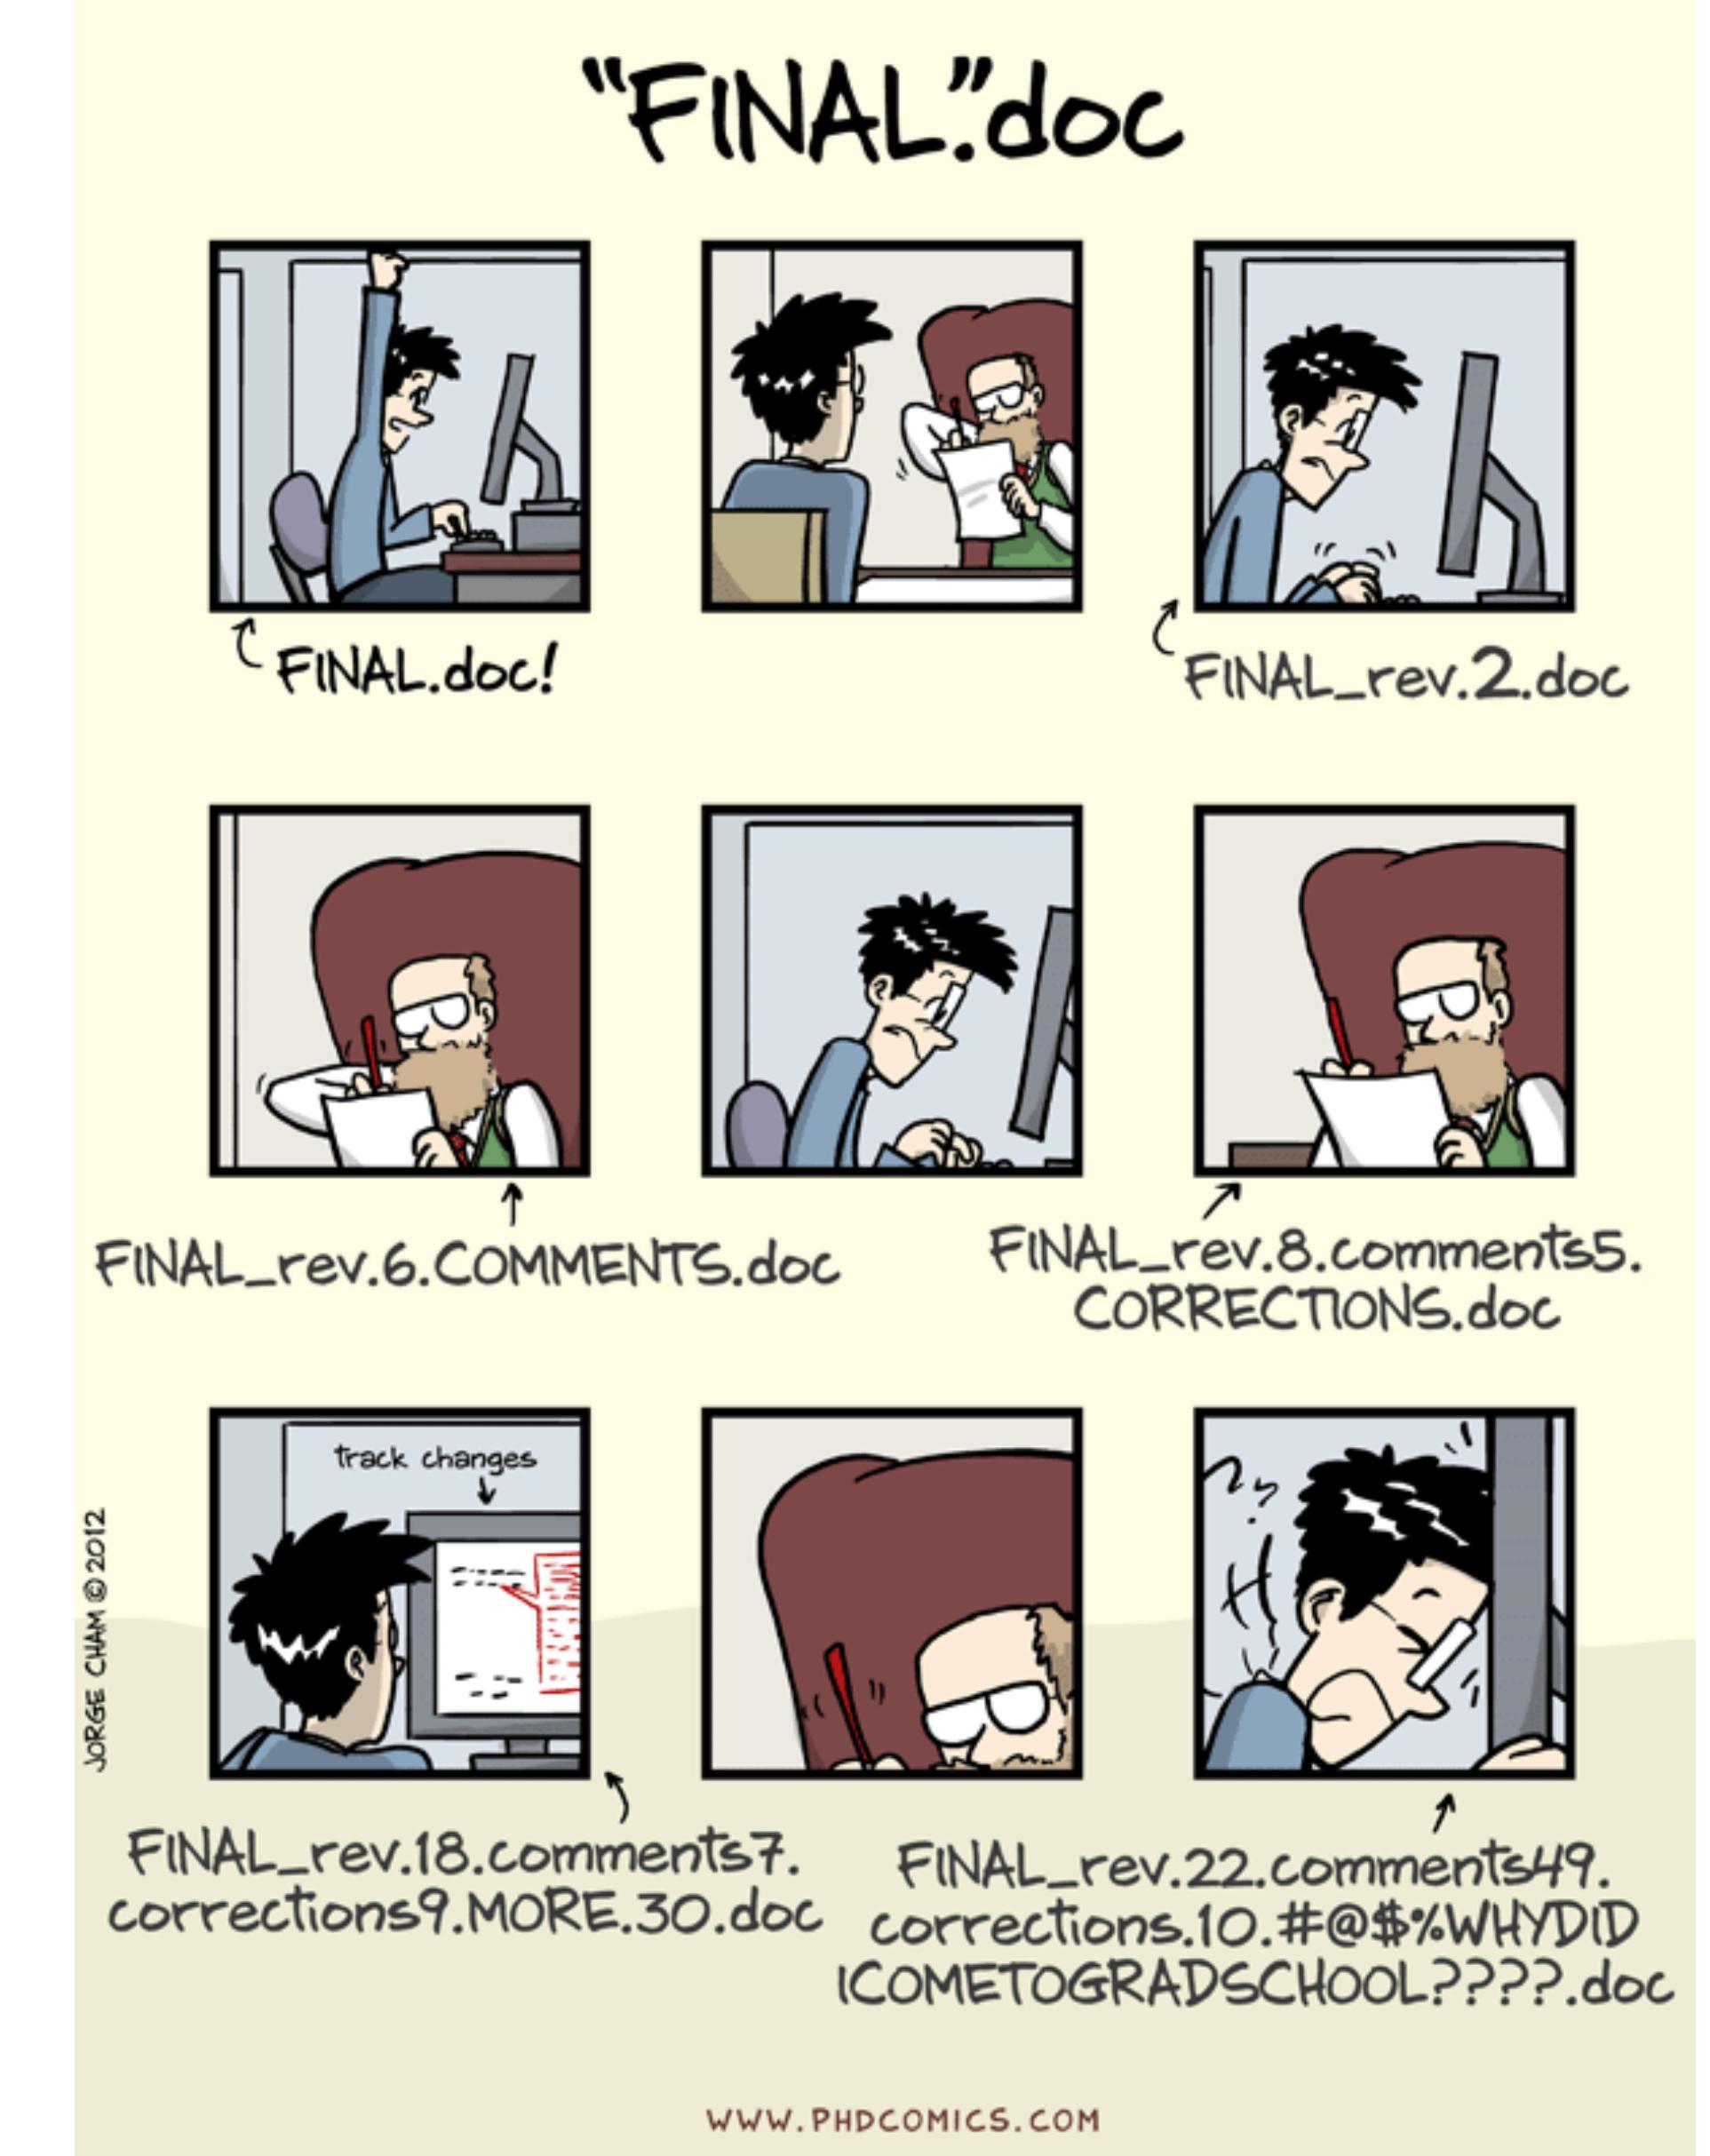
Practical tips

Get acquainted with git / GitHub

Git is a distributed version-control system for tracking changes in any set of files, originally designed to coordinate group work.



Similar systems: CSV, SVN, Mercurial



Get acquainted with git / GitHub

GitHub, Inc. is a subsidiary of Microsoft which provides hosting for software development and version control using Git



GitLab or CodeBerg are alternatives



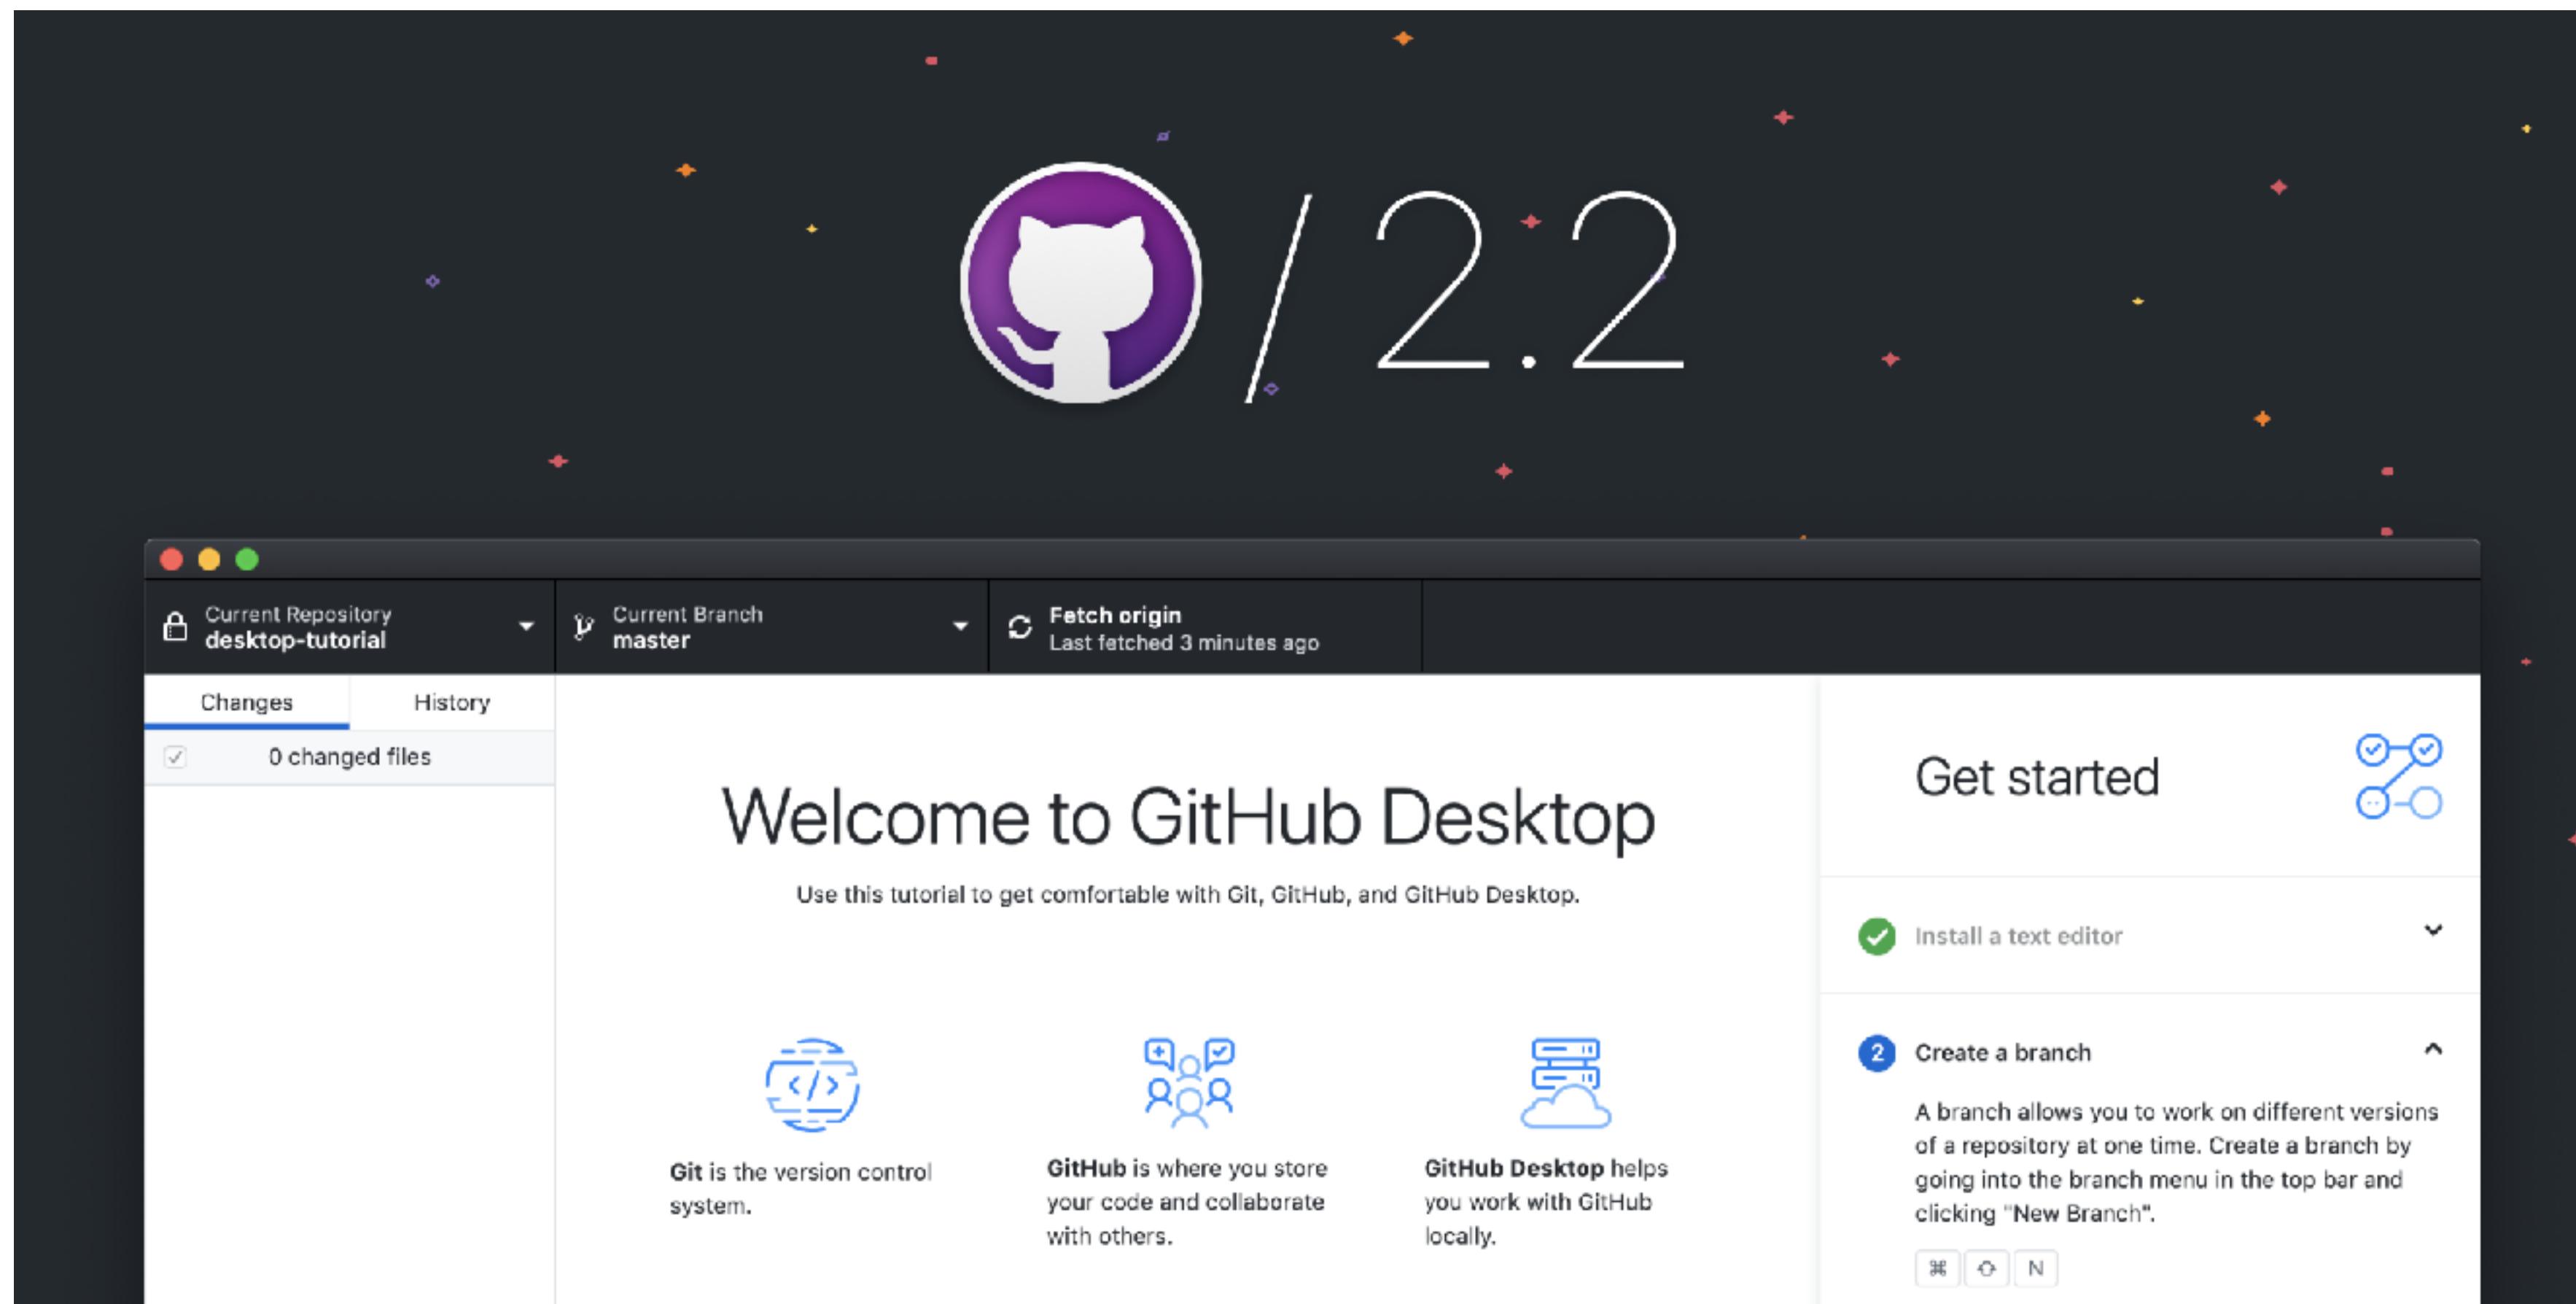
GitLab



Codeberg

Get acquainted with git / GitHub

GitHub Desktop is a Graphical User Interface (GUI) for Github



Before we delve into a project, you MUST get organized!

```
├── LICENSE
├── Makefile      <- Makefile with commands like `make data` or `make train`
├── README.md     <- The top-level README for developers using this project.
└── data
    ├── external   <- Data from third party sources.
    ├── interim    <- Intermediate data that has been transformed.
    ├── processed   <- The final, canonical data sets for modeling.
    └── raw         <- The original, immutable data dump.

── docs          <- A default Sphinx project; see sphinx-doc.org for details

── models        <- Trained and serialized models, model predictions, or model summaries

── notebooks     <- Jupyter notebooks. Naming convention is a number (for ordering),
                    the creator's initials, and a short '-' delimited description, e.g.
                    `1.0-jqp-initial-data-exploration`.

── references    <- Data dictionaries, manuals, and all other explanatory materials.

── reports       <- Generated analysis as HTML, PDF, LaTeX, etc.
    └── figures    <- Generated graphics and figures to be used in reporting

── requirements.txt <- The requirements file for reproducing the analysis environment, e.g.
                      generated with `pip freeze > requirements.txt`

── src
    ├── __init__.py  <- Source code for use in this project.
    │   <- Makes src a Python module
    ├── data         <- Scripts to download or generate data
    │   └── make_dataset.py
    ├── features     <- Scripts to turn raw data into features for modeling
    │   └── build_features.py
    ├── models       <- Scripts to train models and then use trained models to make
                        predictions
    │   ├── predict_model.py
    │   └── train_model.py
    └── visualization <- Scripts to create exploratory and results oriented visualizations
        └── visualize.py

── tox.ini        <- tox file with settings for running tox; see tox.readthedocs.io
```

Folder structure is not a science, but based on years of experience

Cookiecutter Data Science

<https://github.com/drivendata/cookiecutter-data-science>

Before we delve into a project, you MUST get organized!

```
├── LICENSE  
├── Makefile      <- Makefile with commands like `make data` or `make train`  
├── README.md     <- The top-level README for developers using this project.  
└── data           
    ├── external    <- Data from third party sources.  
    ├── interim     <- Intermediate data that has been transformed.  
    ├── processed   <- The final, canonical data sets for modeling.  
    └── raw         <- The original, immutable data dump.  
  
── docs          <- A default Sphinx project; see sphinx-doc.org for details  
  
── models        <- Trained and serialized models, model predictions, or model summaries  
  
── notebooks     <- Jupyter notebooks. Naming convention is a number (for ordering),  
                  the creator's initials, and a short '-' delimited description, e.g.  
                  `1.0-jqp-initial-data-exploration`.  
  
── references    <- Data dictionaries, manuals, and all other explanatory materials.  
  
── reports       <- Generated analysis as HTML, PDF, LaTeX, etc.  
    └── figures     <- Generated graphics and figures to be used in reporting  
  
── requirements.txt <- The requirements file for reproducing the analysis environment, e.g.  
                      generated with `pip freeze > requirements.txt`  
  
── src           <- Source code for use in this project.  
    ├── __init__.py  <- Makes src a Python module  
    └── data         <- Scripts to download or generate data  
        └── make_dataset.py  
  
    ├── features     <- Scripts to turn raw data into features for modeling  
        └── build_features.py  
  
    ├── models       <- Scripts to train models and then use trained models to make  
                      predictions  
        └── predict_model.py  
        └── train_model.py  
  
    └── visualization <- Scripts to create exploratory and results oriented visualizations  
        └── visualize.py  
  
tox.ini      <- tox file with settings for running tox; see tox.readthedocs.io
```

MUST haves:

- README.md
- Data: raw, interim, processed
- Notebooks / Code
- References
- Reports / Figures

Cookiecutter Data Science

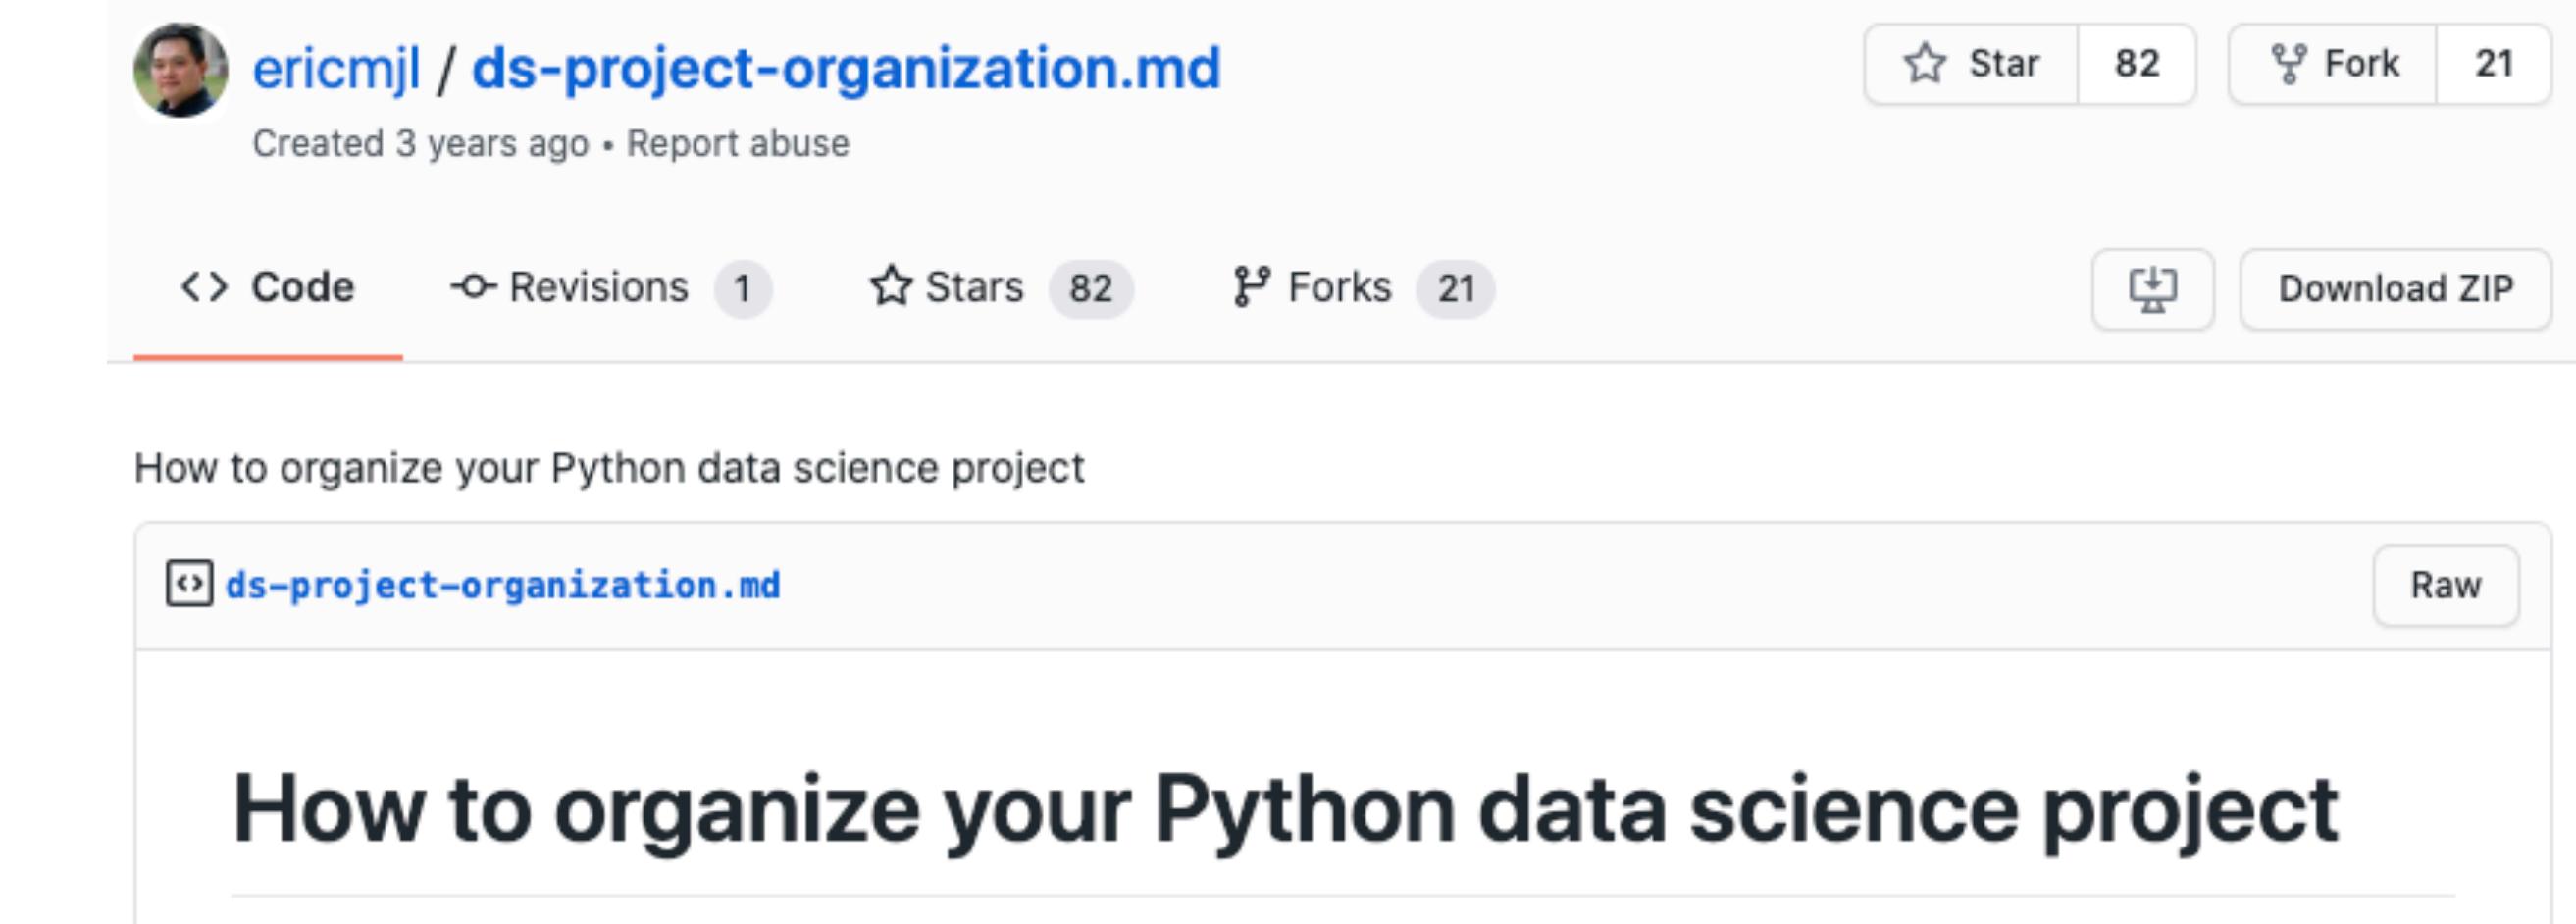
<https://github.com/drivendata/cookiecutter-data-science>

Before we delve into a project, you MUST get organized!

How to Create a Professional Github Data Science Repository

Good practices in repository structure, documenting jupyter notebooks, and writing an informative README

 Ahilan Srivishnumohan Jul 25, 2020 · 7 min read ★



The screenshot shows a GitHub repository page for the file `ds-project-organization.md`. The repository was created 3 years ago by `ericmjl`. It has 82 stars, 21 forks, and 1 revision. The page title is "How to organize your Python data science project". The file content is displayed in a code editor-like interface.

ericmjl / [ds-project-organization.md](#)

Created 3 years ago · Report abuse

Code Revisions 1 Stars 82 Forks 21

How to organize your Python data science project

[ds-project-organization.md](#) Raw

How to organize your Python data science project

<https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510>

<https://towardsdatascience.com/how-to-create-a-professional-github-data-science-repository-84e9607644a2>

Data is usually NOT stored on Github

Github is for code and work documents, not data.

There are hard limitations: 100 MB files, 2 GB push, ~5 GB repo

If you have hundreds of MB data, add the data folder to .gitignore

Data is usually NOT stored on Github

Github is for code and work documents, not data.

There are hard limitations: 100 MB files, 2 GB push, ~5 GB repo

If you have hundreds of MB data, add the data folder to .gitignore

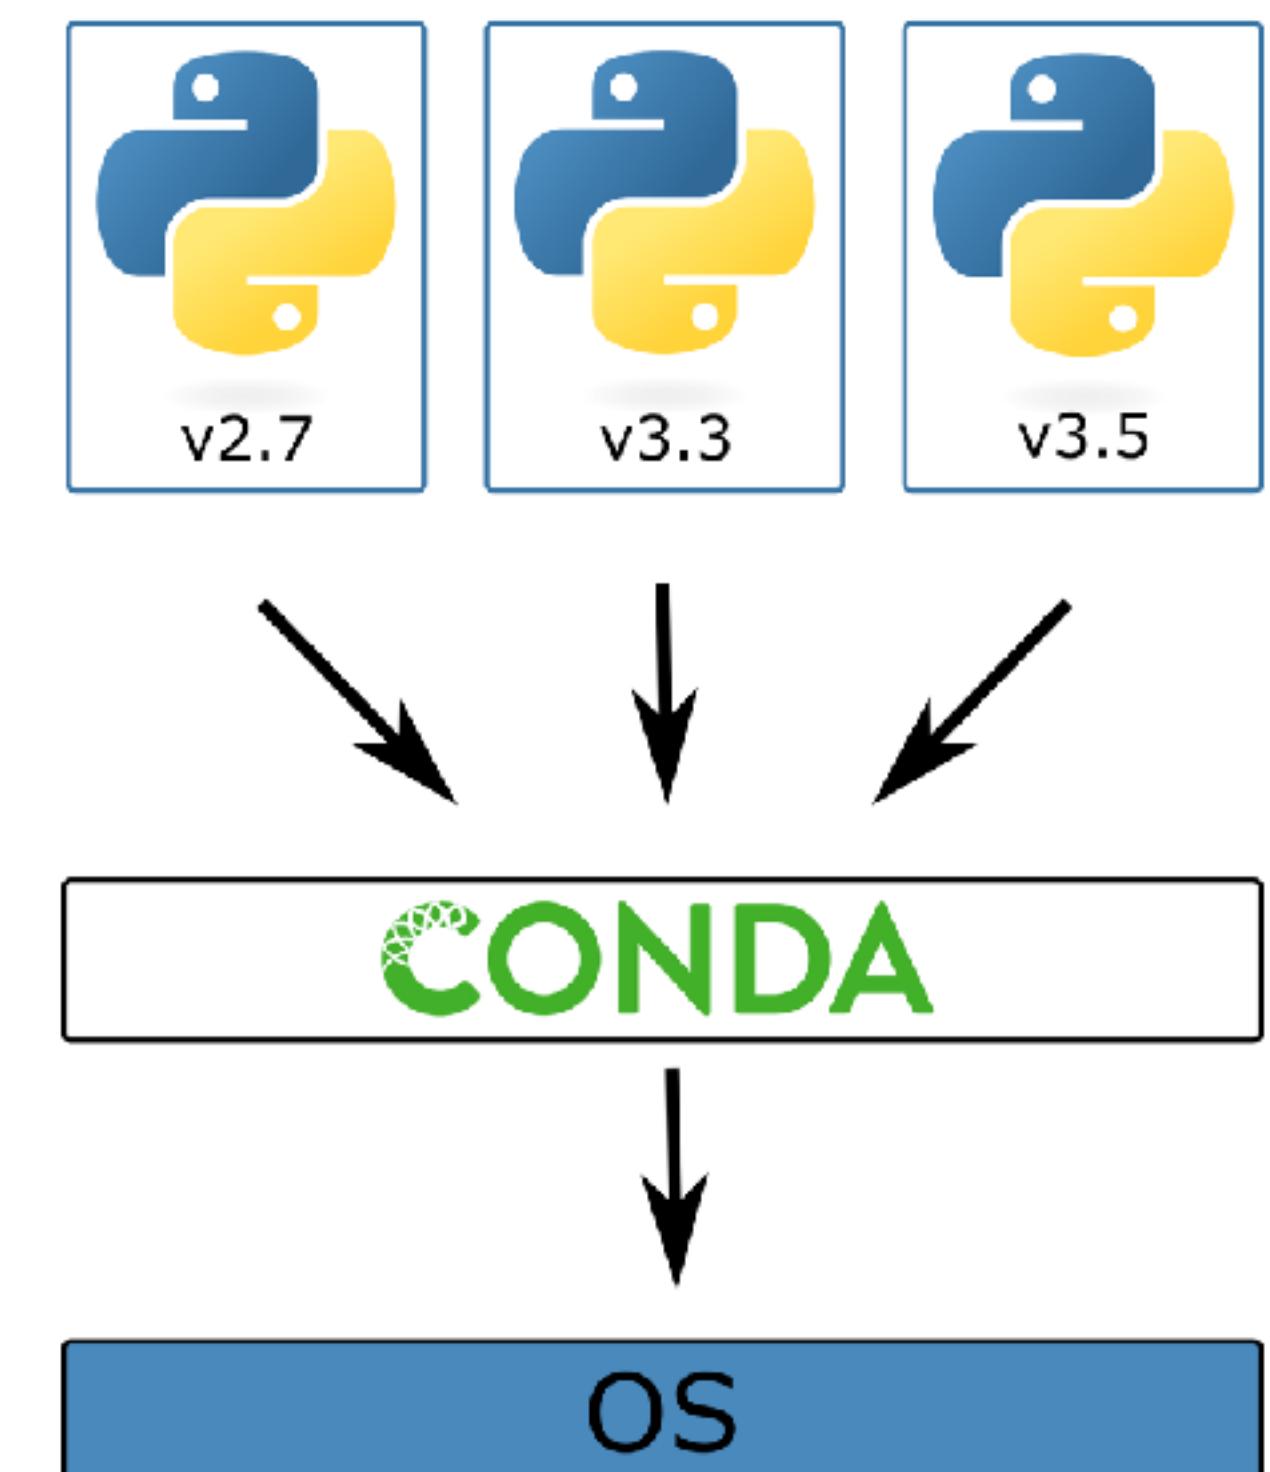
To make projects reproducible,
we store data here, for example:



Advanced organization

Often we use a **virtual environment** (venv) for each project

Why?



Advanced organization

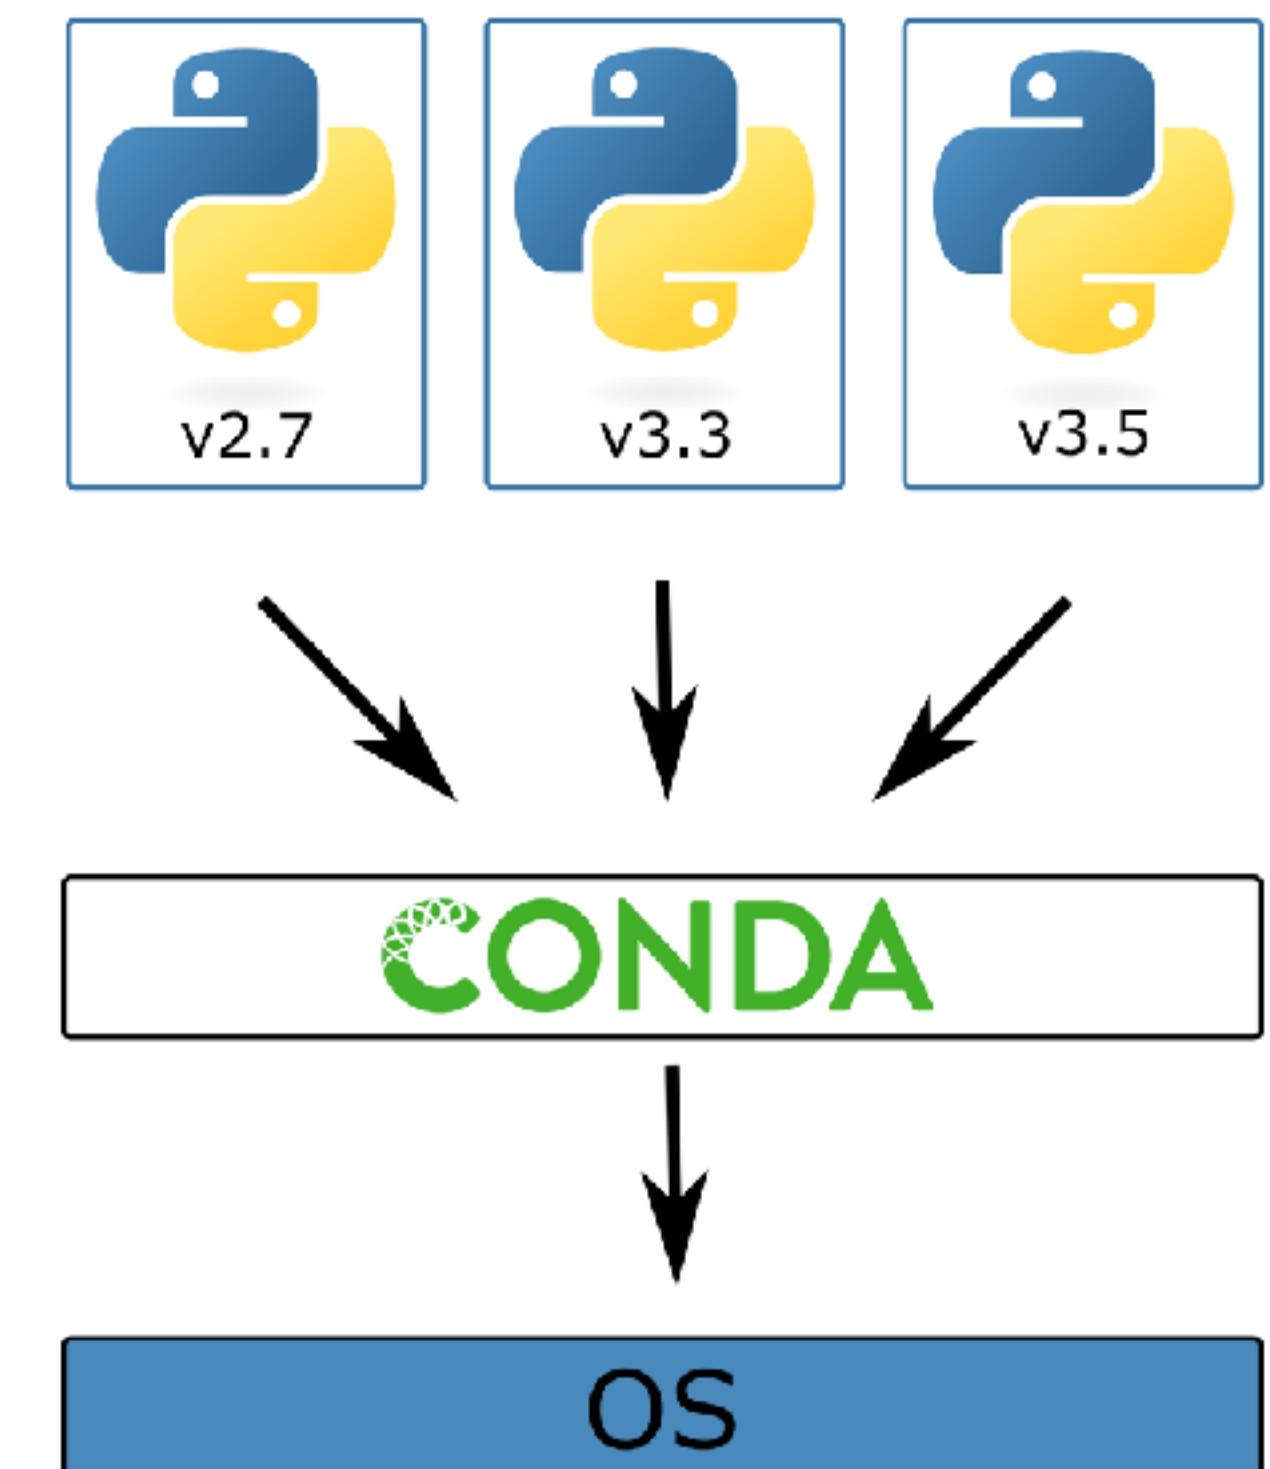
Often we use a **virtual environment** (venv) for each project

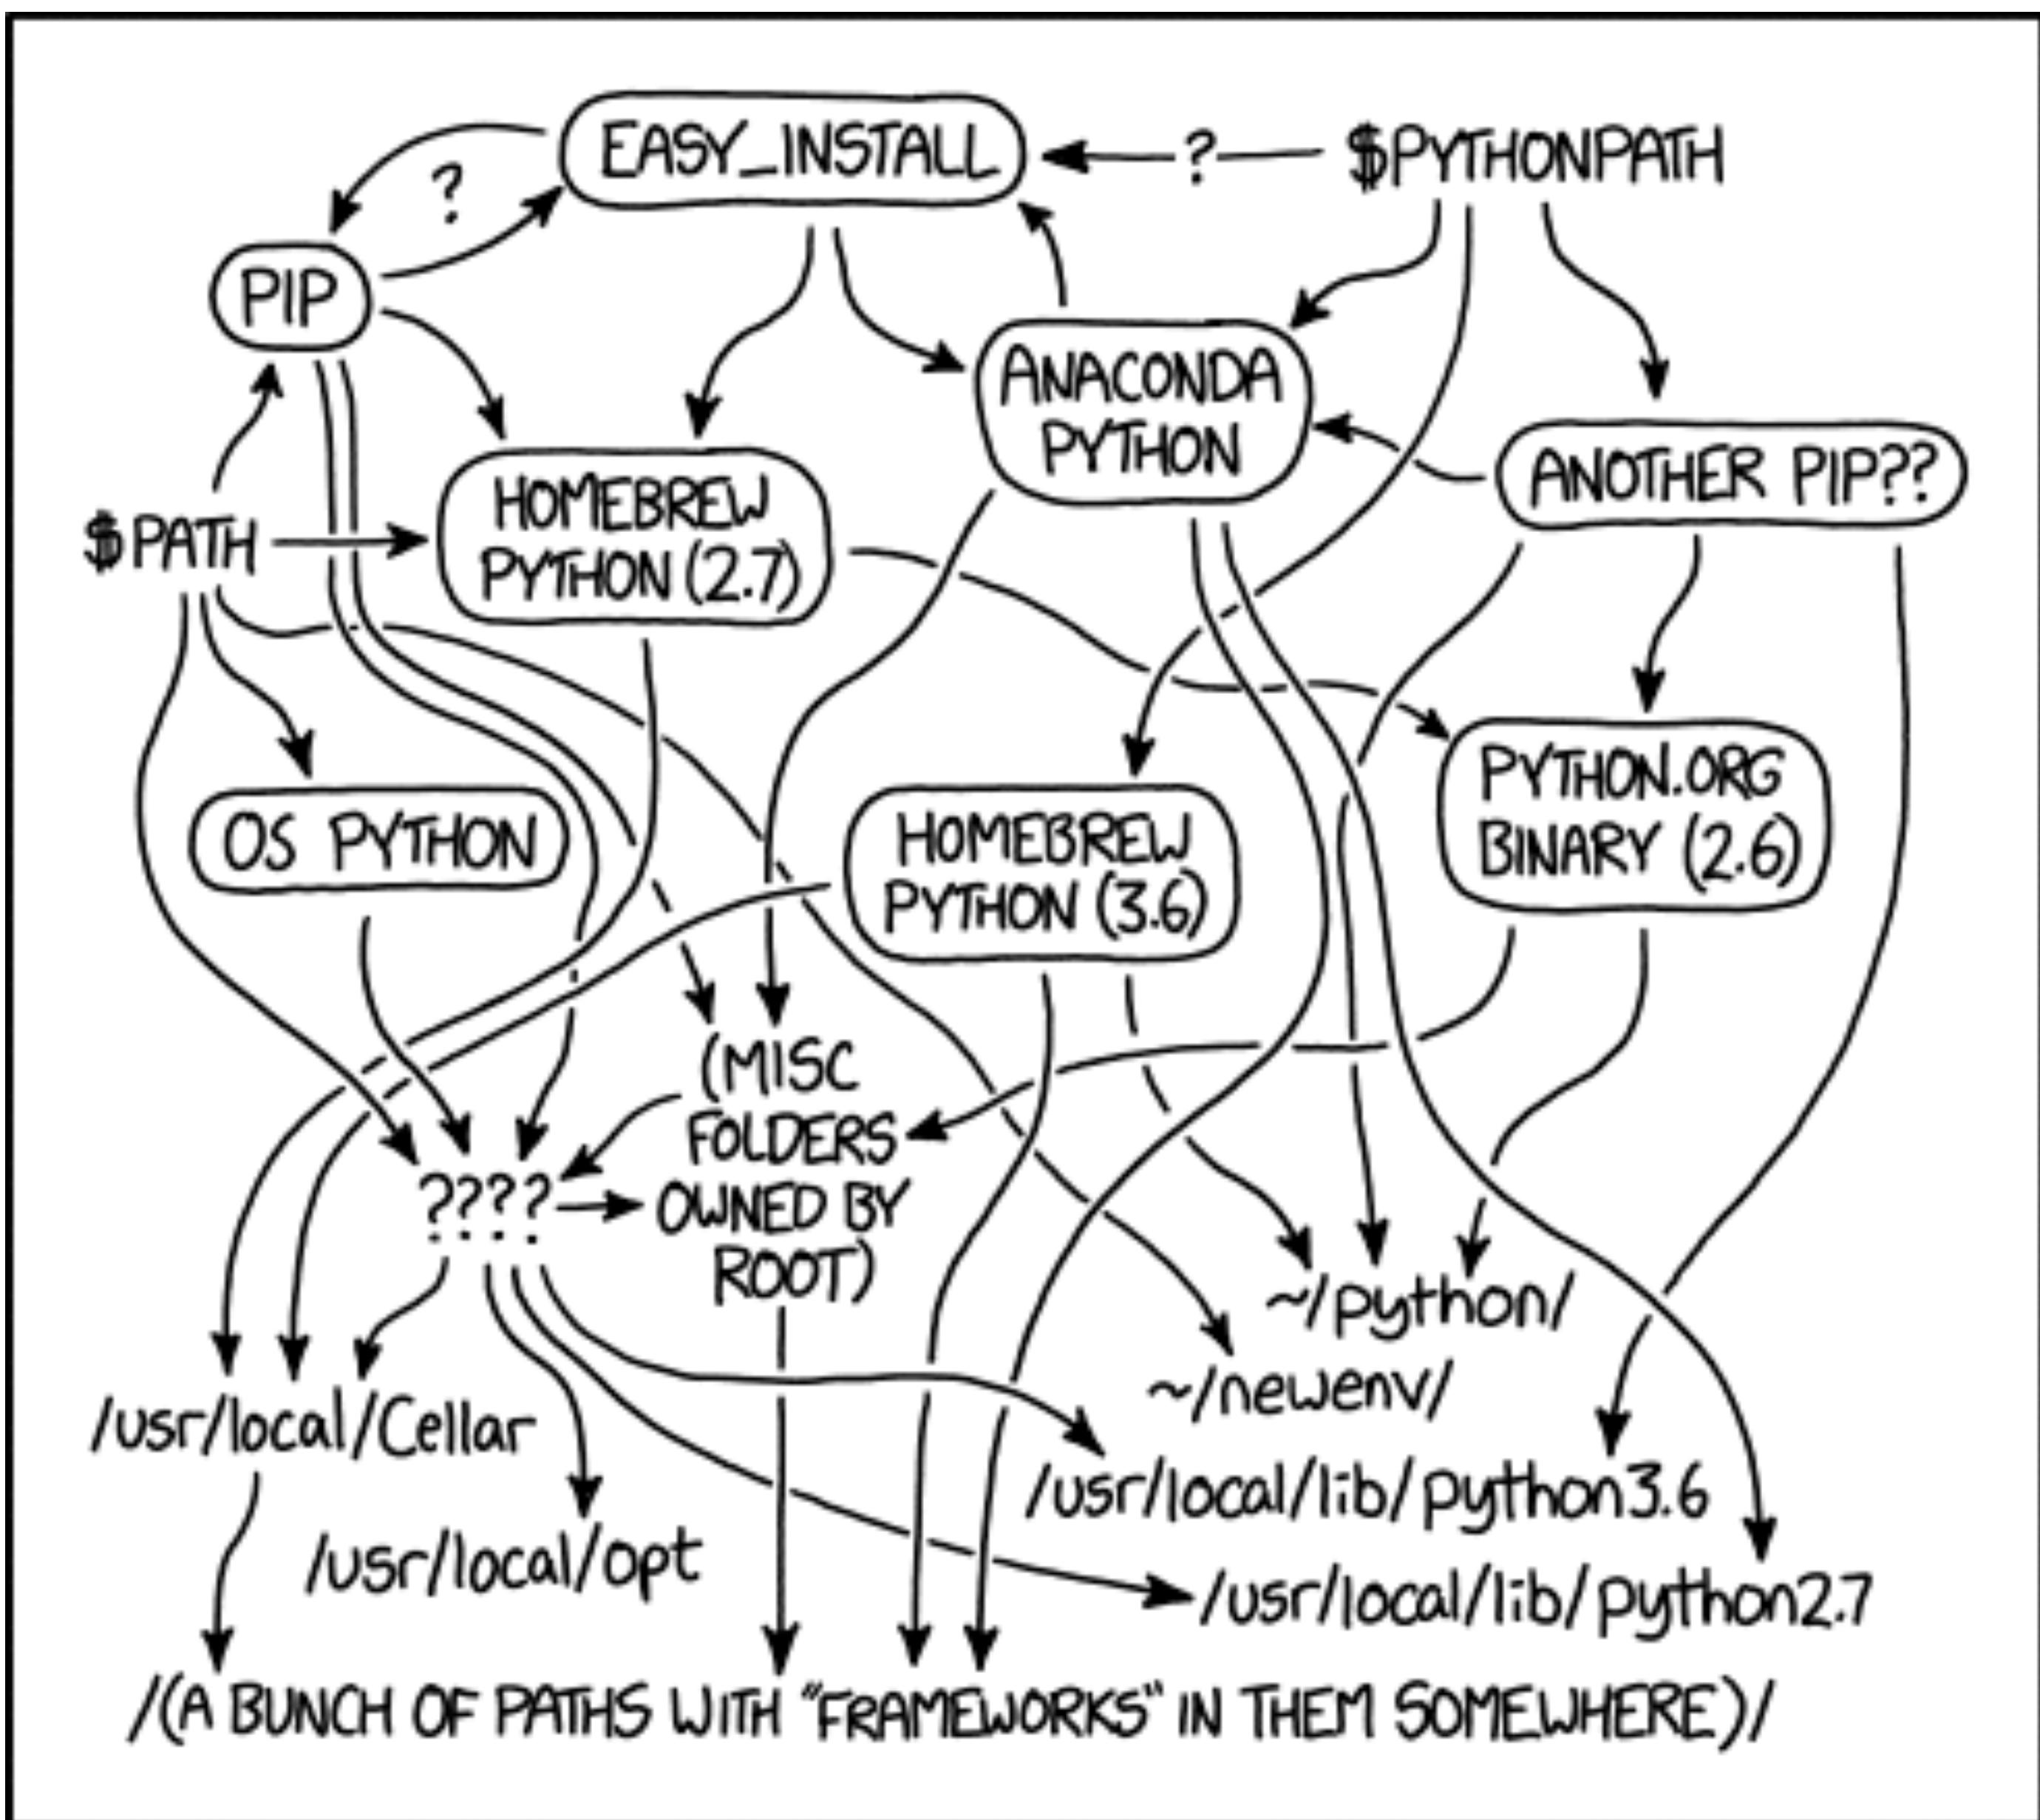
Why?

1) Reproducibility:

- Python changes all the time
- Python packages change all the time
- requirements.txt differ between projects

2) To not mess up your Python install





MY PYTHON ENVIRONMENT HAS BECOME SO DEGRADED
THAT MY LAPTOP HAS BEEN DECLARED A SUPERFUND SITE.

Advanced organization

```
requirements.txt      *  
1  matplotlib>=3.3.3  
2  numpy>=1.19.4  
3  pandas>=1.0.3  
4  pyproj>=2.6.1.post1  
5  geojson>=2.5.0  
6  shapely>=1.7.0  
7  csv>=1.0  
8  networkx>=2.5  
9  igraph>=0.8.3  
10 fiona>=1.8.18  
11 osmnx==0.16.2  
12 geopandas>=0.8.1  
13 tqdm>=4.55.0  
14 haversine>=2.3.0
```

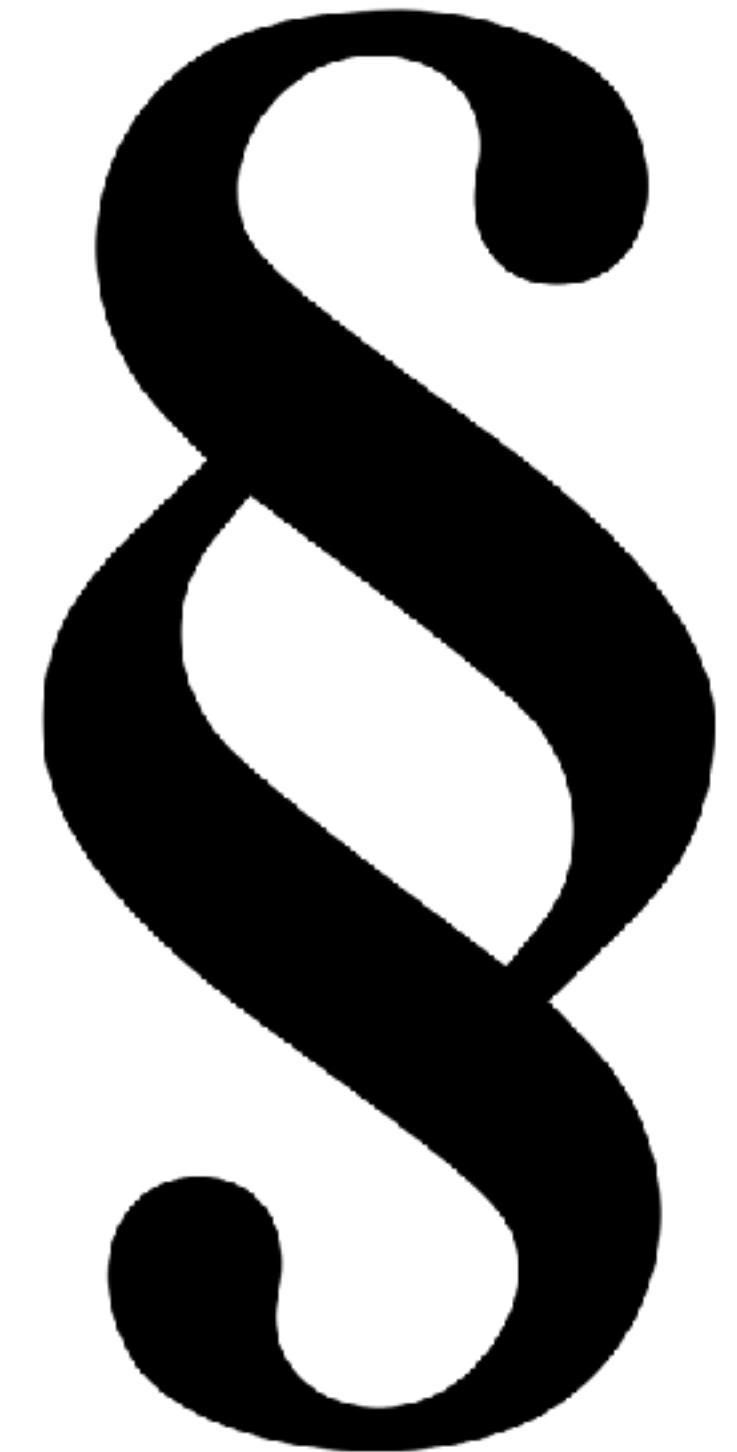
```
pip freeze > requirements.txt
```

Advanced organization

LICENSE

If you publish your project, the license is the legal document telling others how they can use it.

No serious company will ever risk building on your code if you don't have a license.



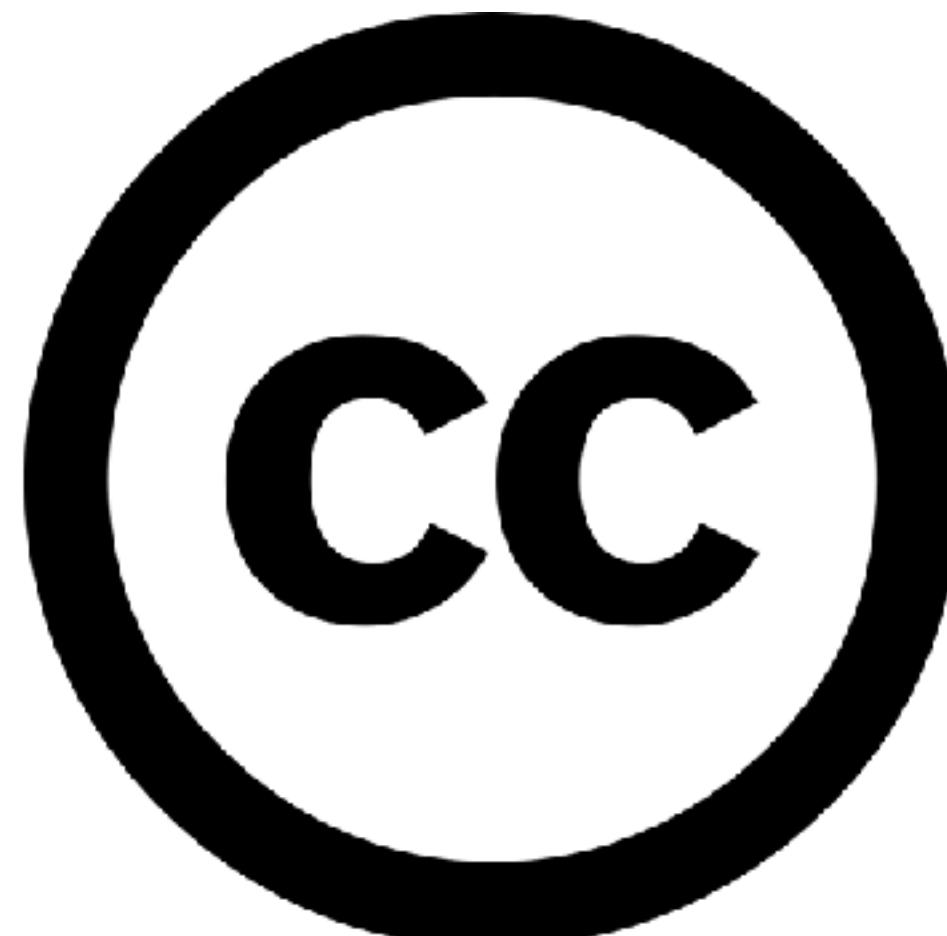
Advanced organization

LICENSE

If you publish your project, the license is the legal document telling others how they can use it.

Choose an established license,
don't write it yourself.

Consider copyleft and creative
commons to protect freedoms.



<https://creativecommons.org/>

<https://www.gnu.org/philosophy/free-sw.en.html>

<https://en.wikipedia.org/wiki/JSLint#License>

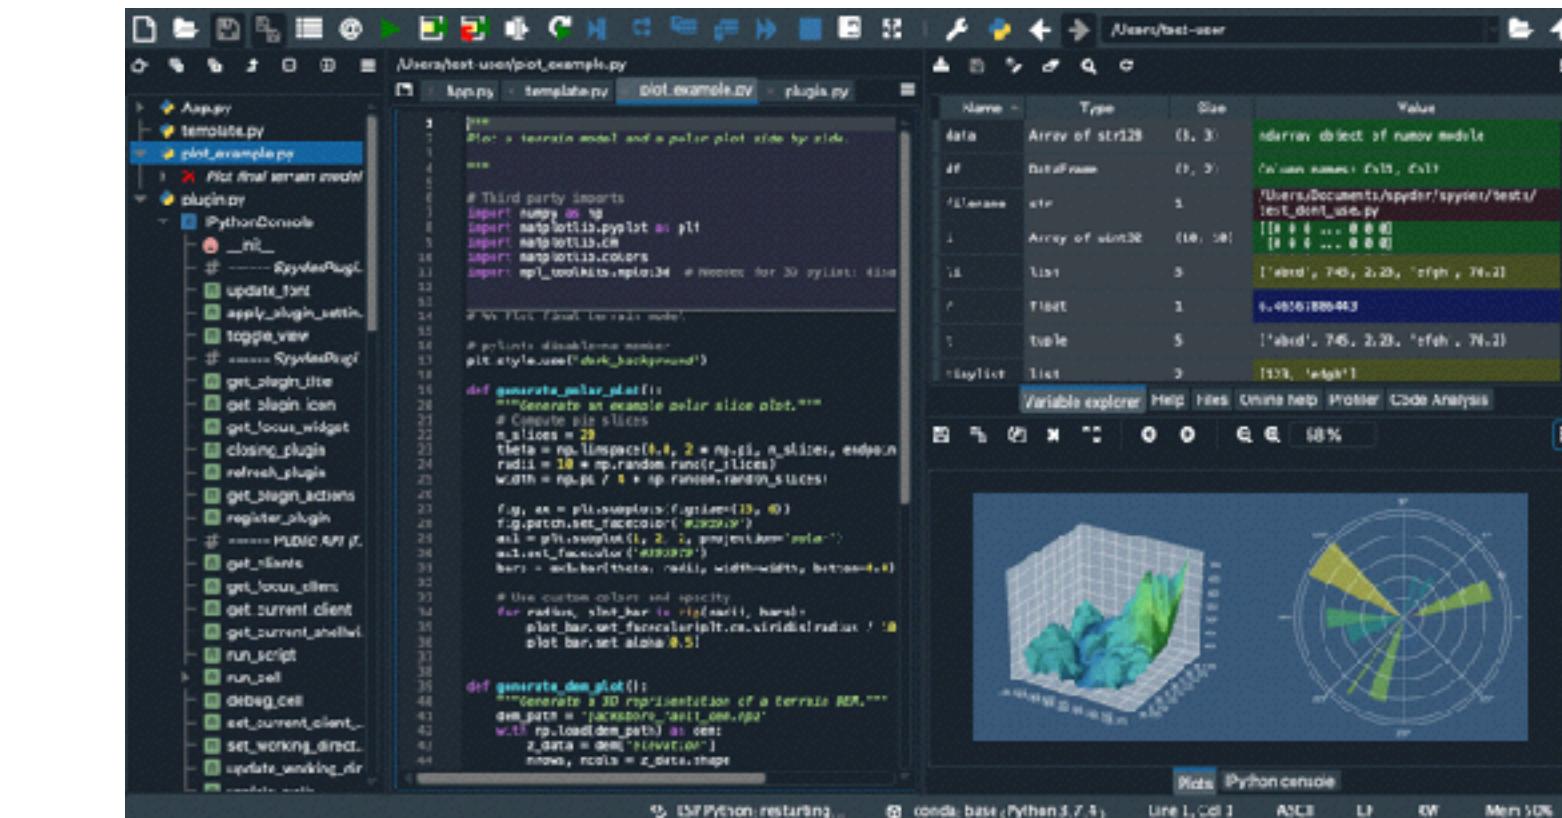
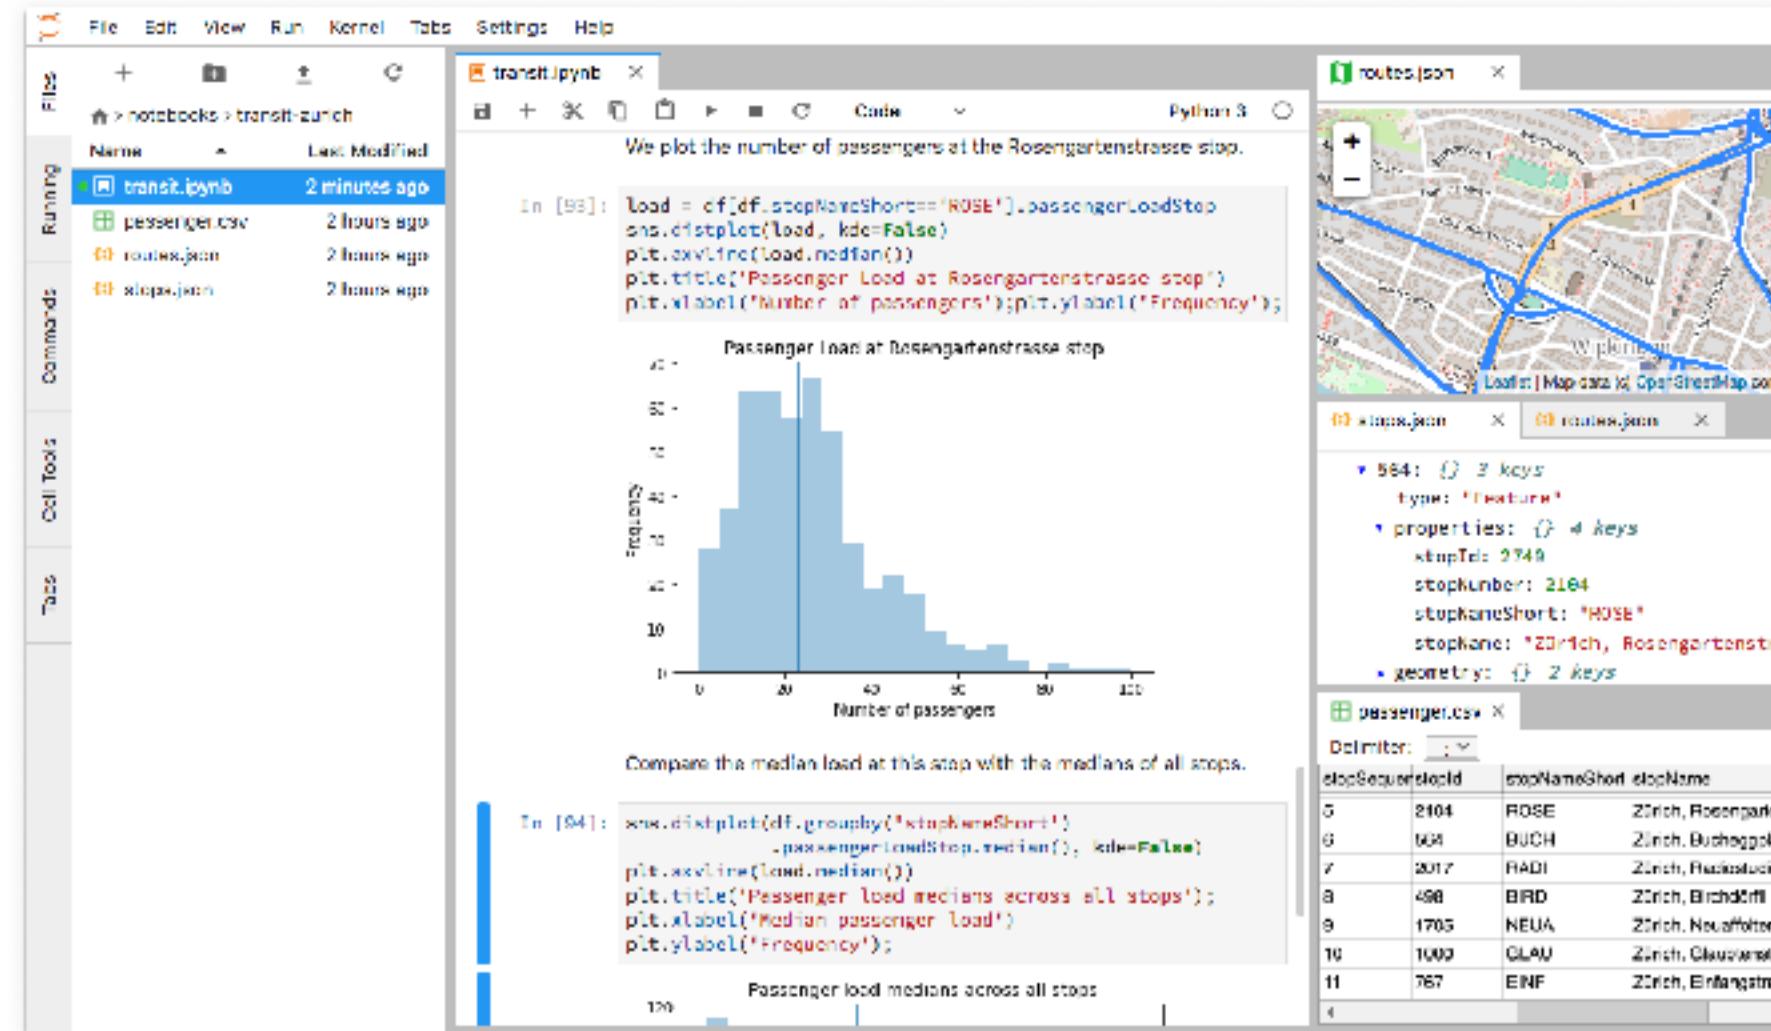
Advanced organization

DOCS

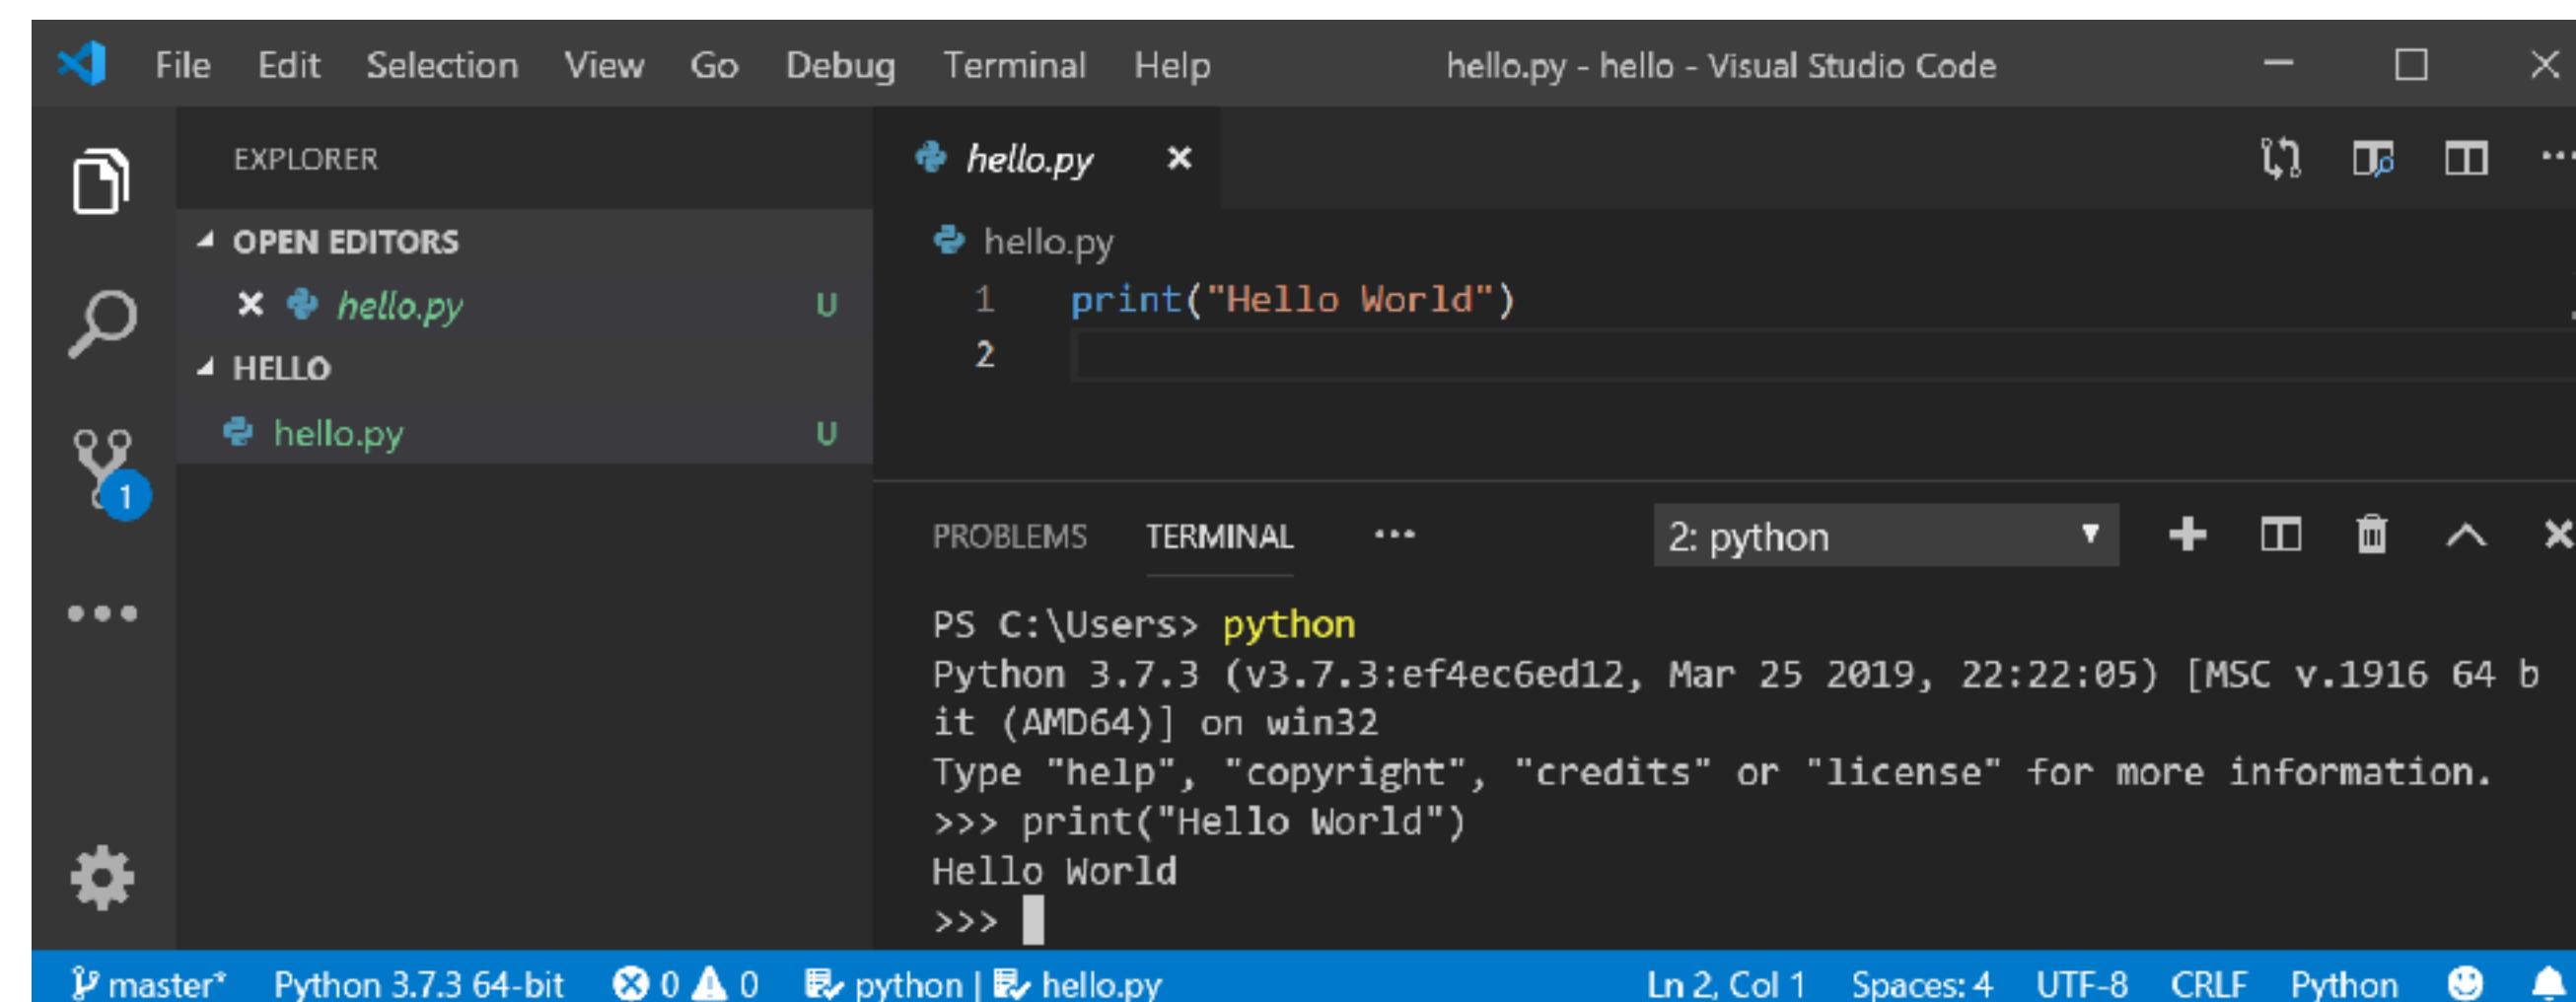
If you want others to use and build on your code, you need excellent documentation. It makes or breaks its success.

Expect investing >20% of the whole project time on docs.

Advanced organization: IDEs (integrated development environments)



Visual Studio Code



Clustering

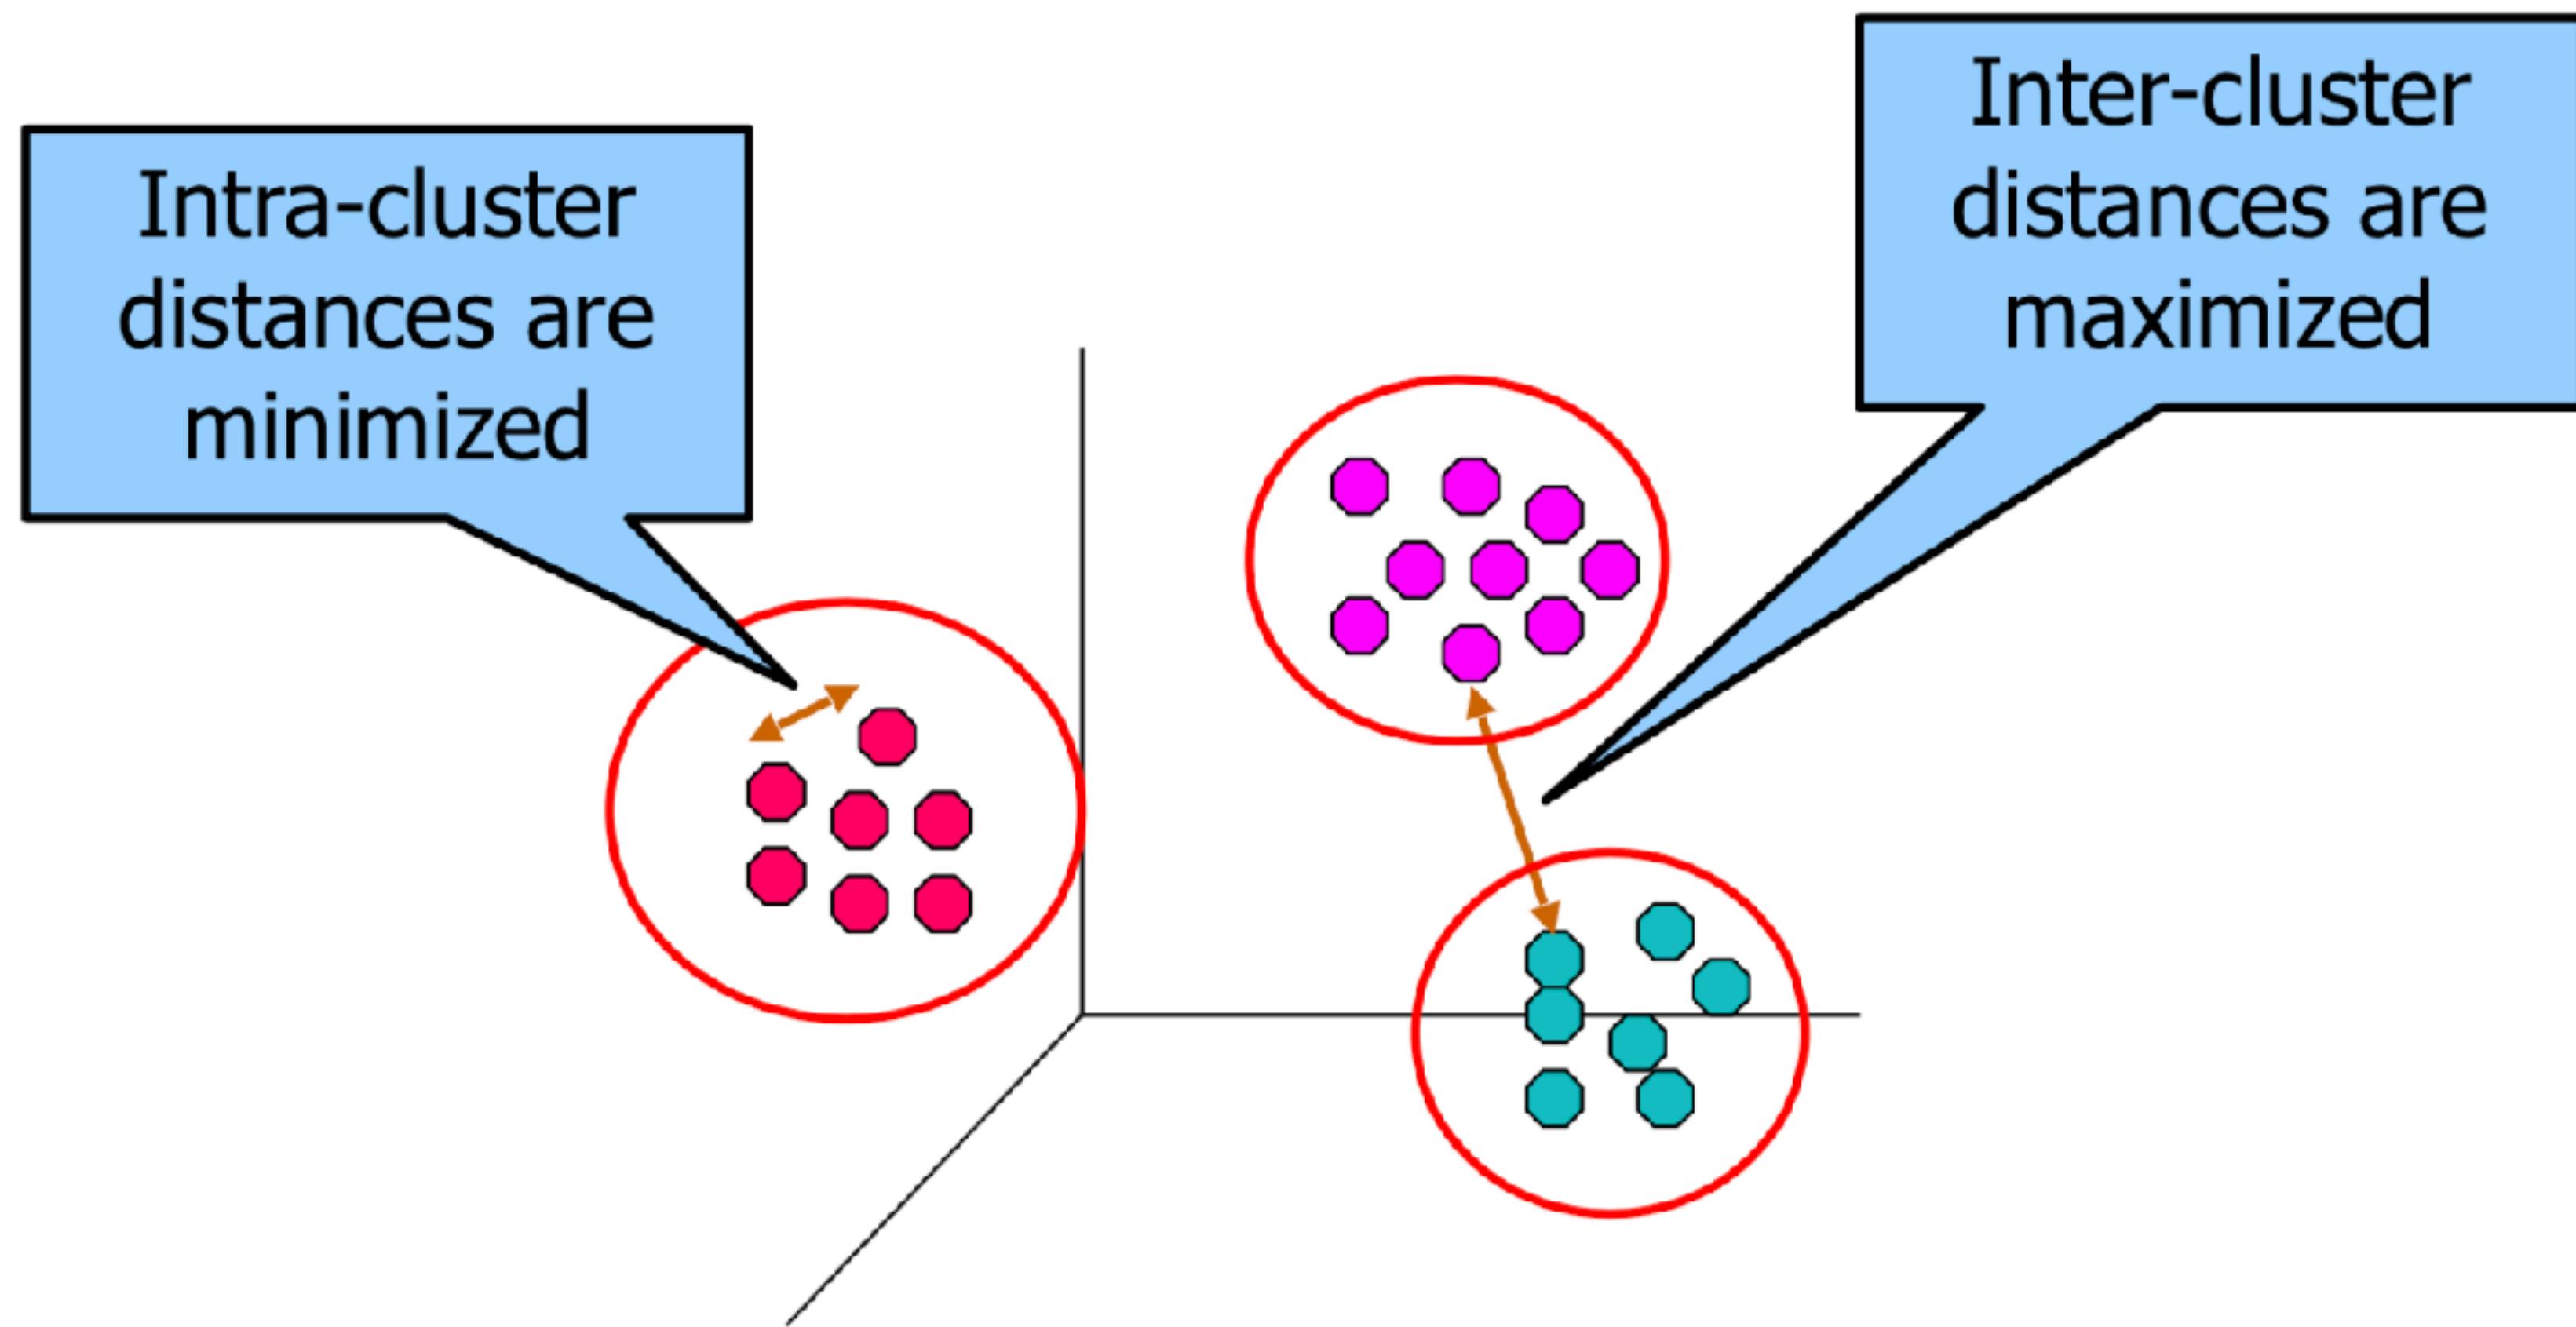
Clustering Example

Image segmentation

Goal: Break up the image into meaningful or perceptually similar regions



Clustering means grouping similar objects



Clustering = Distance + Method

Clustering vs. Classification

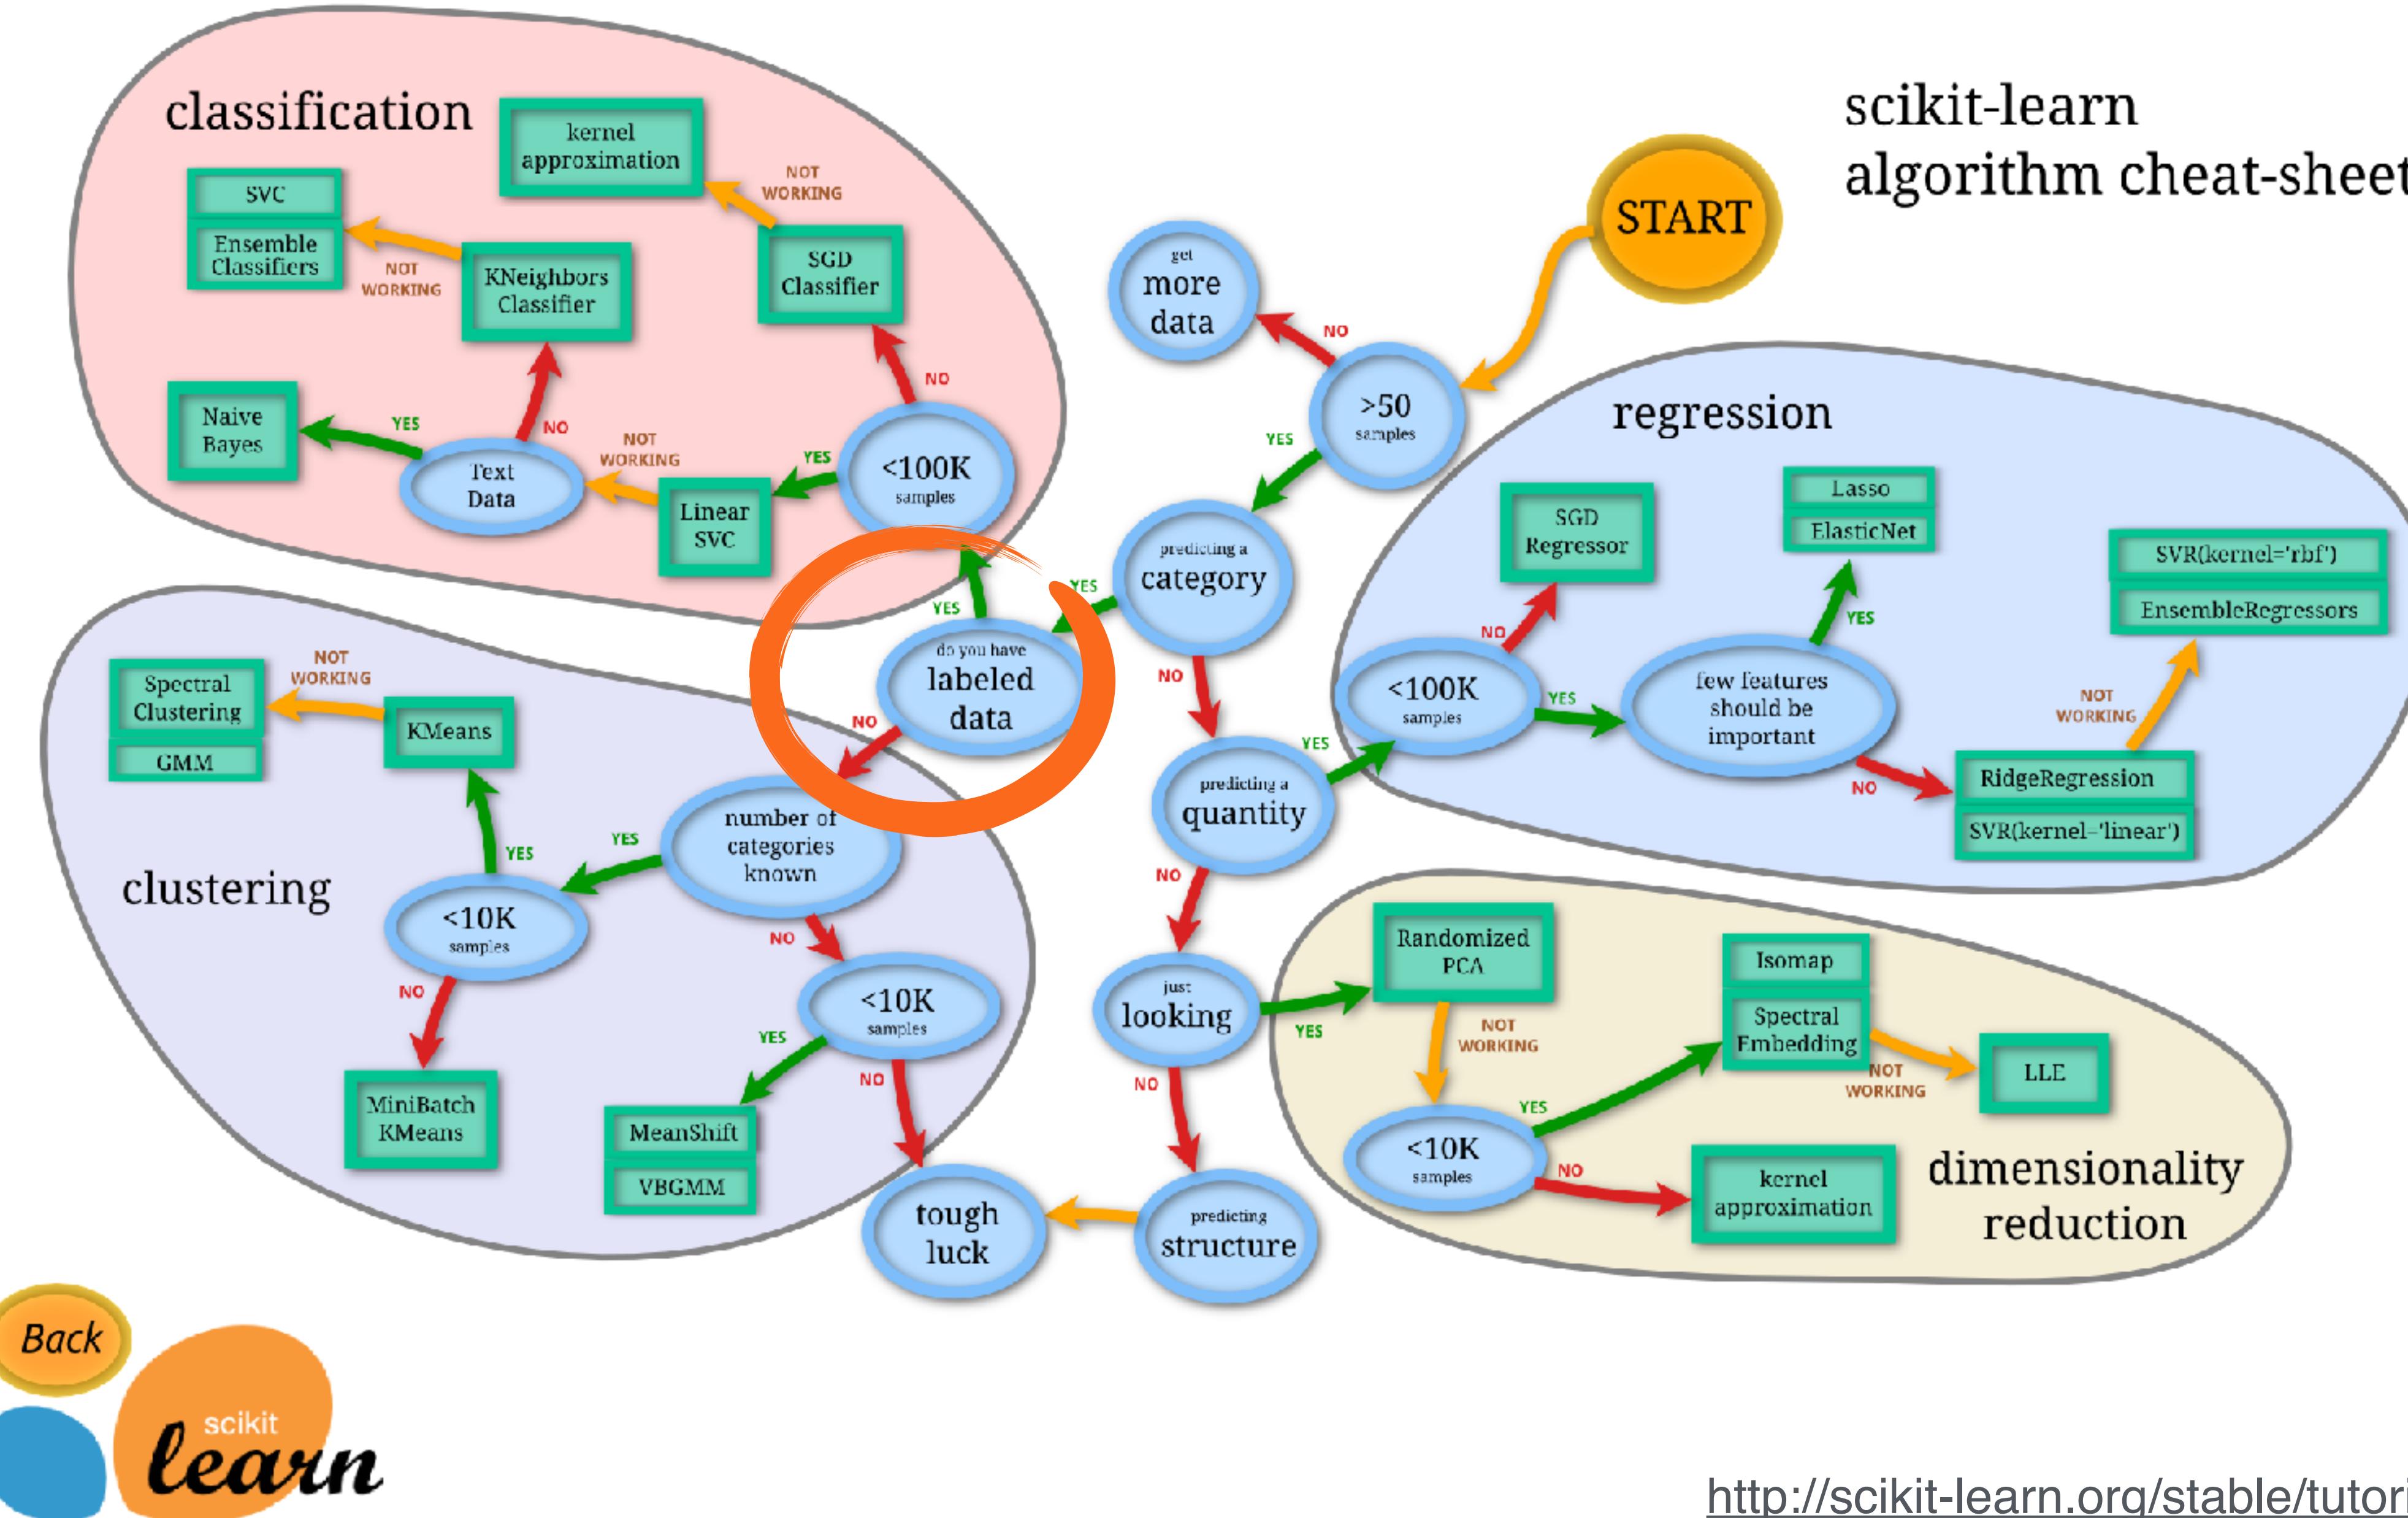
Clustering vs. Classification

unsupervised learning technique
that groups data points based on
their similarities

supervised learning technique
that assigns data points to
predefined categories

labeled data: Classification

no labeled data: Clustering

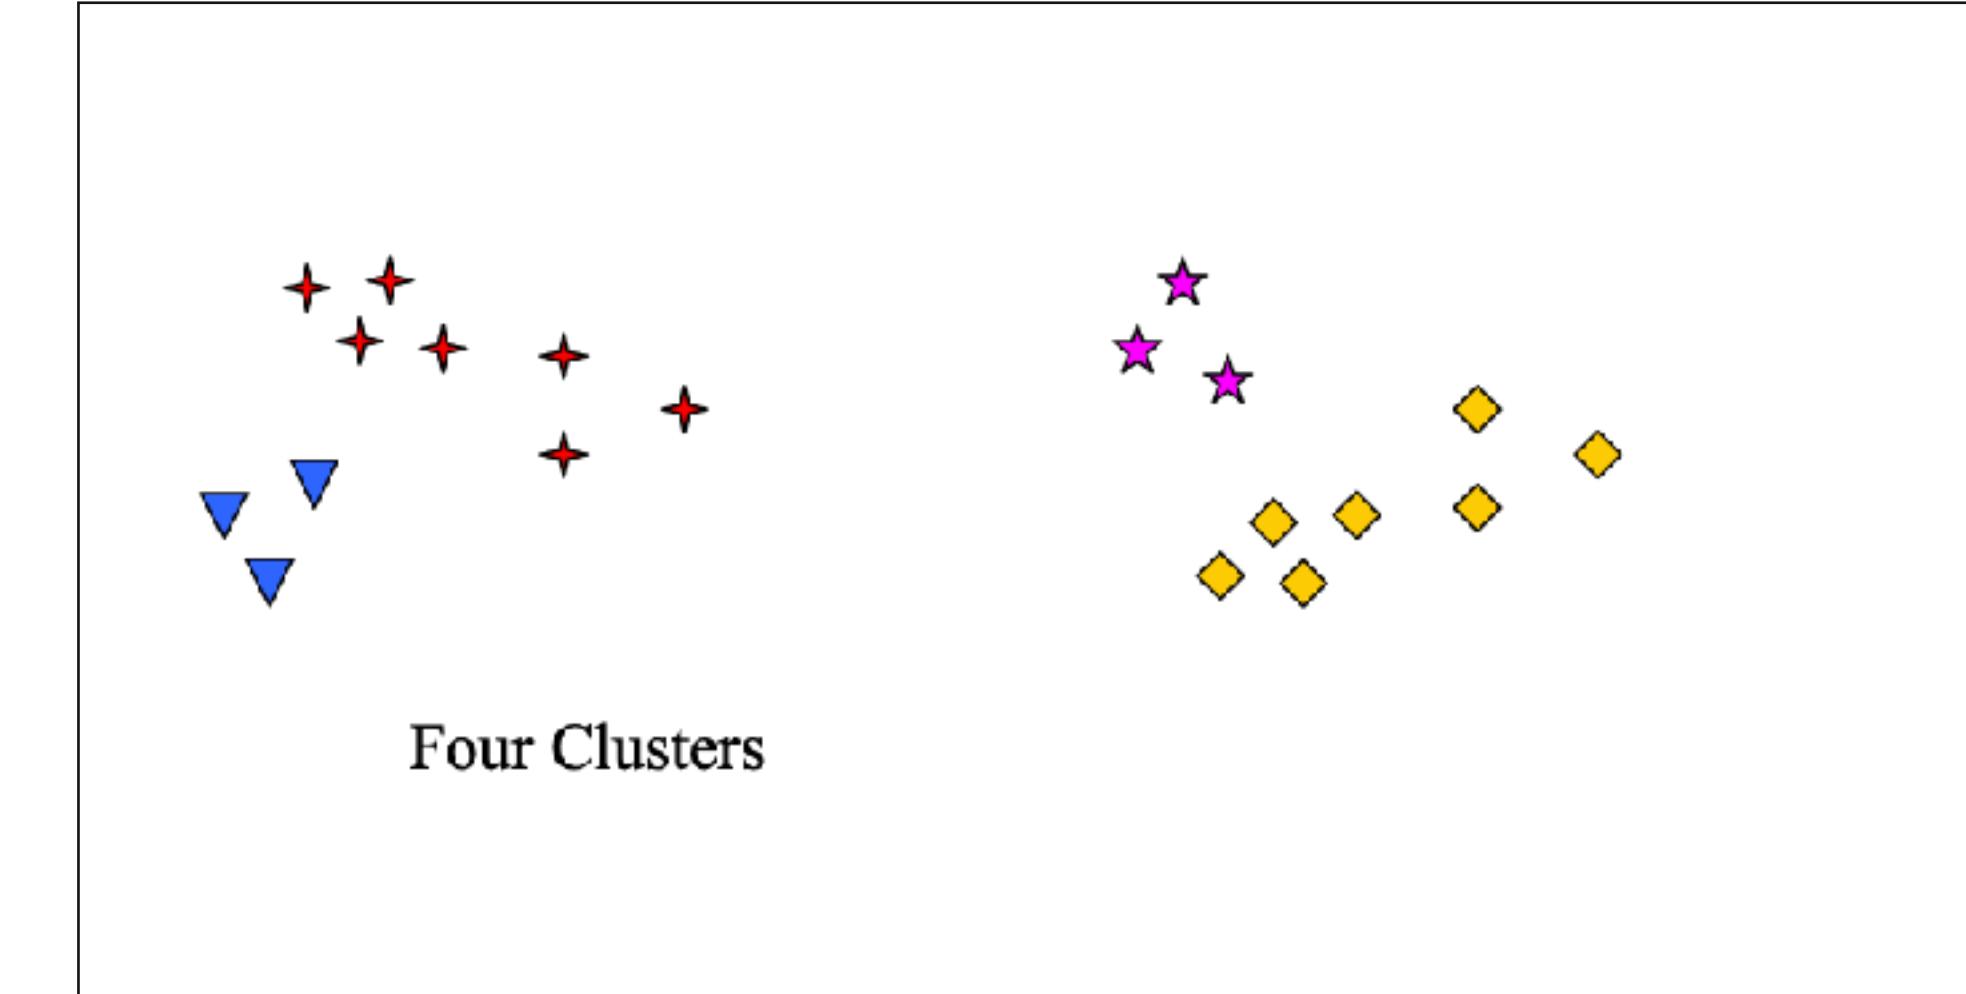
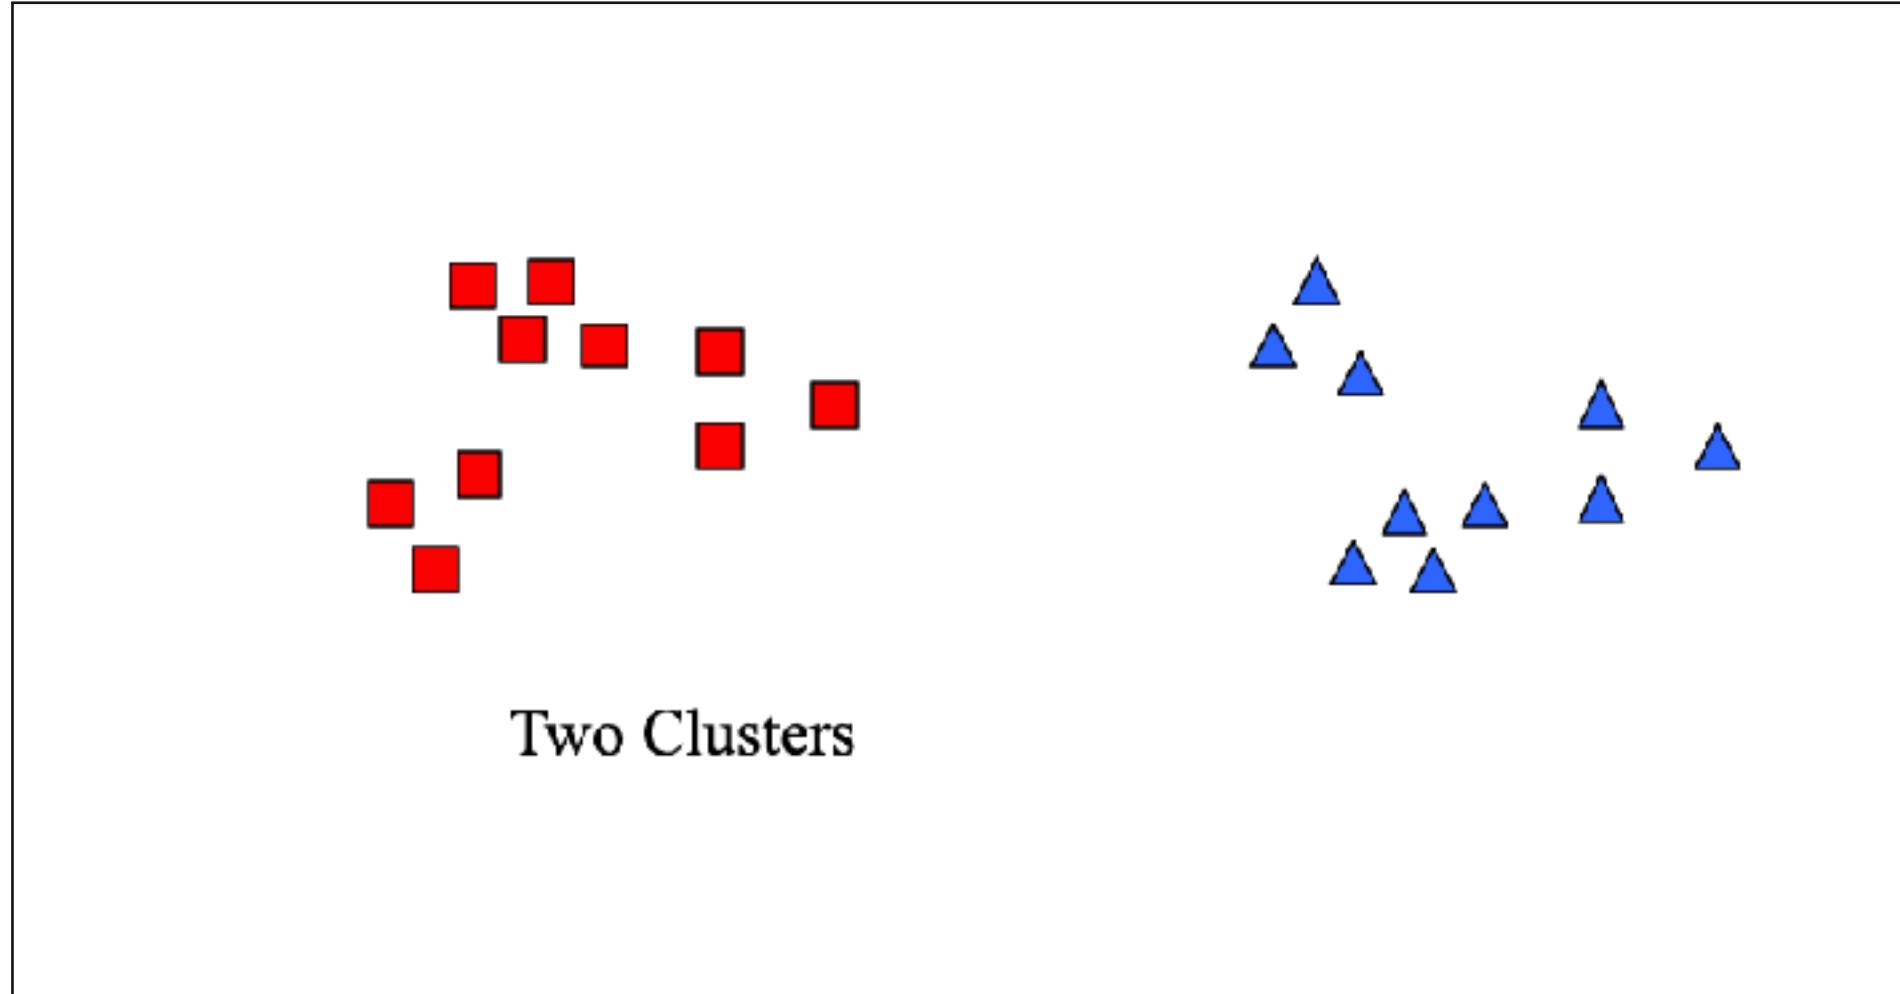
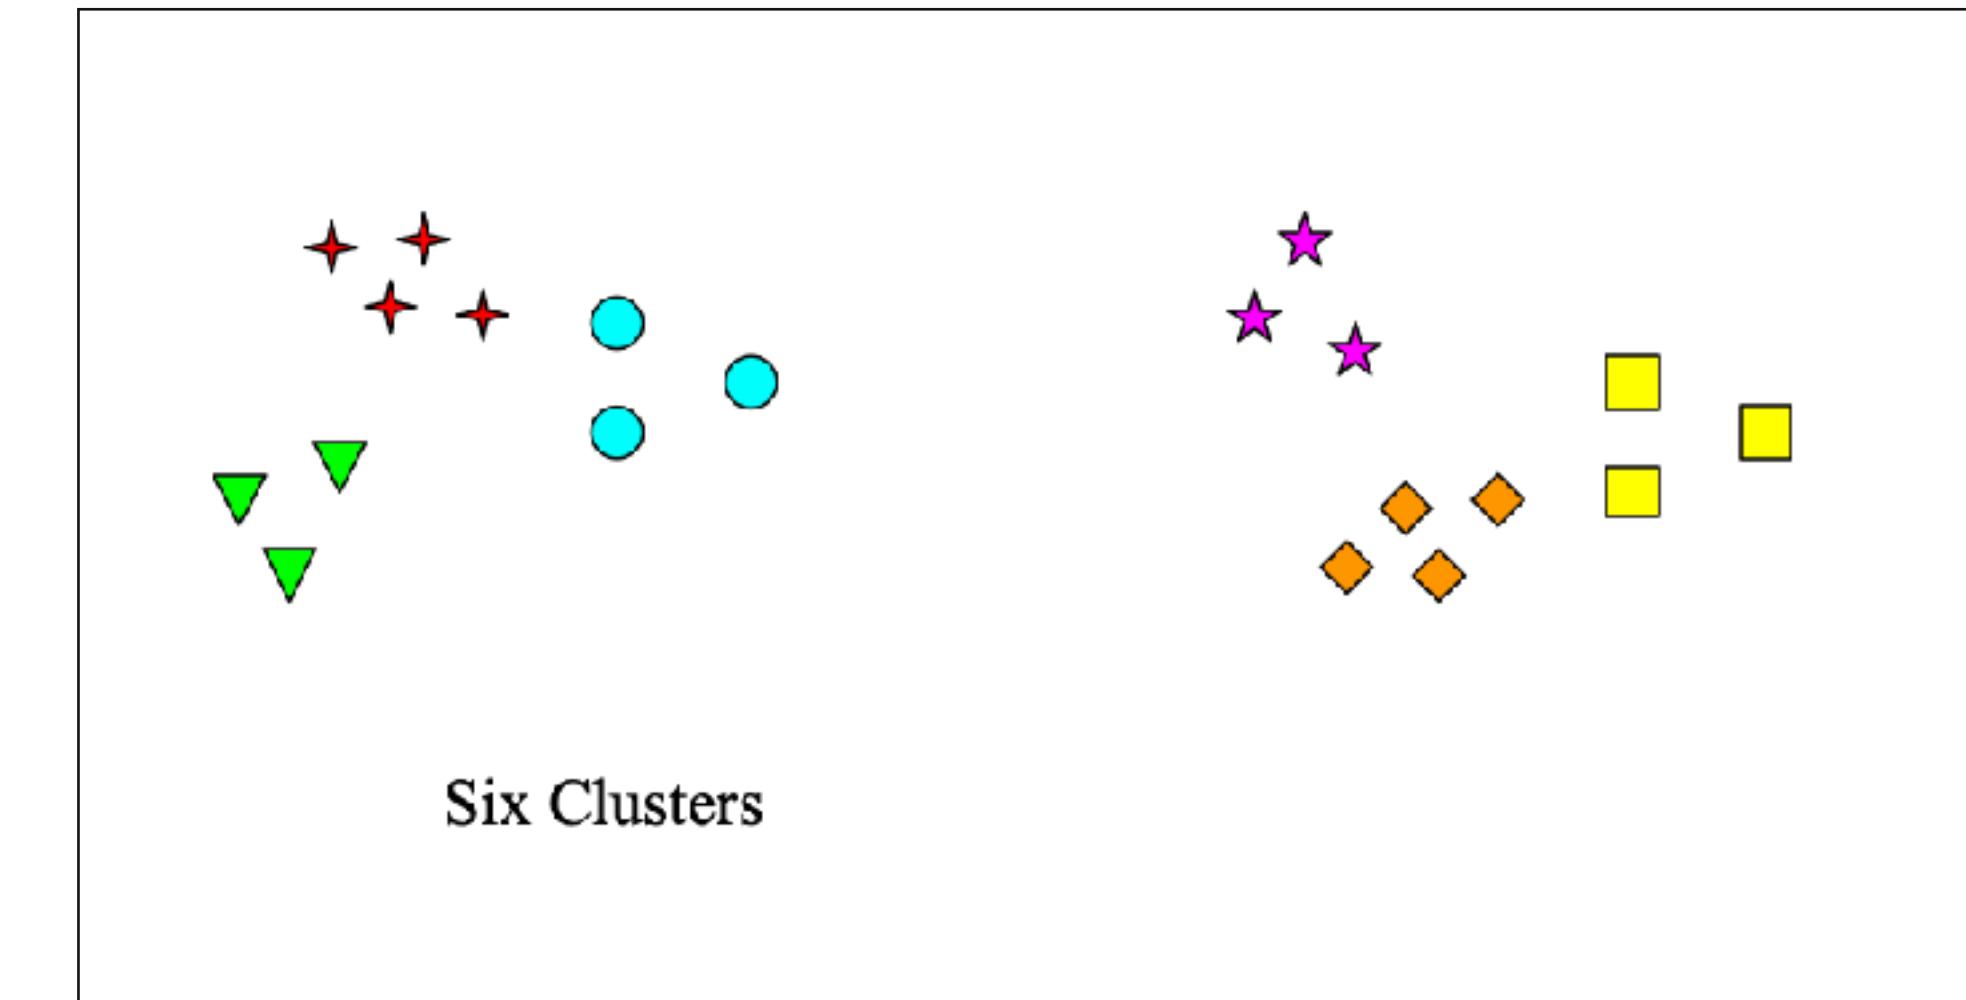
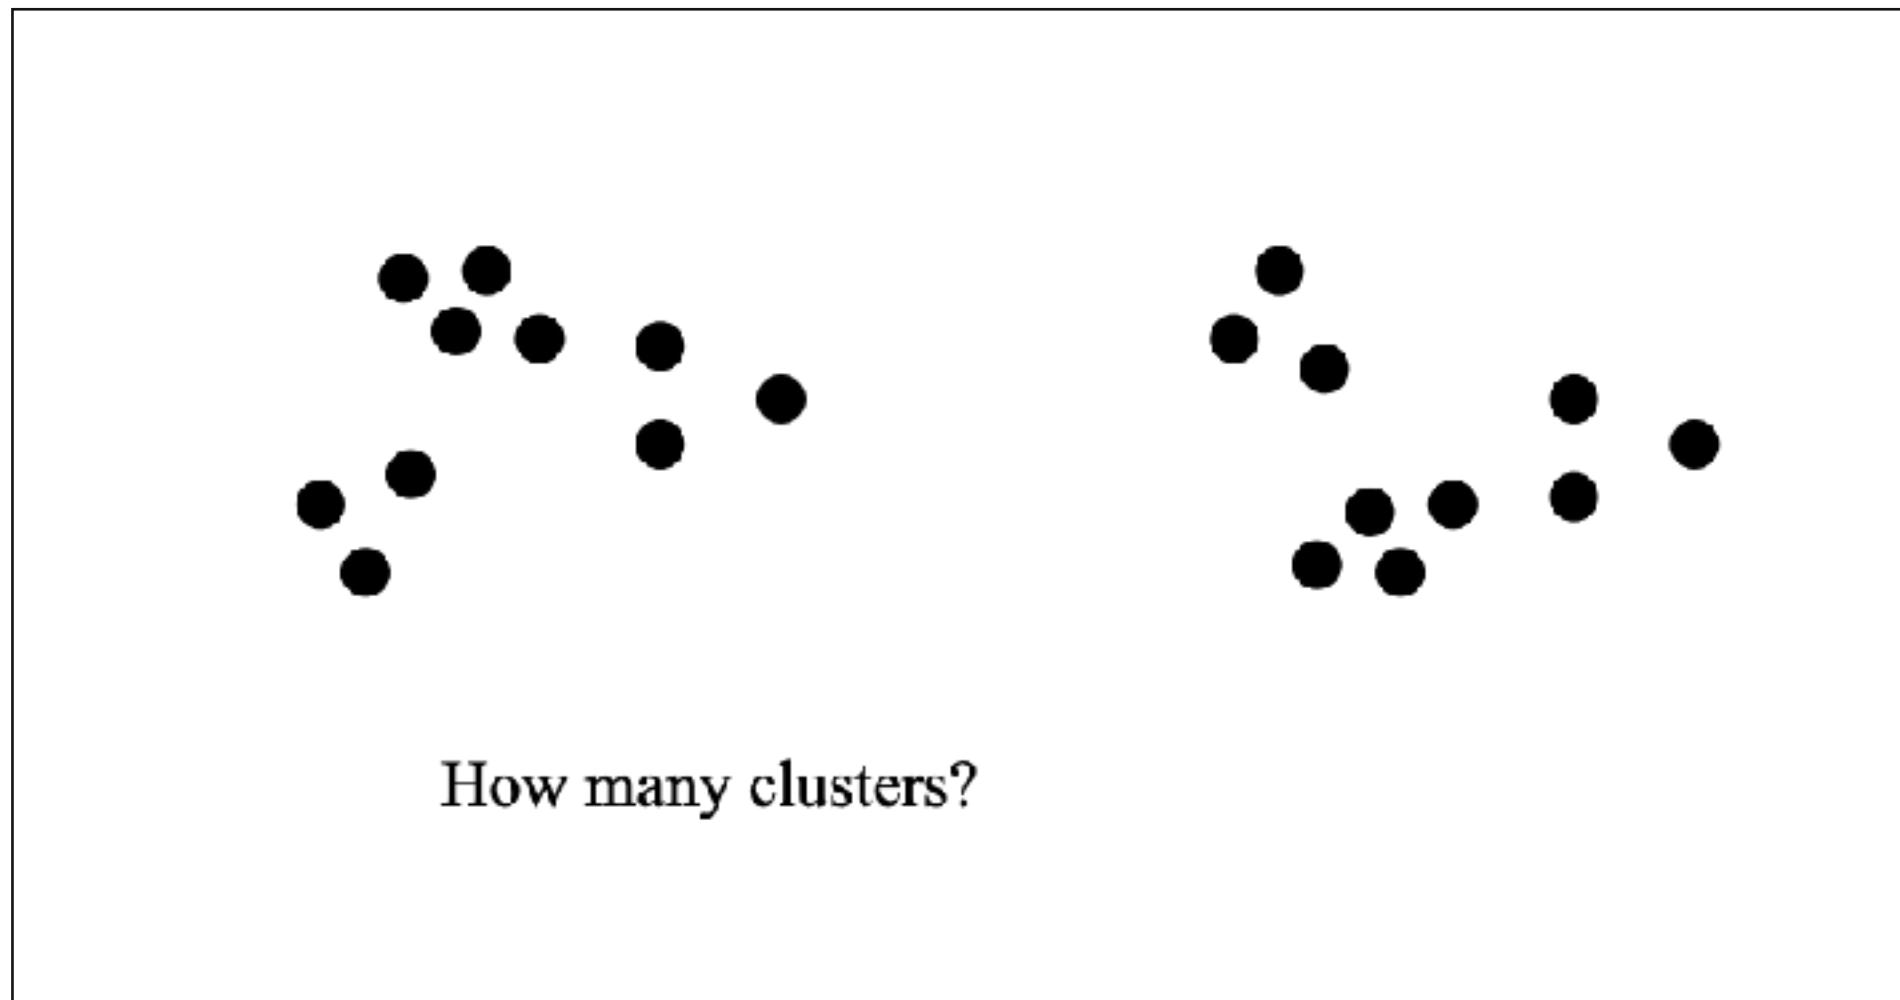


Clustering is Ambiguous

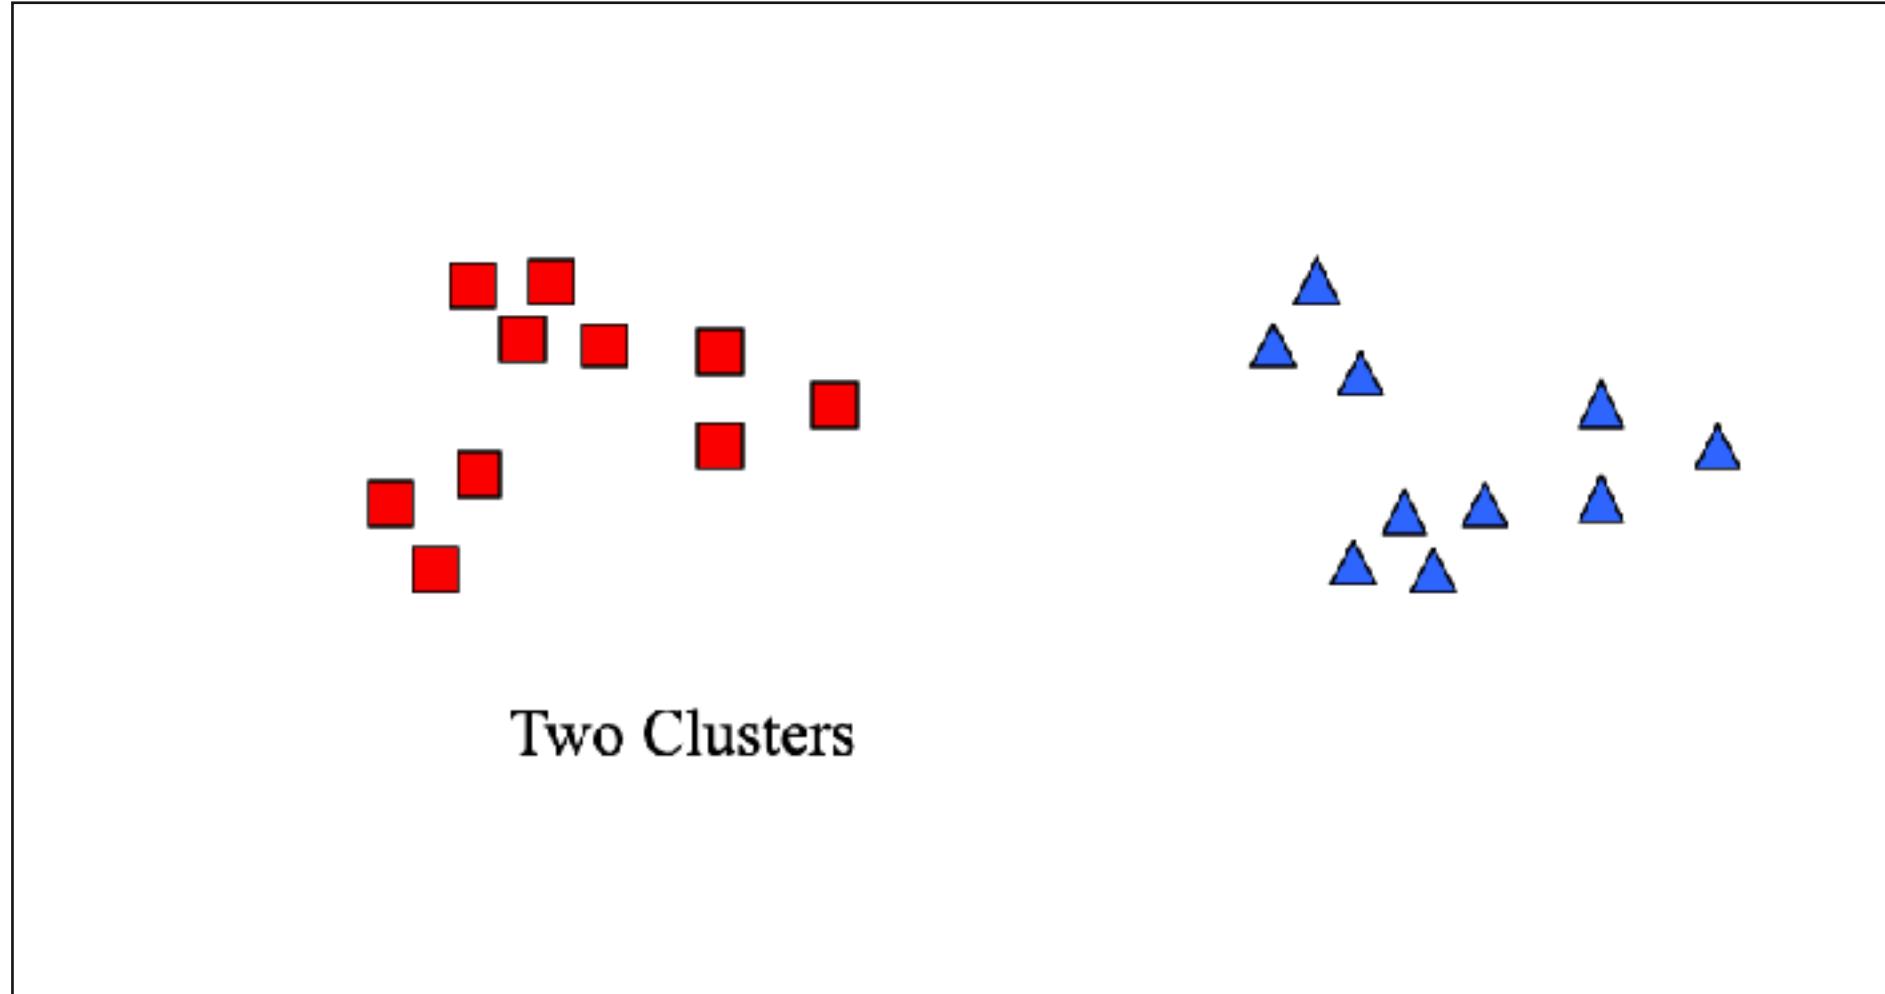
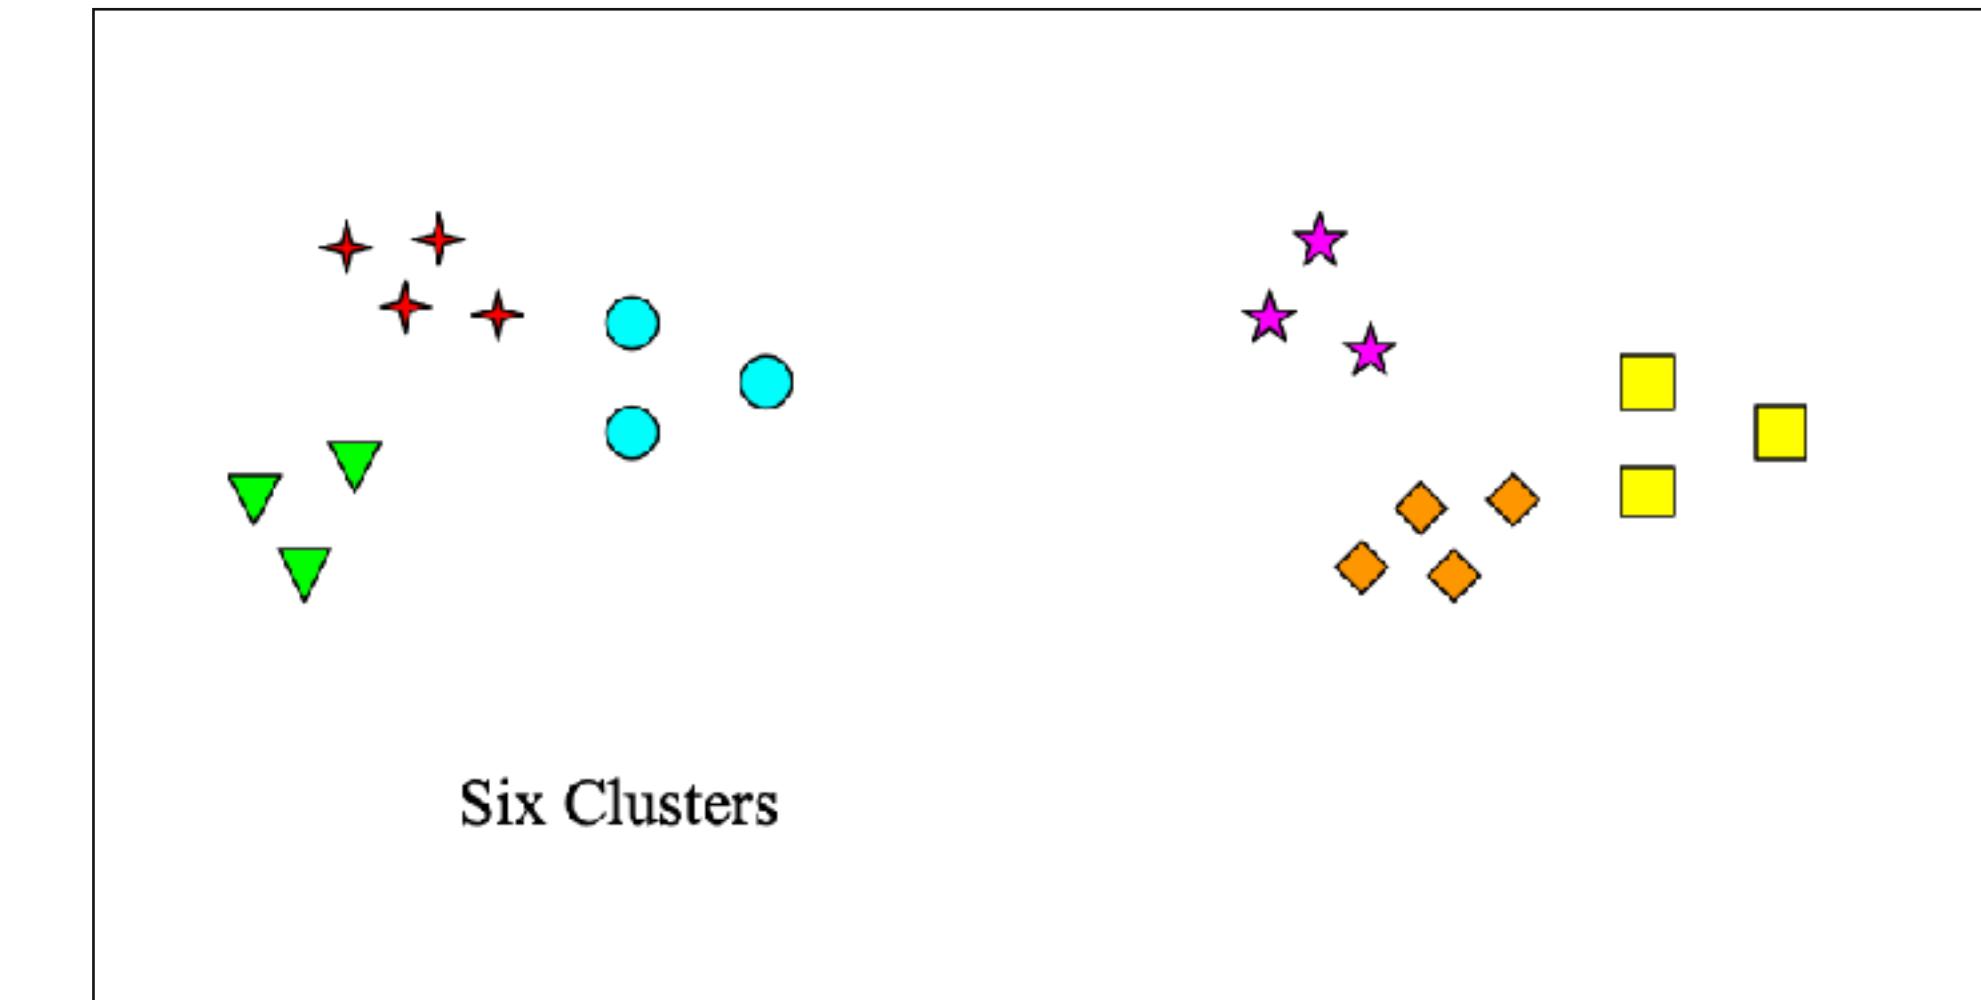
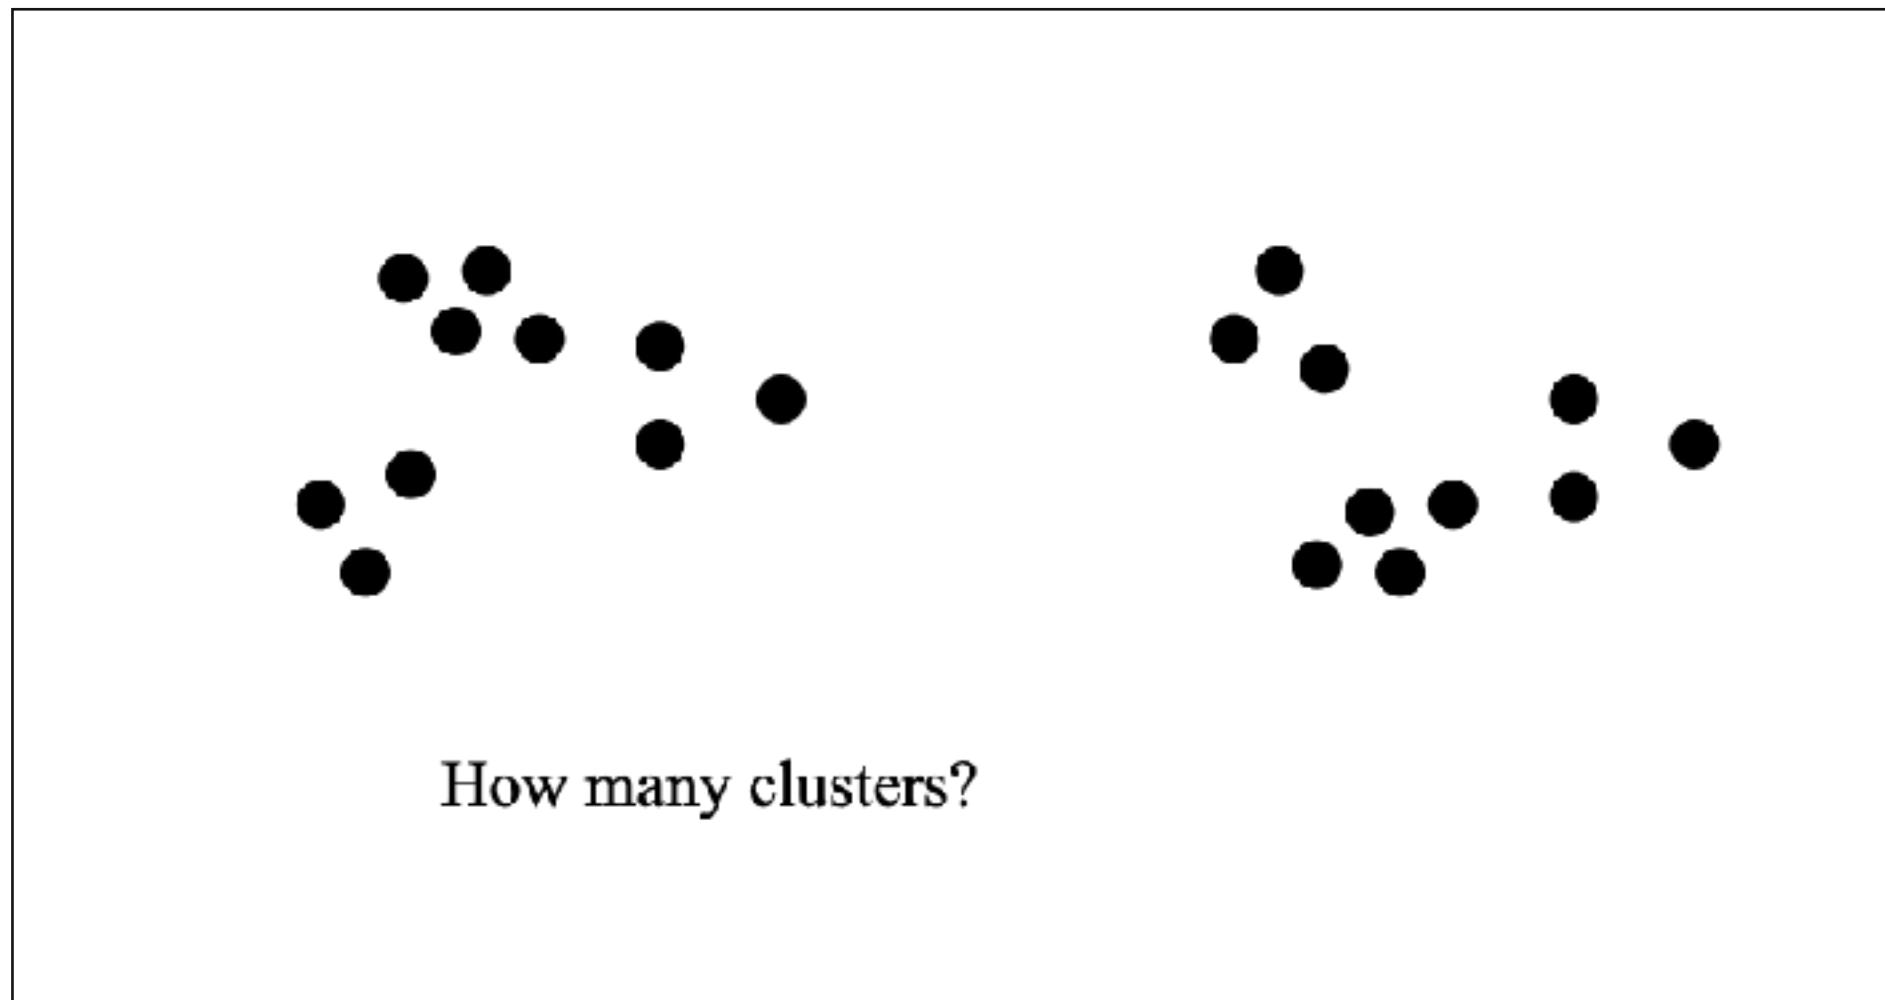


How many clusters?

Clustering is Ambiguous



Clustering is Ambiguous



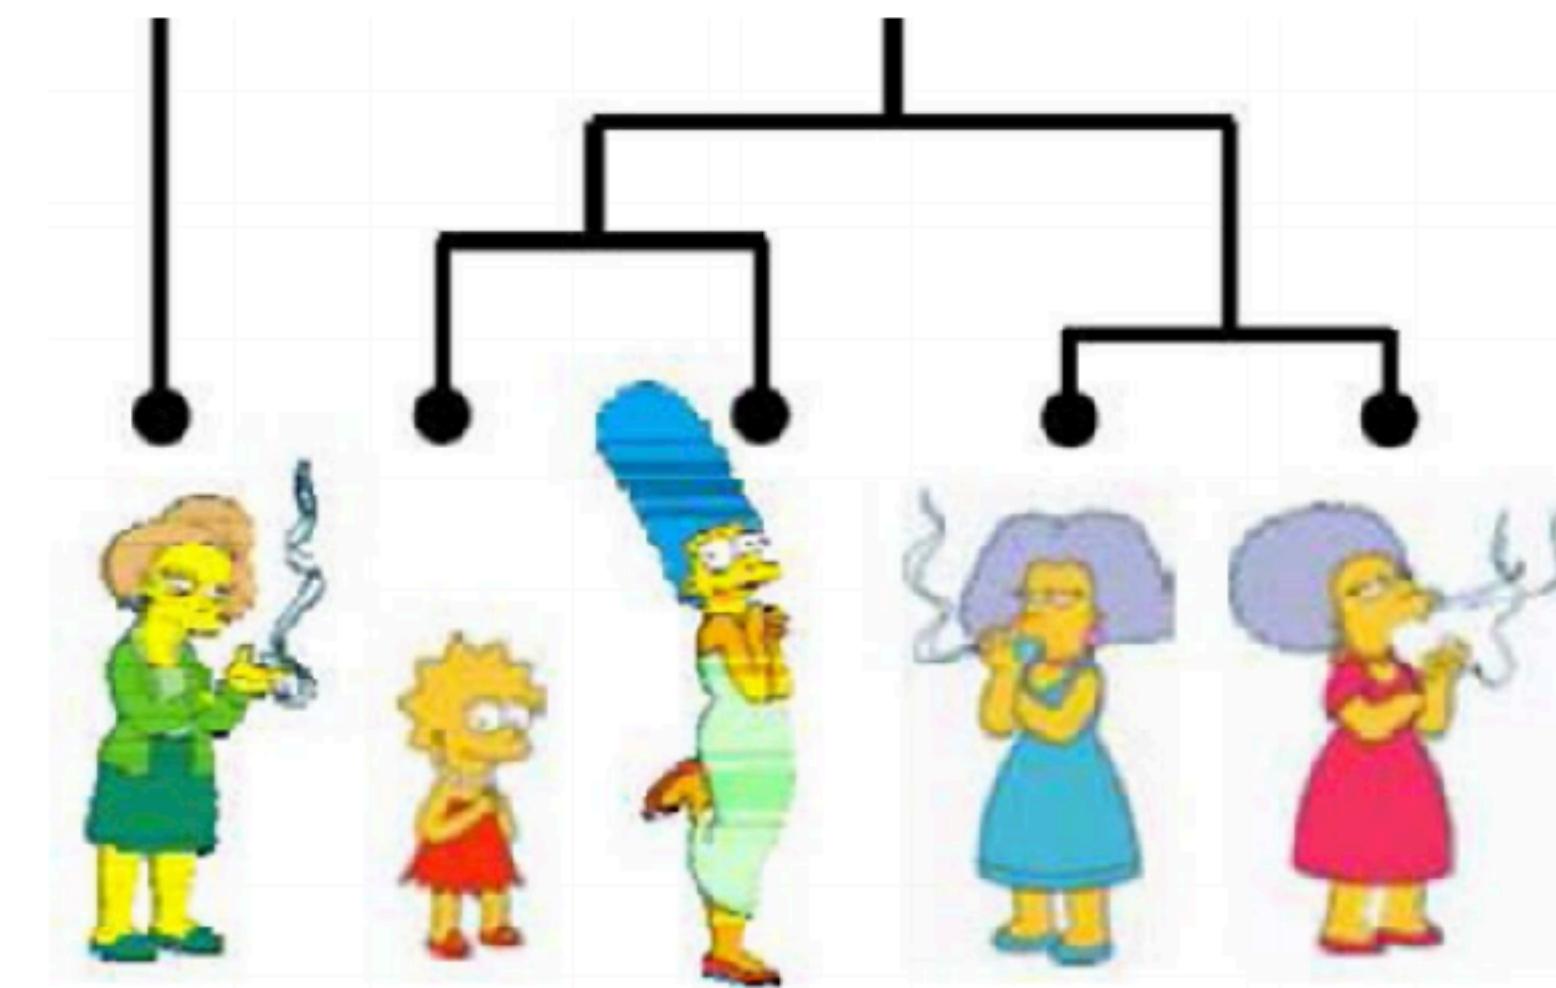
There is no universally applicable clustering technique

Clustering Algorithms

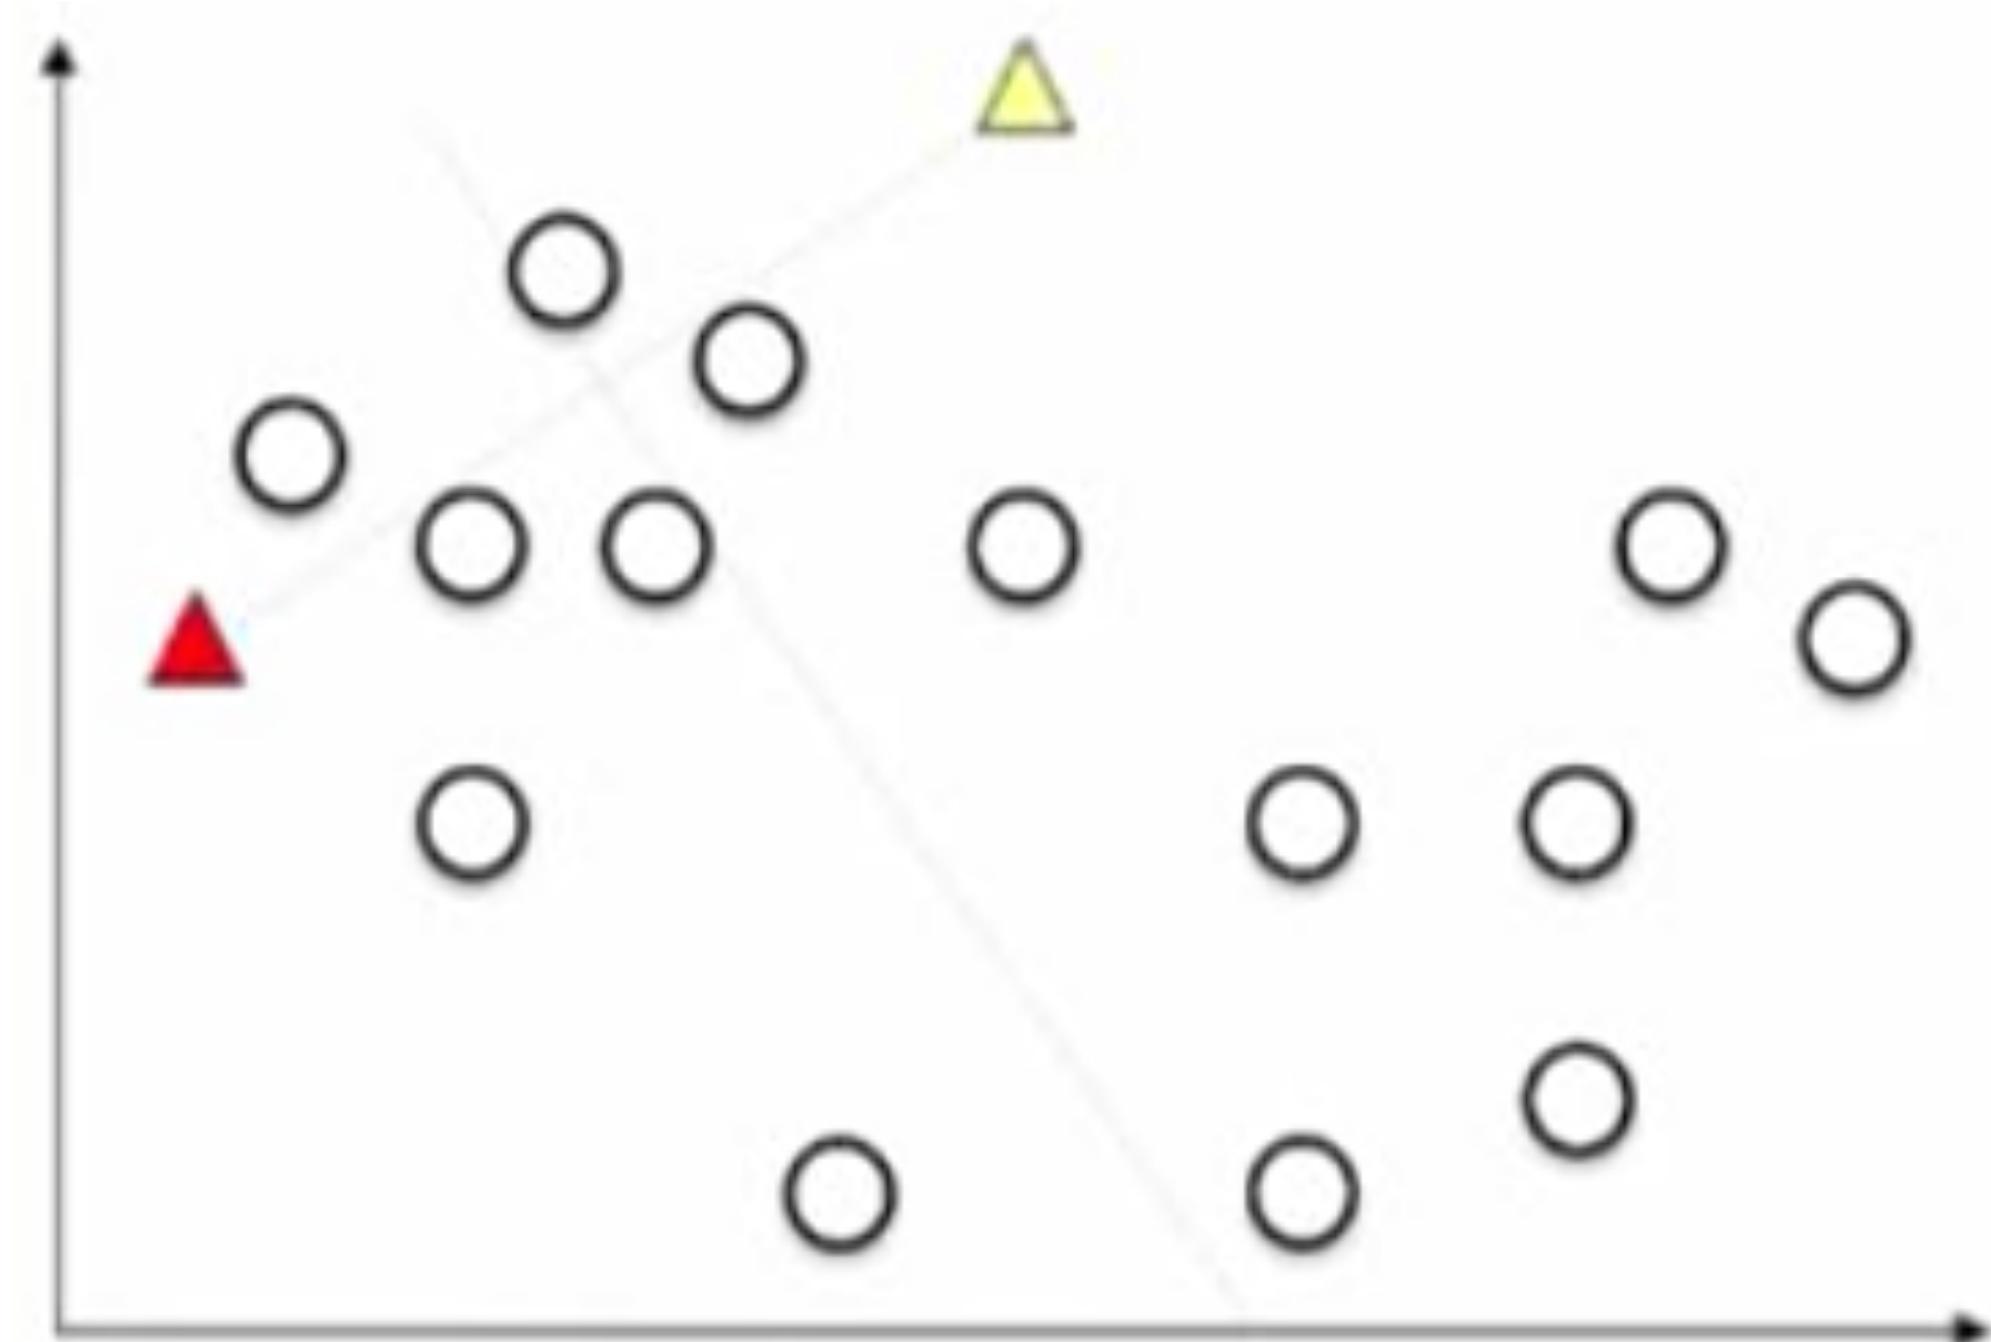
- Partition algorithms (Flat)
 - K-means
 - Mixture of Gaussian
 - Spectral Clustering



- Hierarchical algorithms
 - Bottom up – agglomerative
 - Top down – divisive



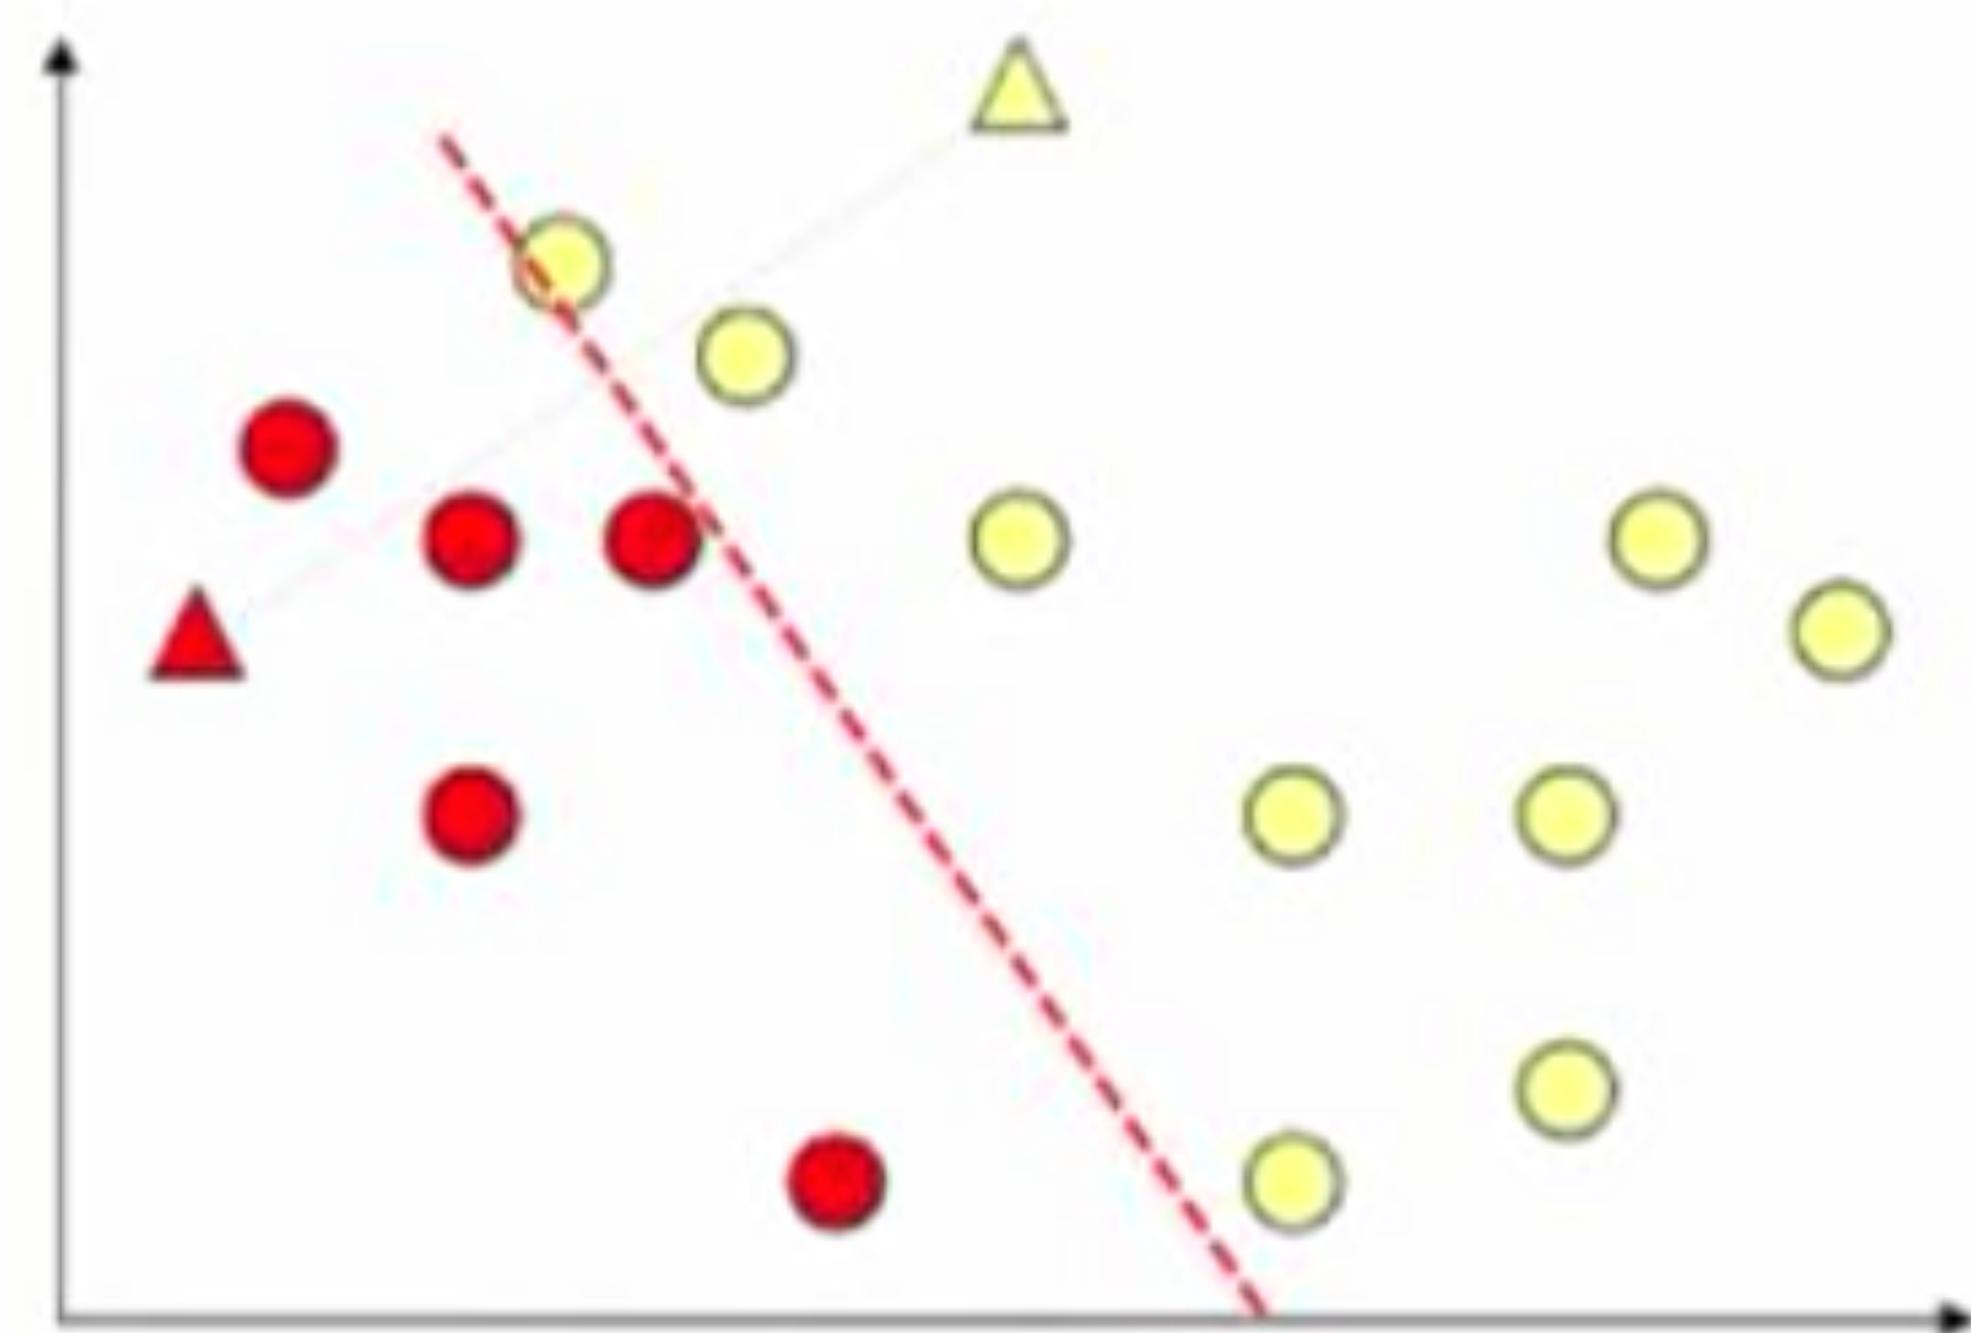
K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

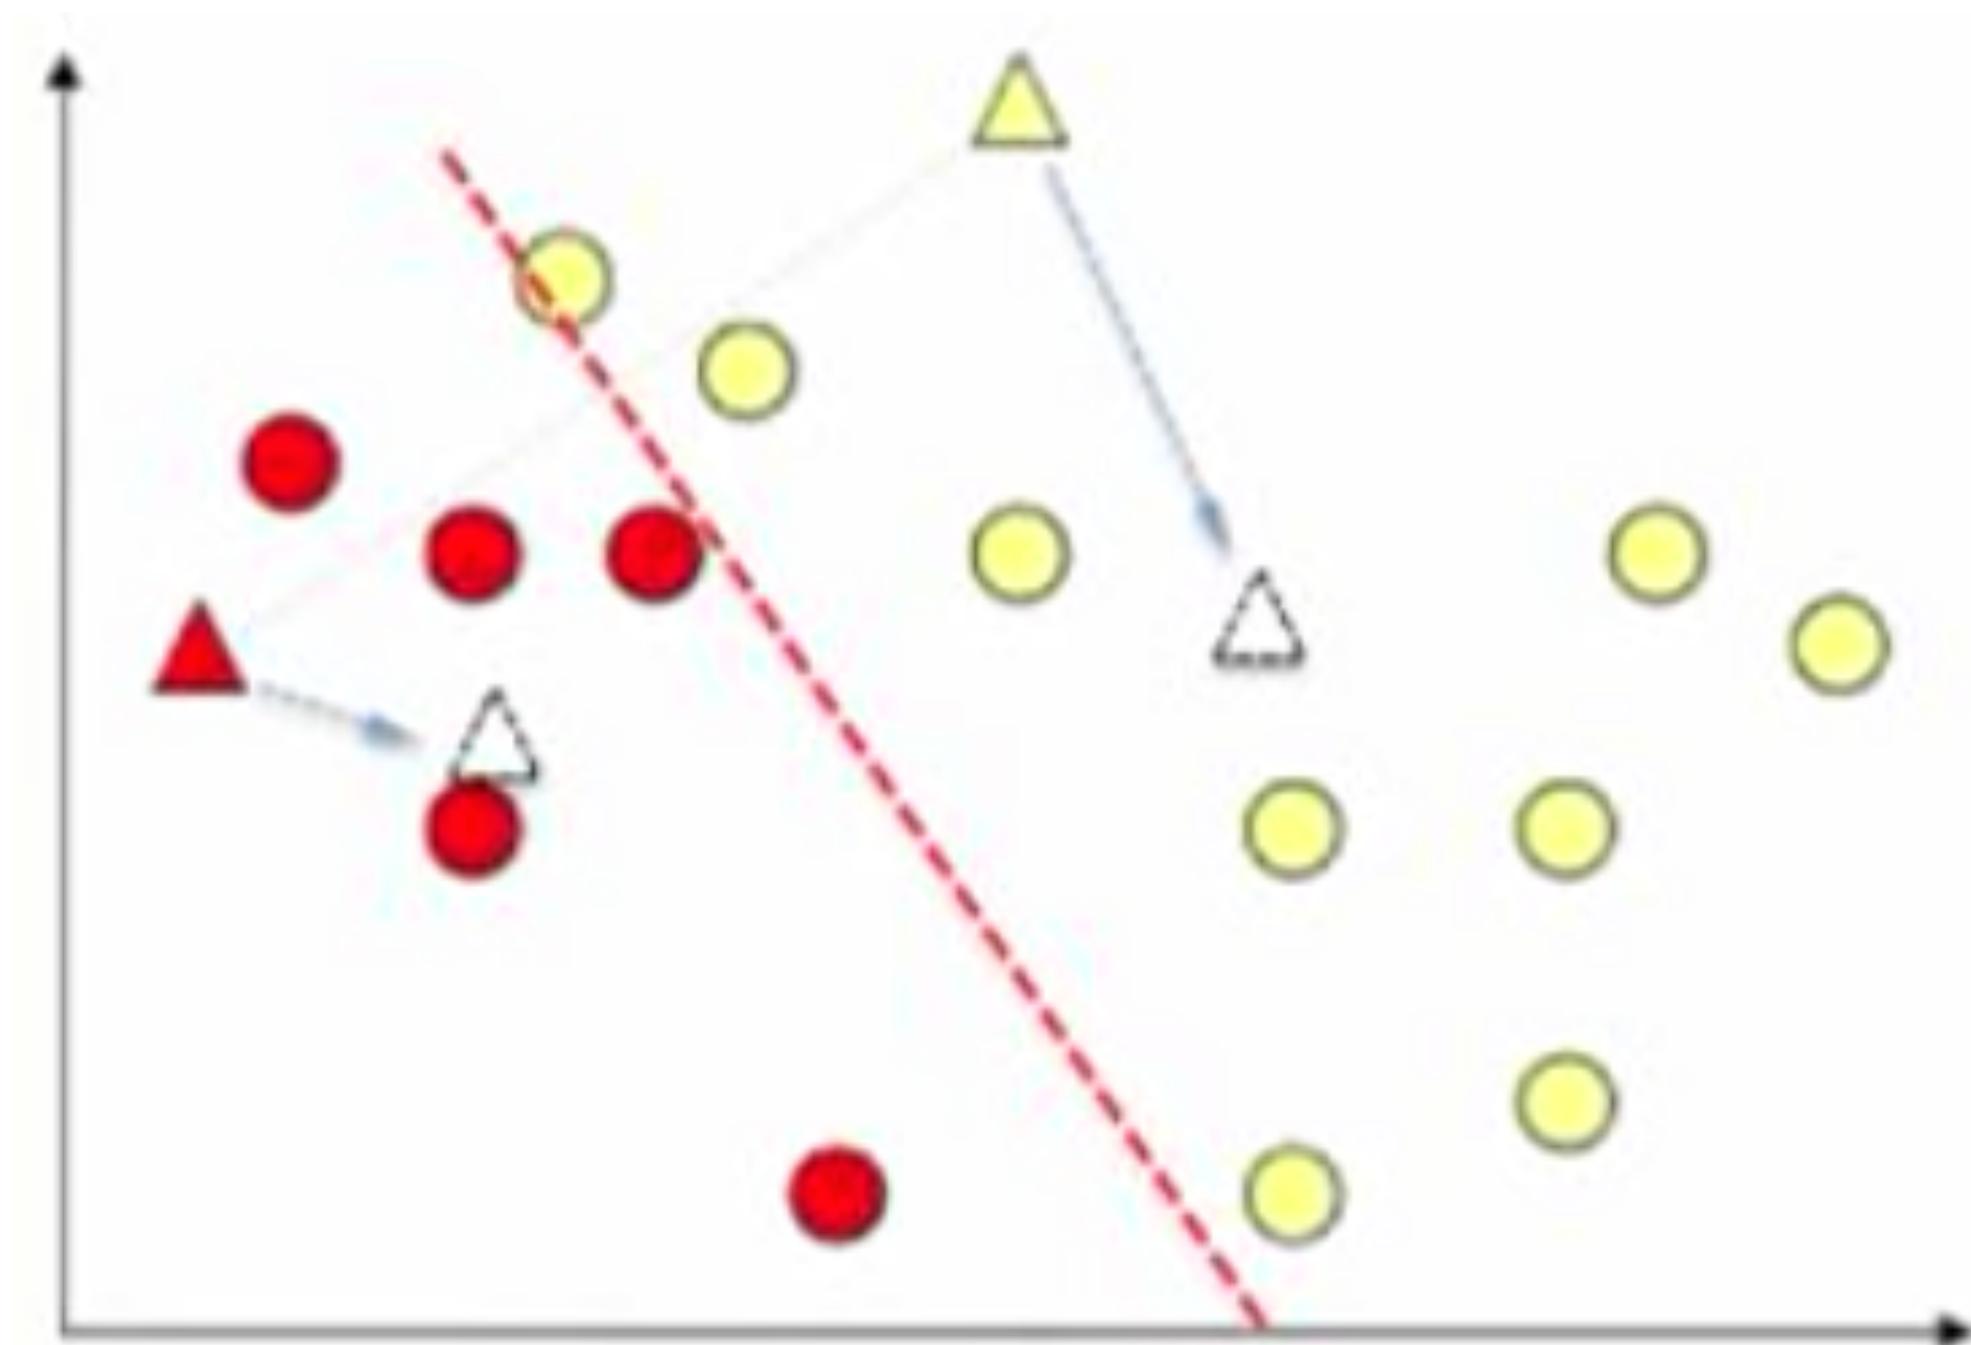
K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

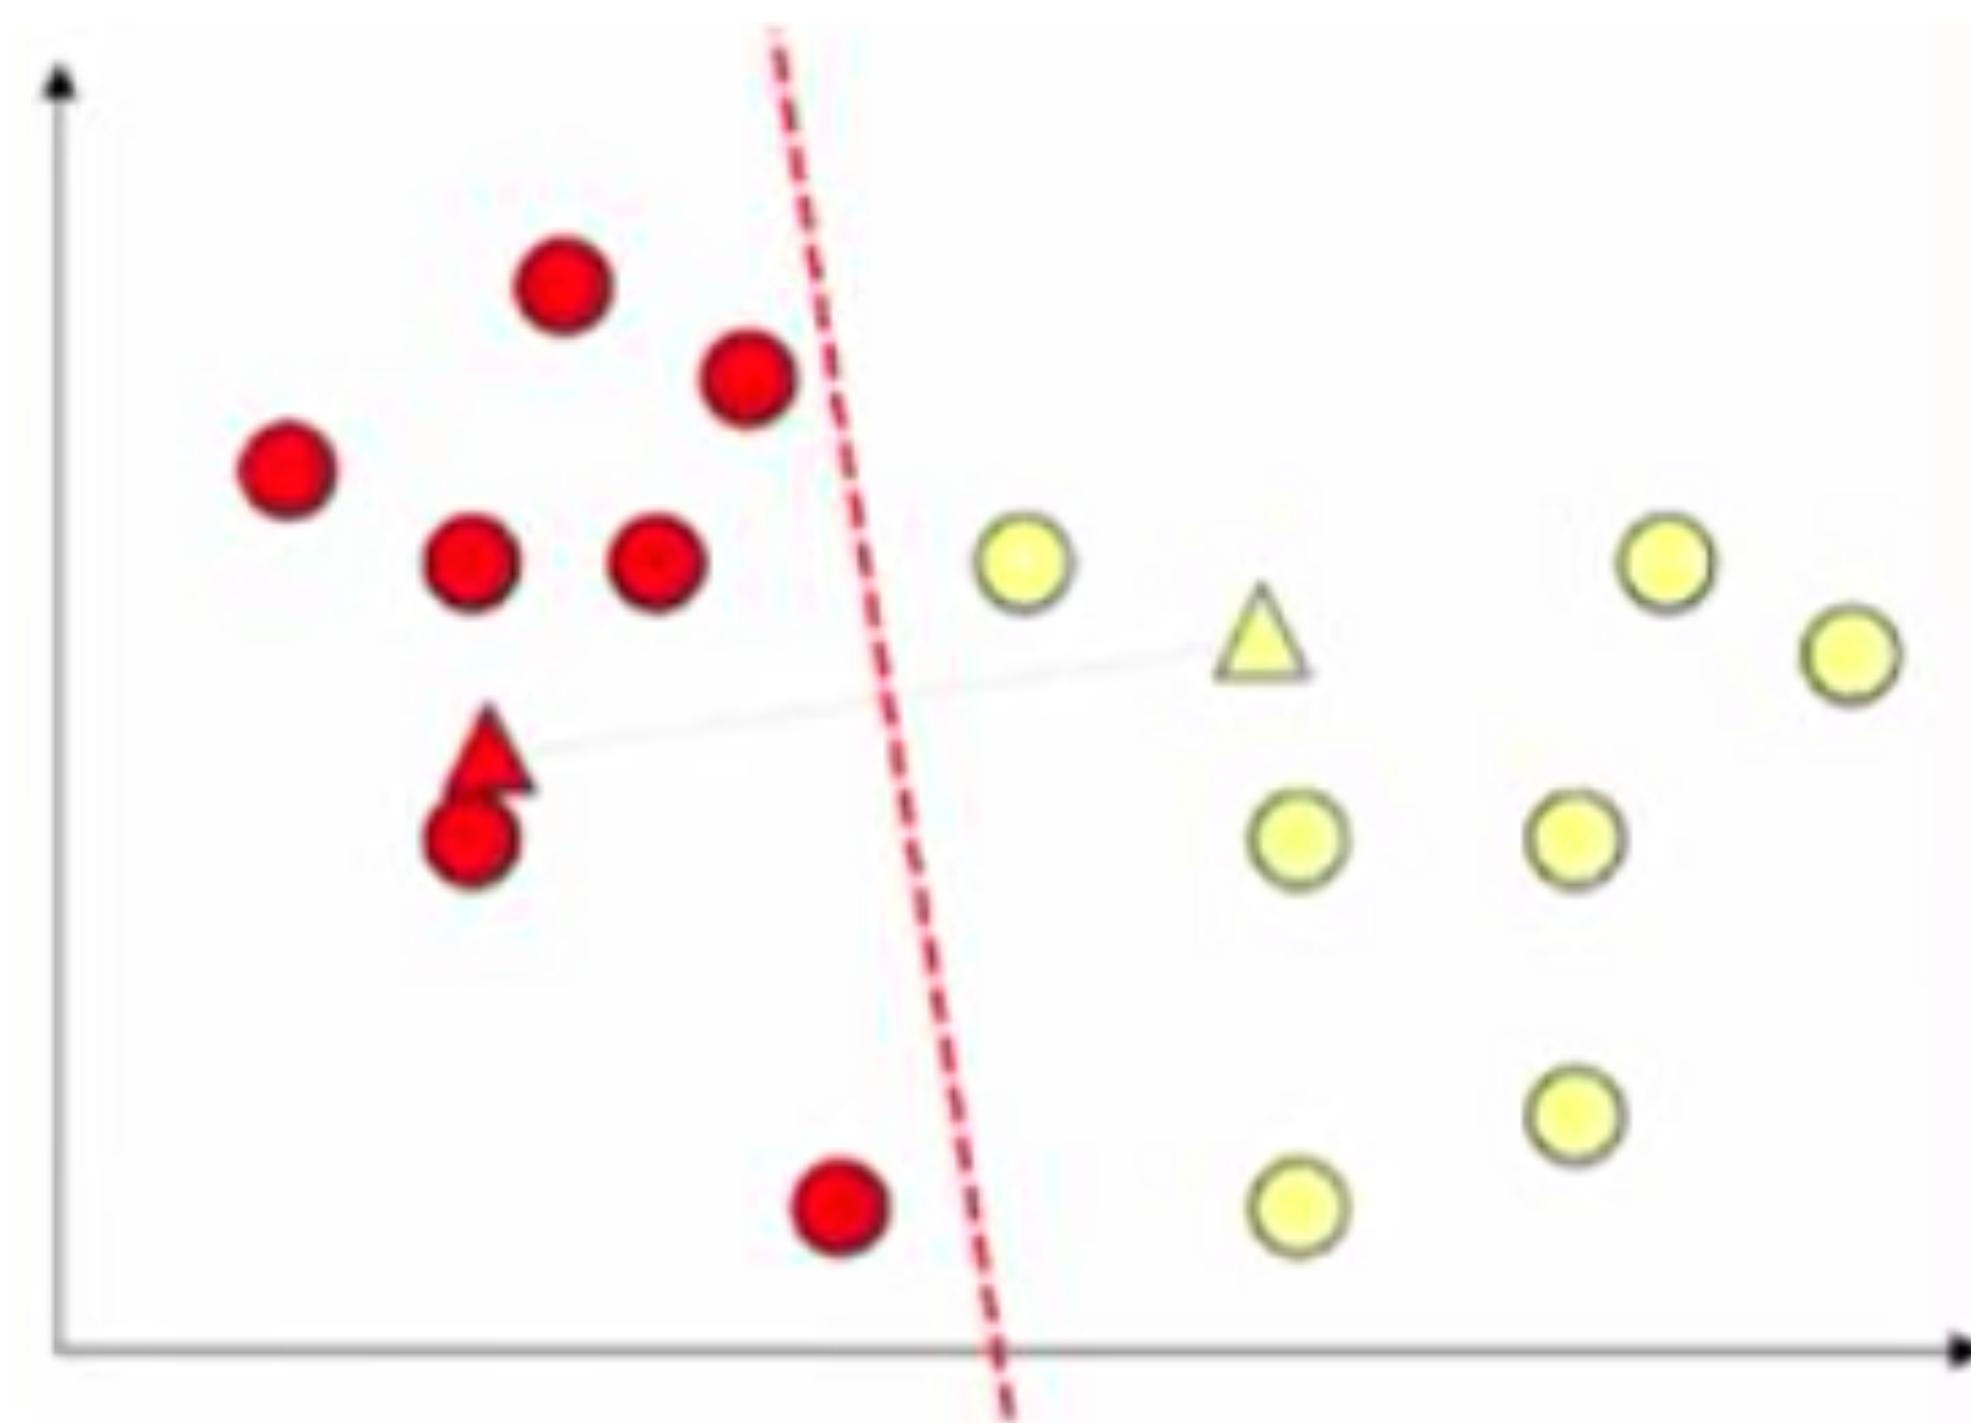
K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

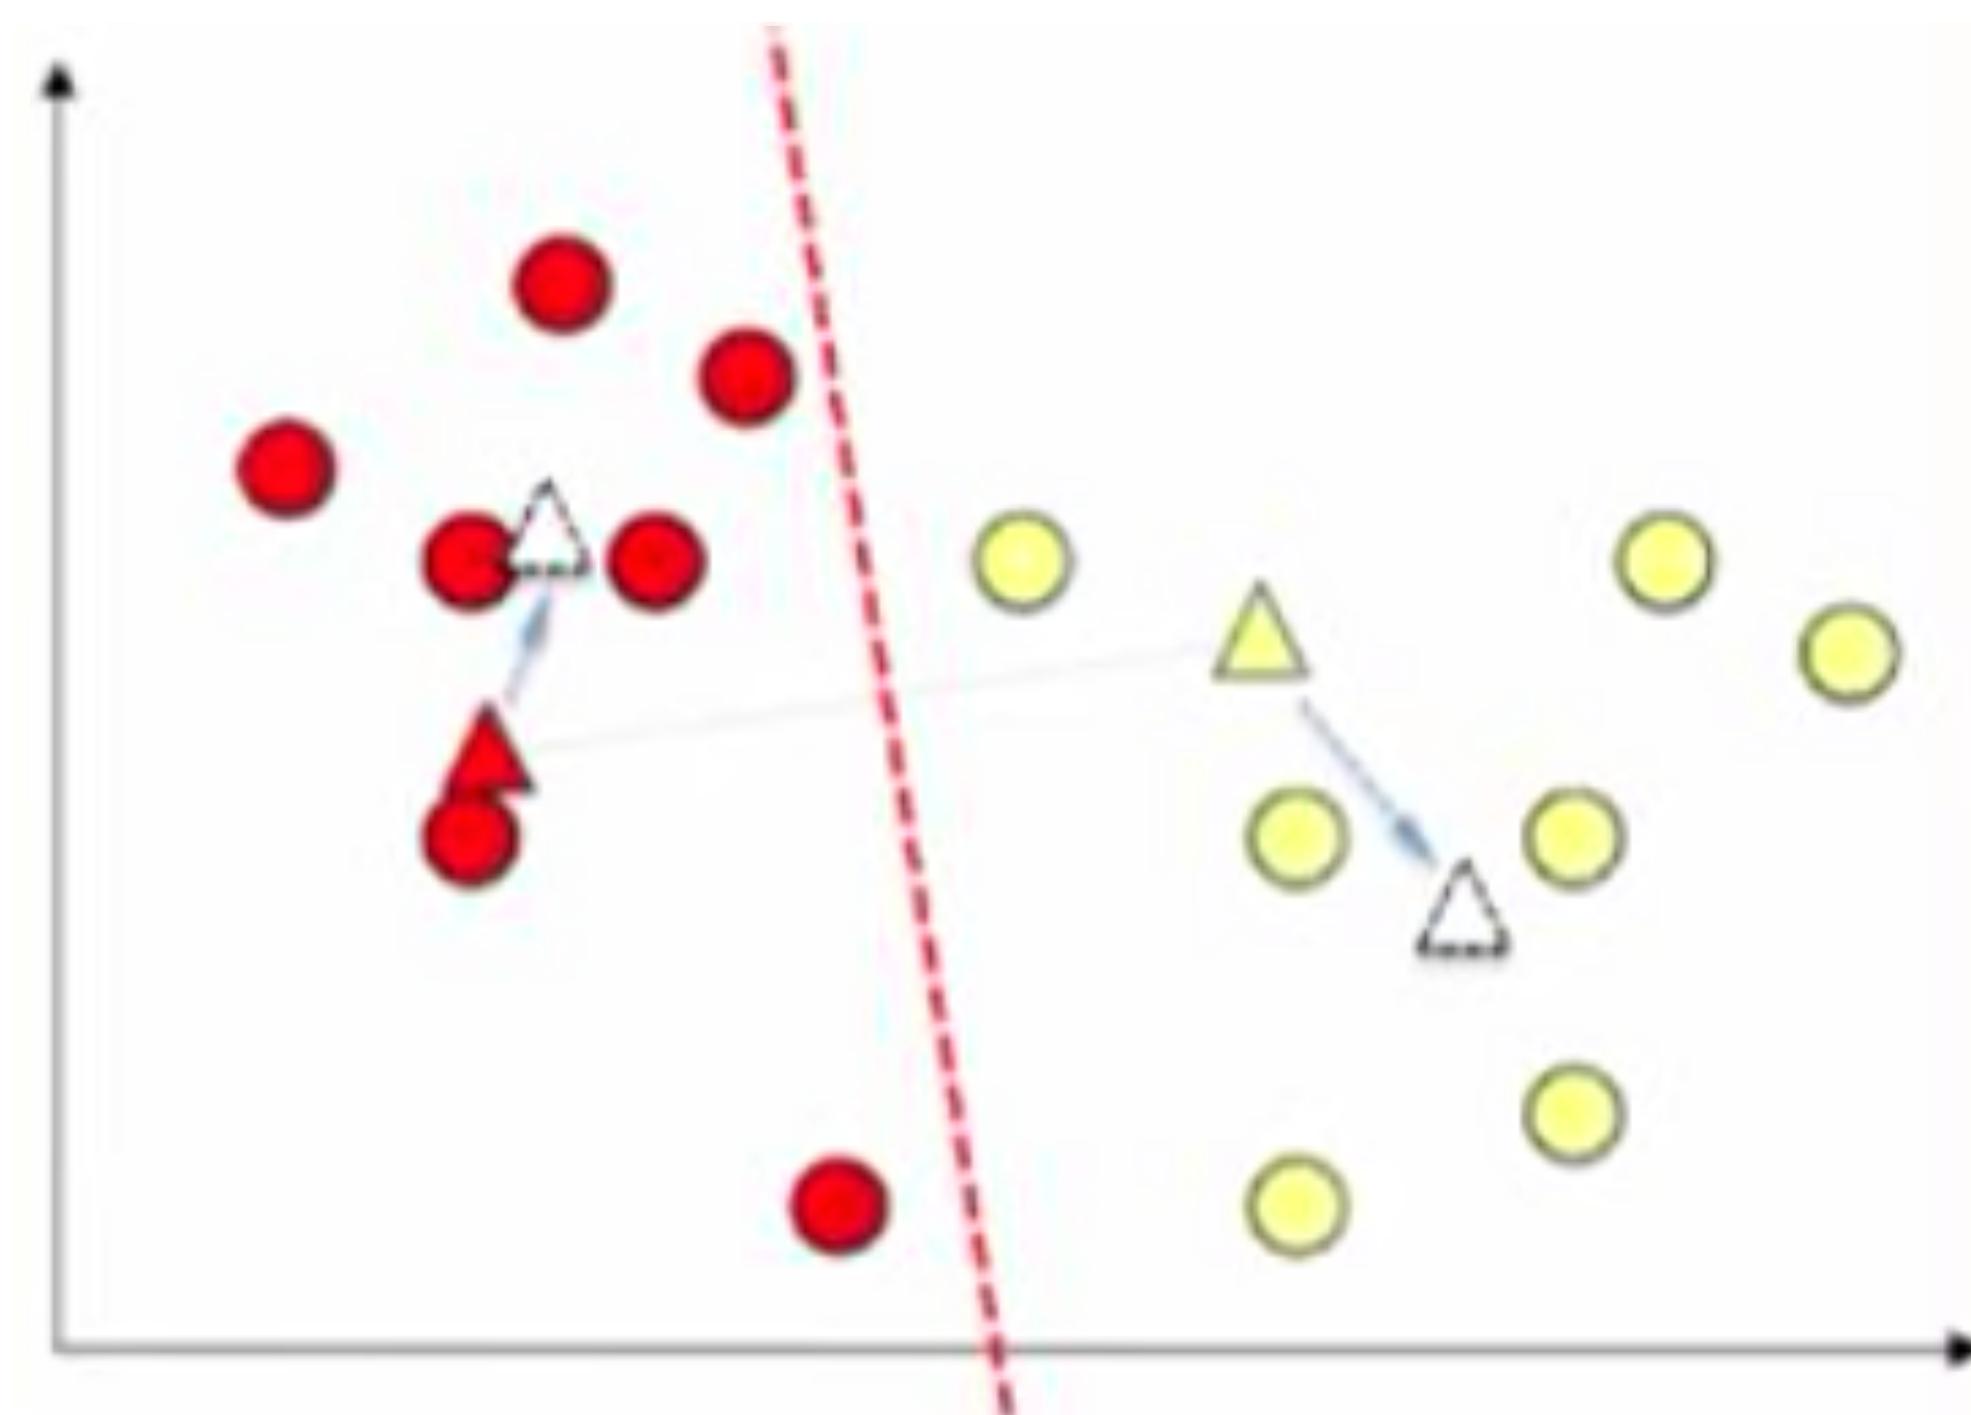
K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

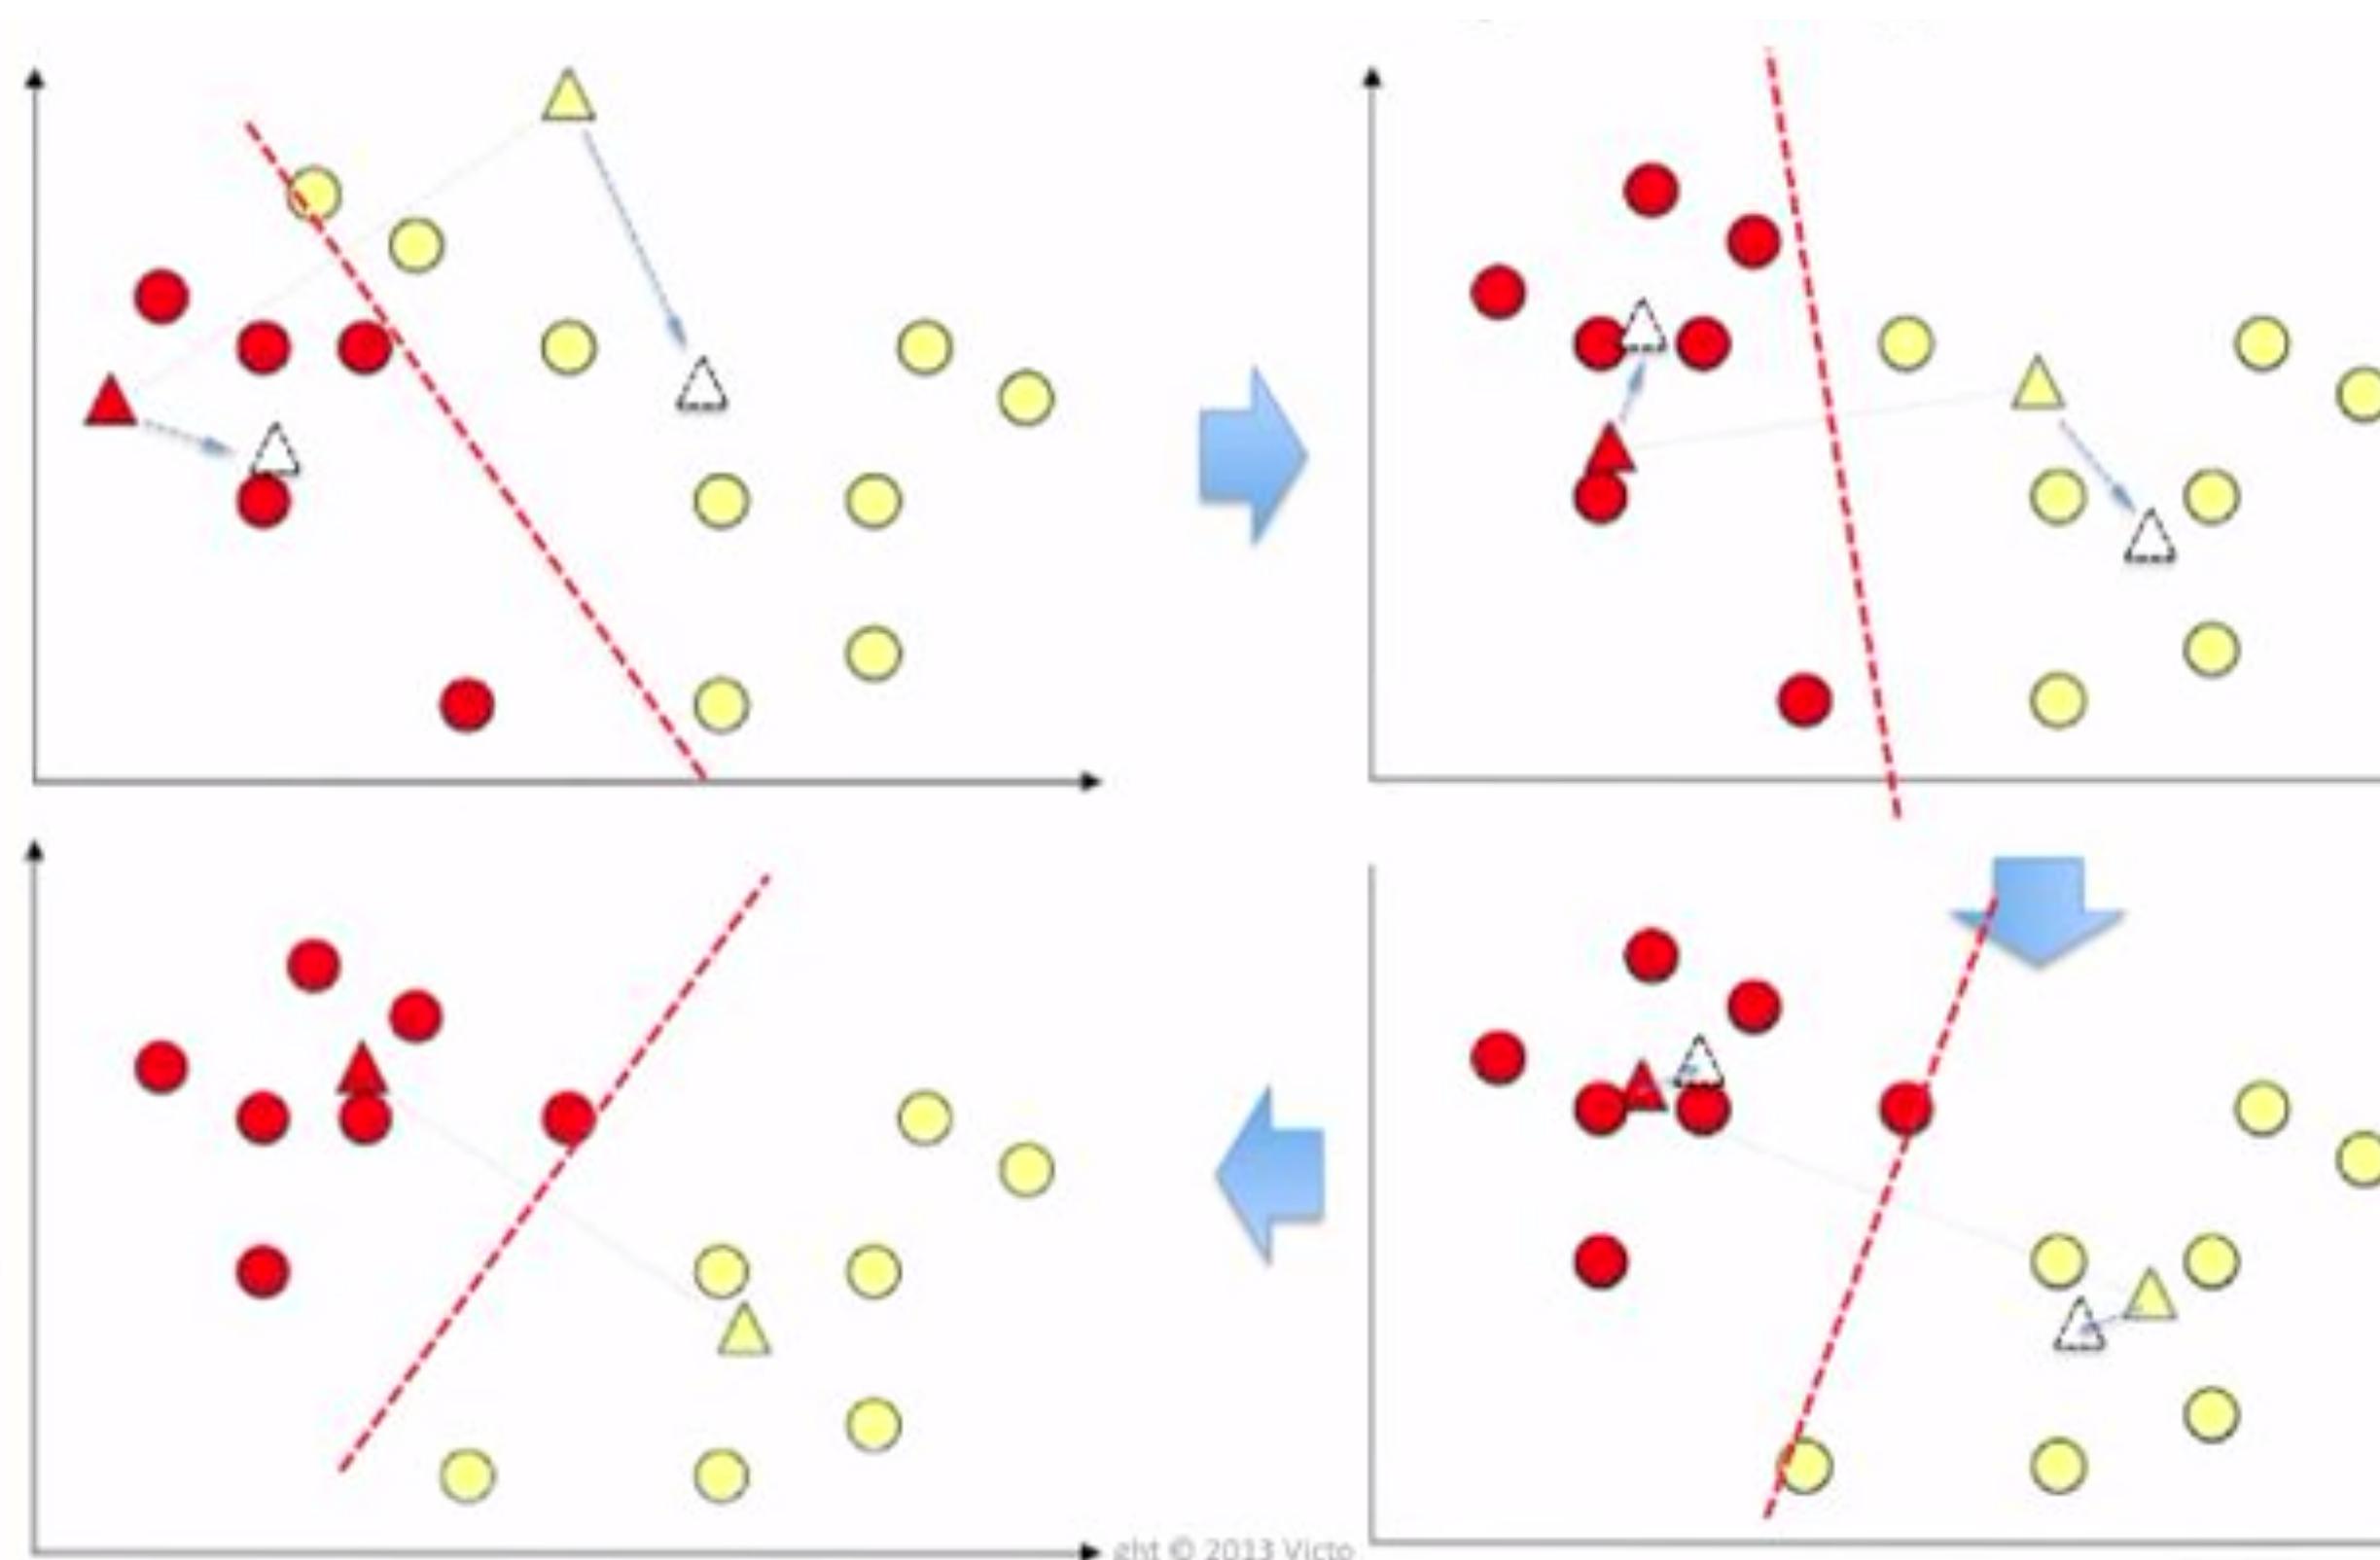
K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

K-means clustering: Example



Algorithm 8.1 Basic K-means algorithm.

- 1: Select K points as initial centroids.
 - 2: **repeat**
 - 3: Form K clusters by assigning each point to its closest centroid.
 - 4: Recompute the centroid of each cluster.
 - 5: **until** Centroids do not change.
-

K-means clustering

Minimizes the sum of squared errors (SSE)

$$\text{SSE} = \sum_{i=1}^K \sum_{\mathbf{x} \in C_i} \text{dist}(\mathbf{c}_i, \mathbf{x})^2$$

The centroid of the i th cluster is the mean of the points

$$\mathbf{c}_i = \frac{1}{m_i} \sum_{\mathbf{x} \in C_i} \mathbf{x}$$

Distance can be Euclidian, Manhattan, Minkowski,..

K-means clustering

Advantages

Easy to implement

Fast: $O(n)$

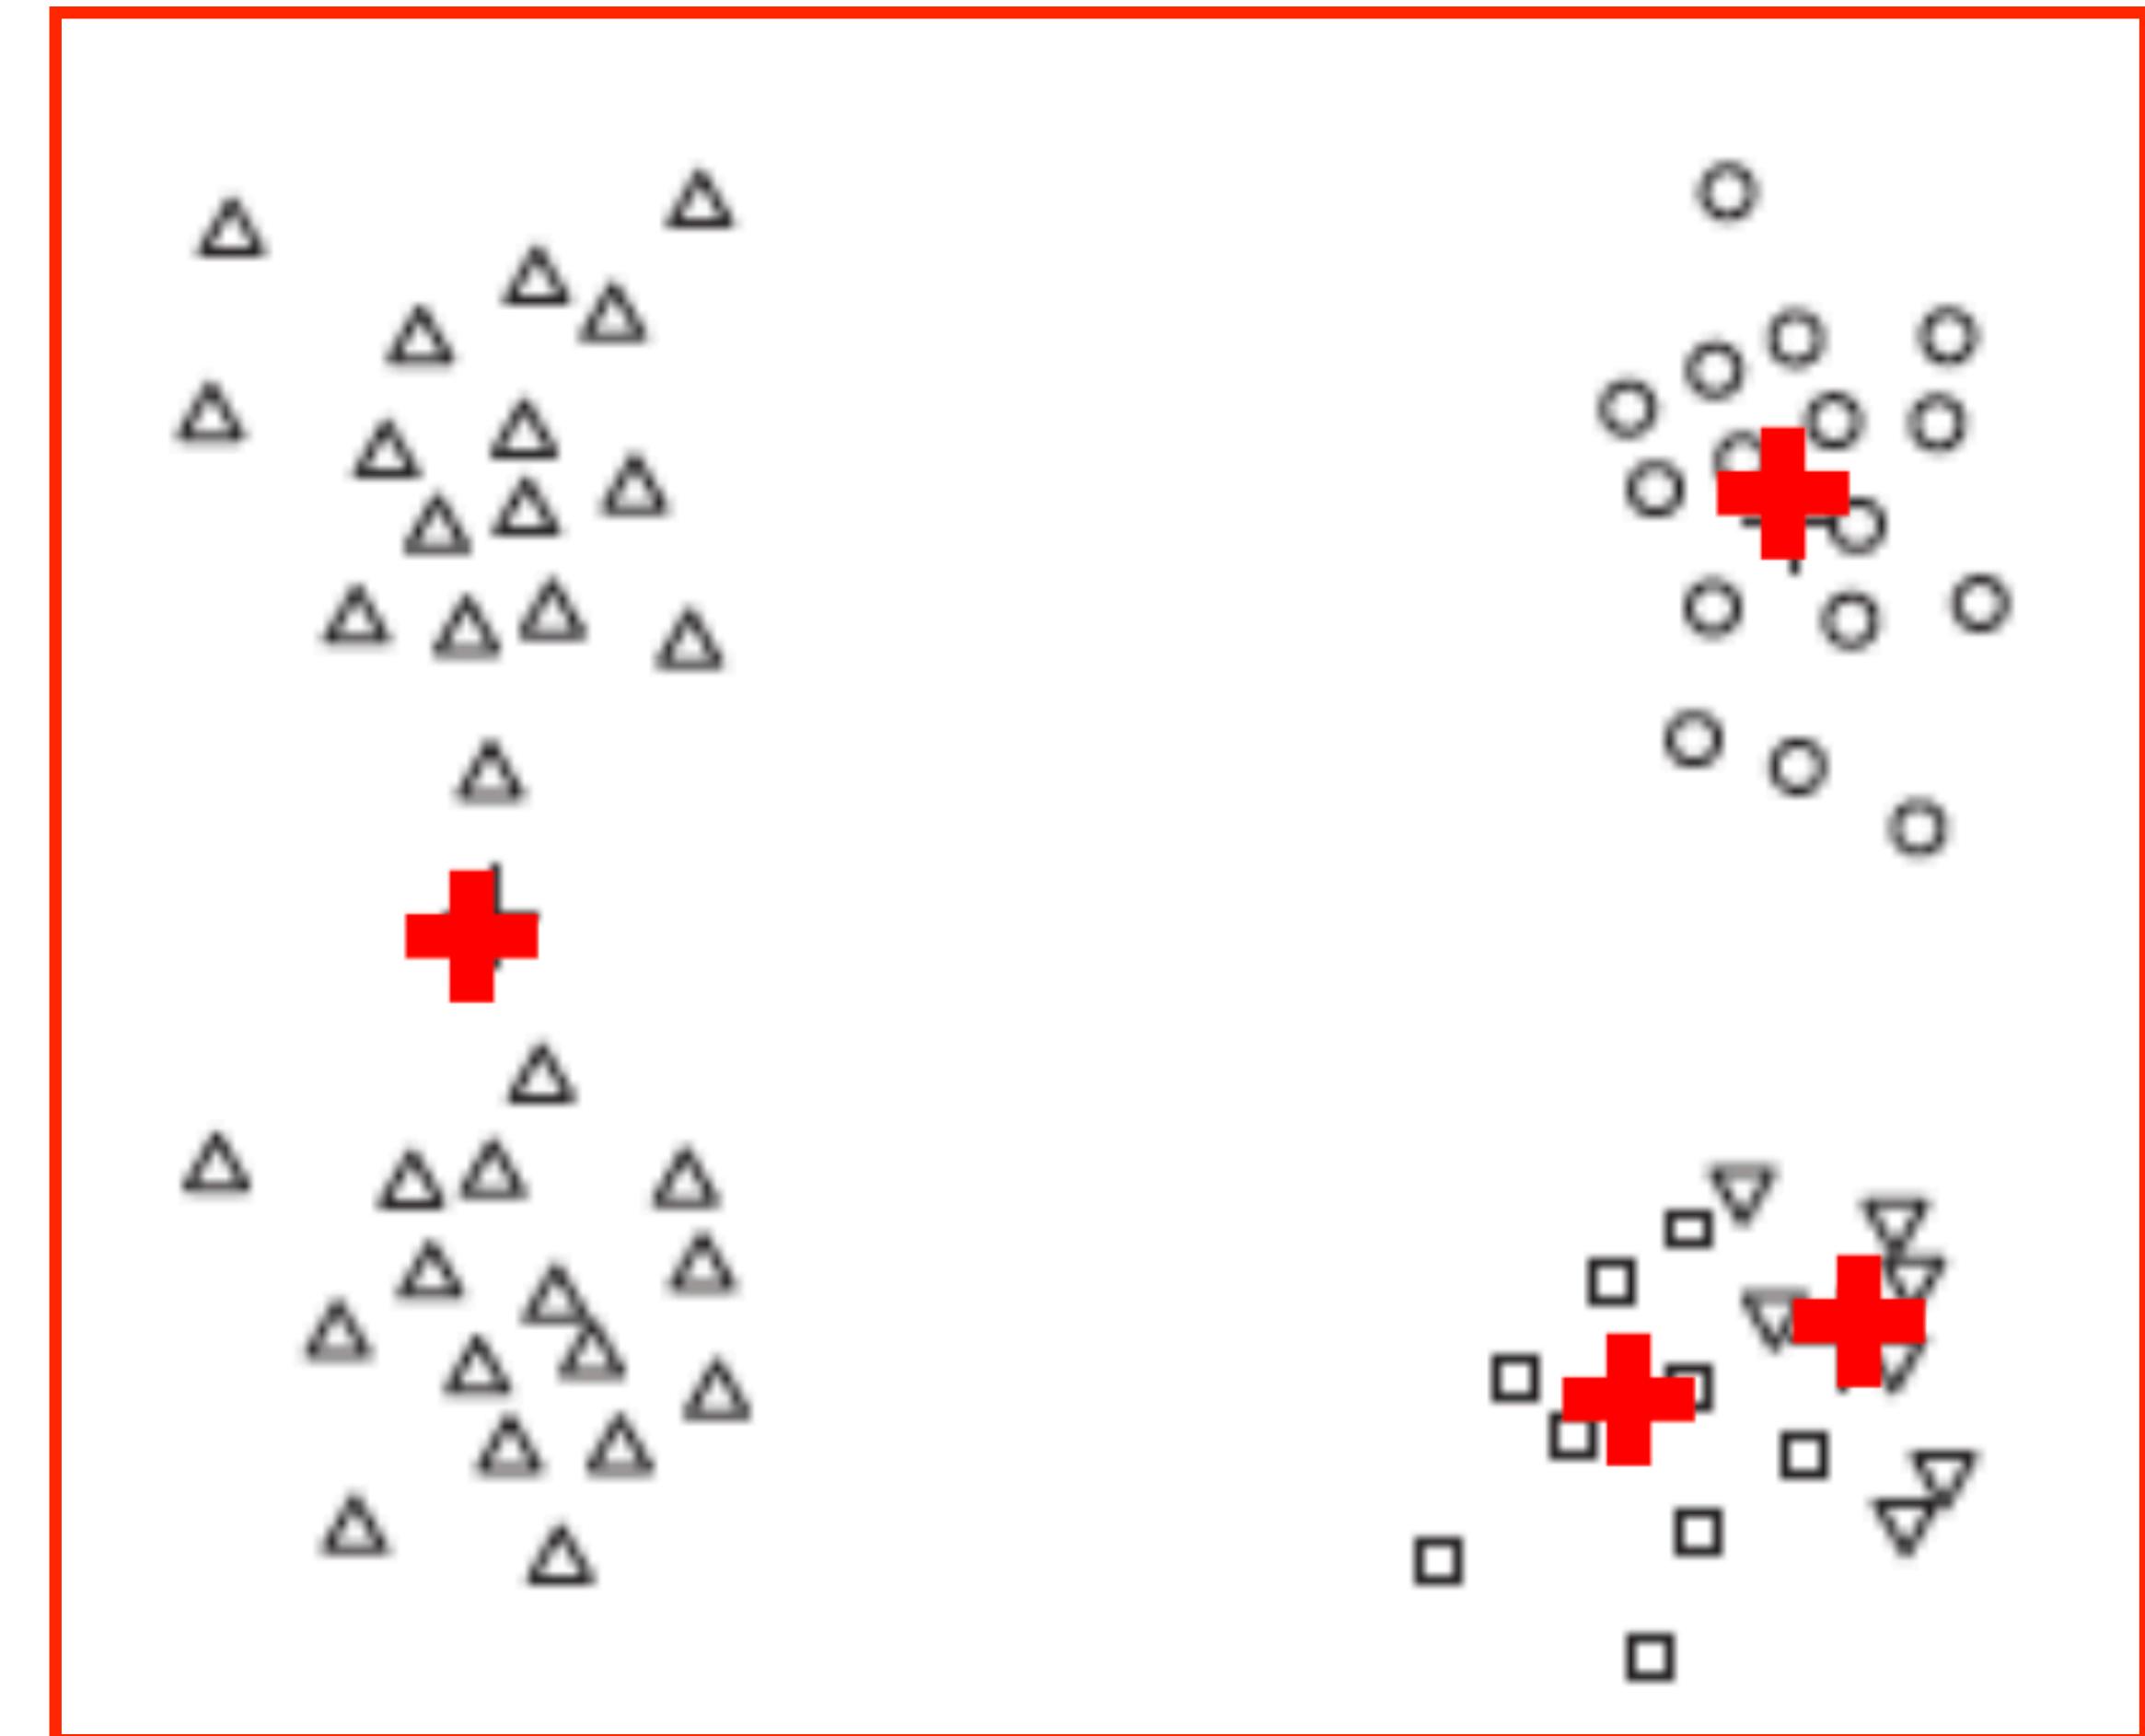
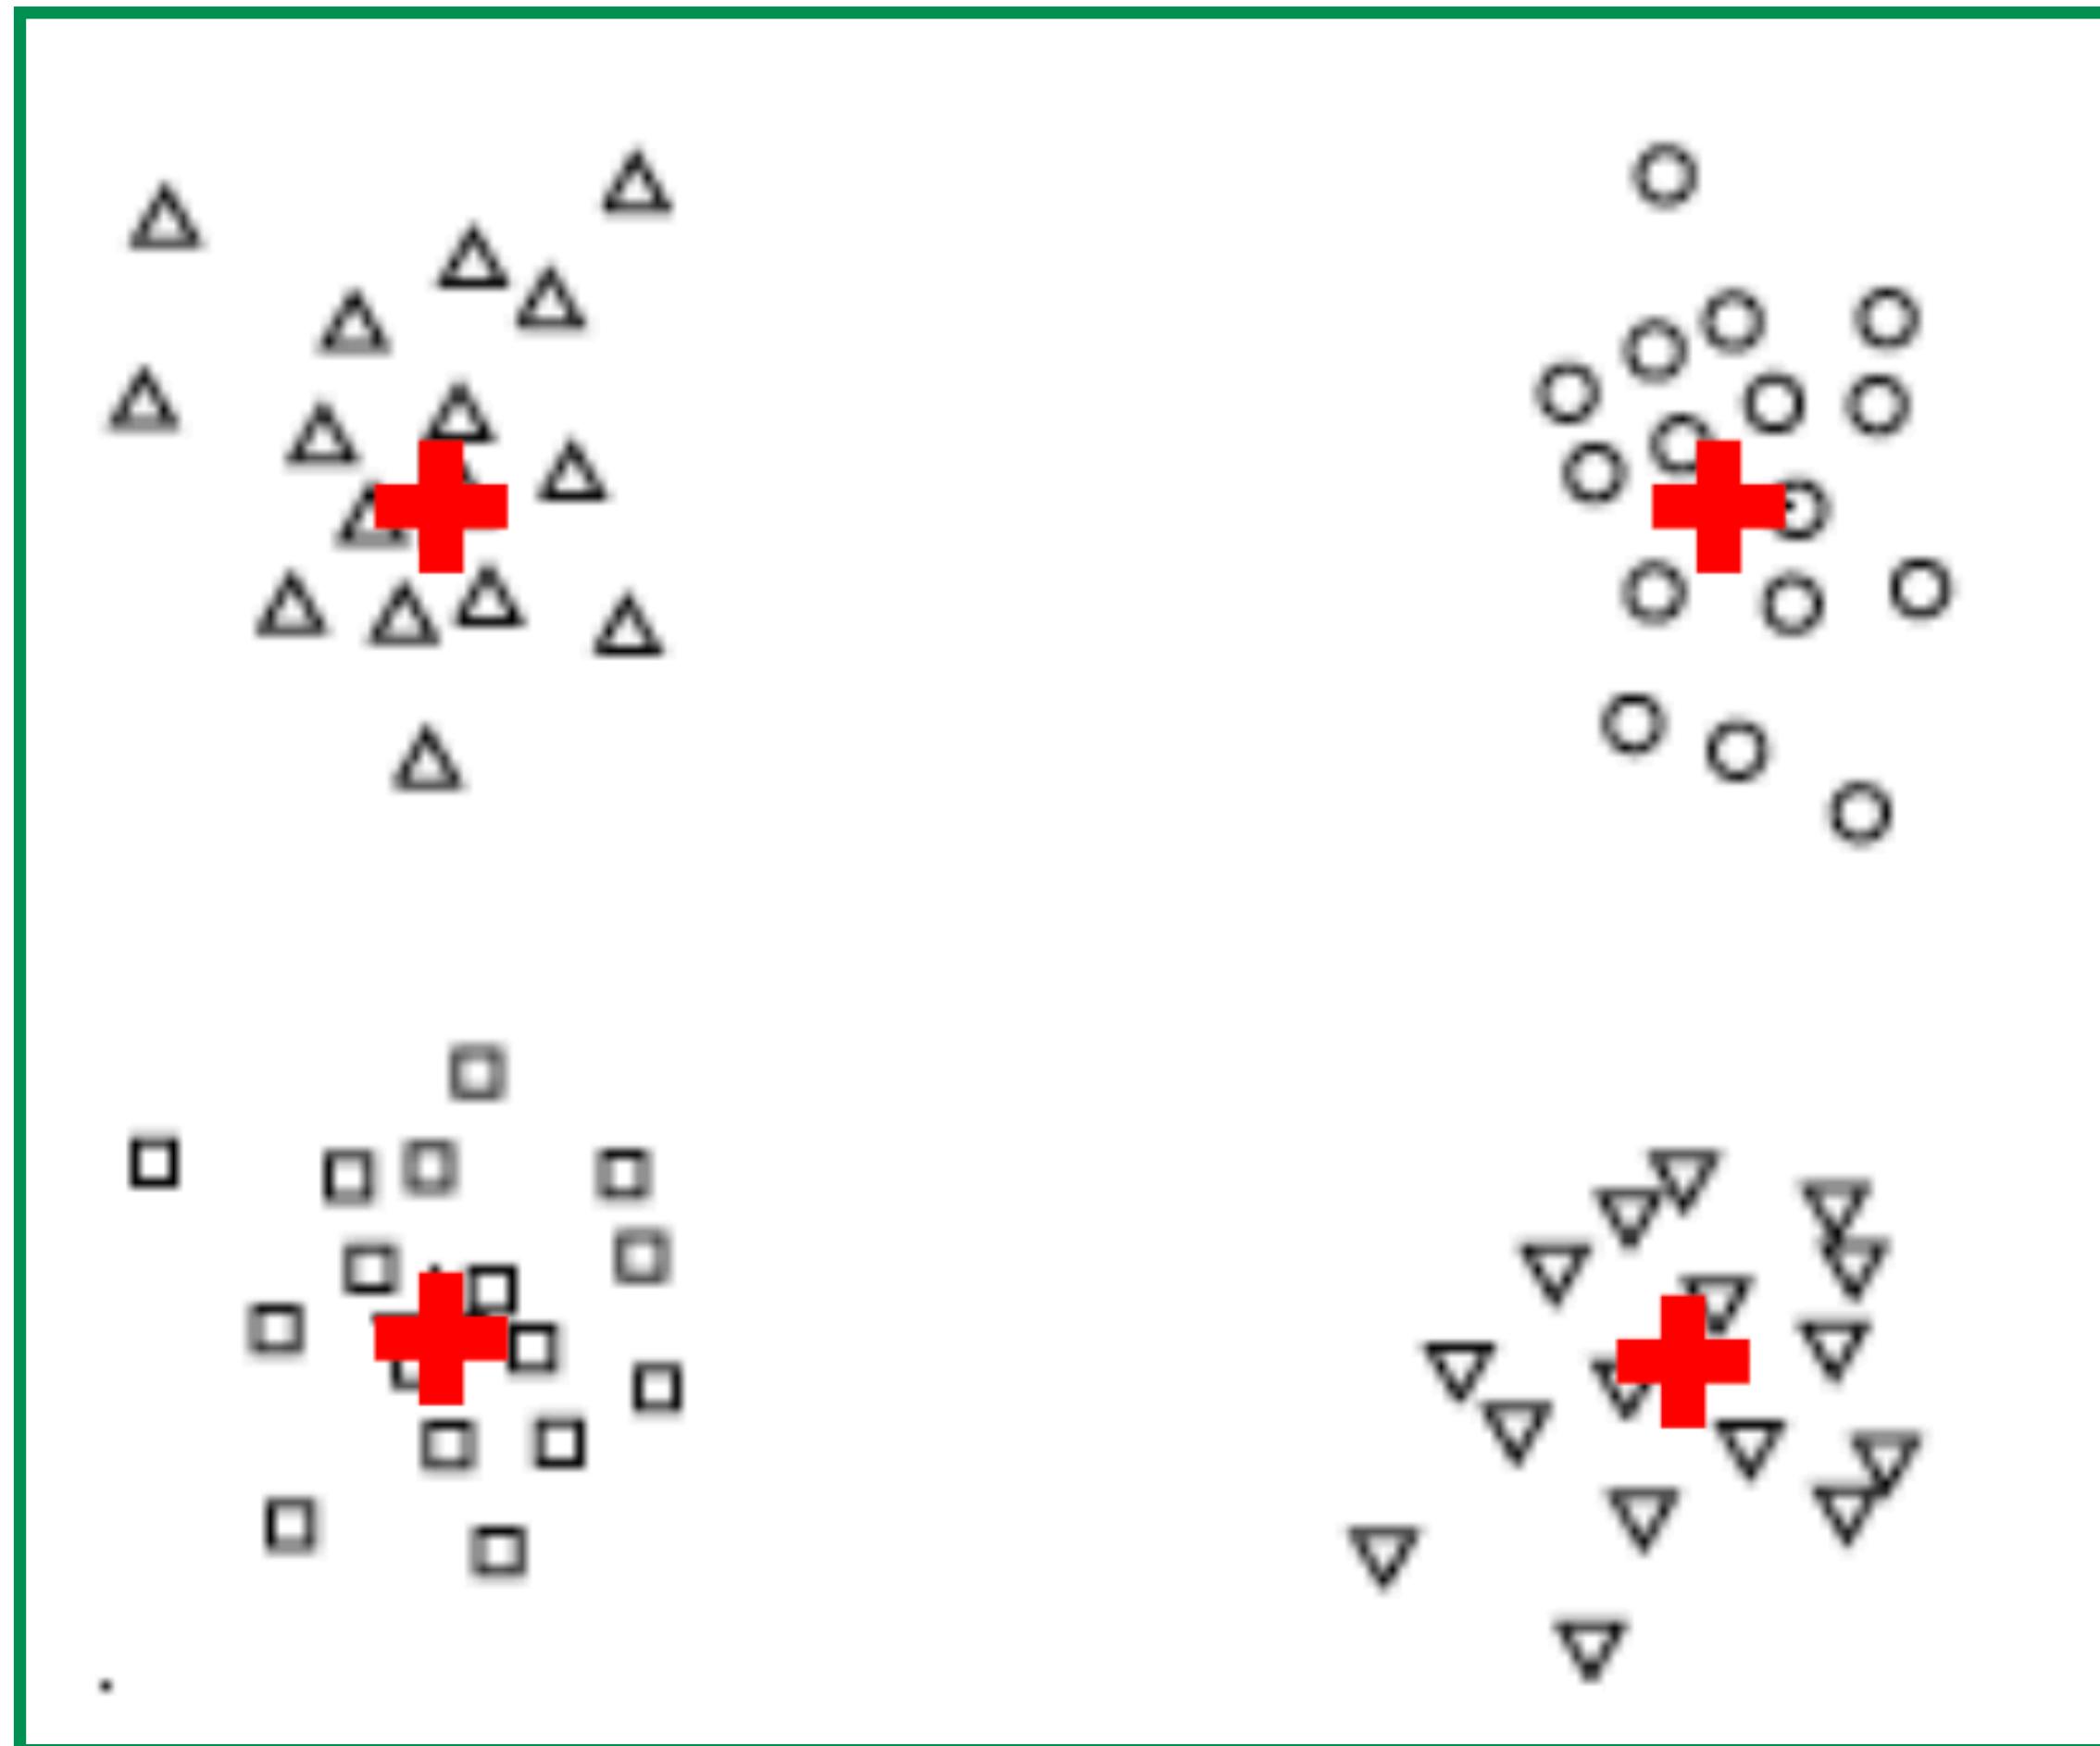
Disadvantages

Number of partitions needs to be known

Sensitivity to initial conditions

Not effective under several conditions

K-means can get stuck



K-means is often not working well



Figure 8.9. Clusters of different size

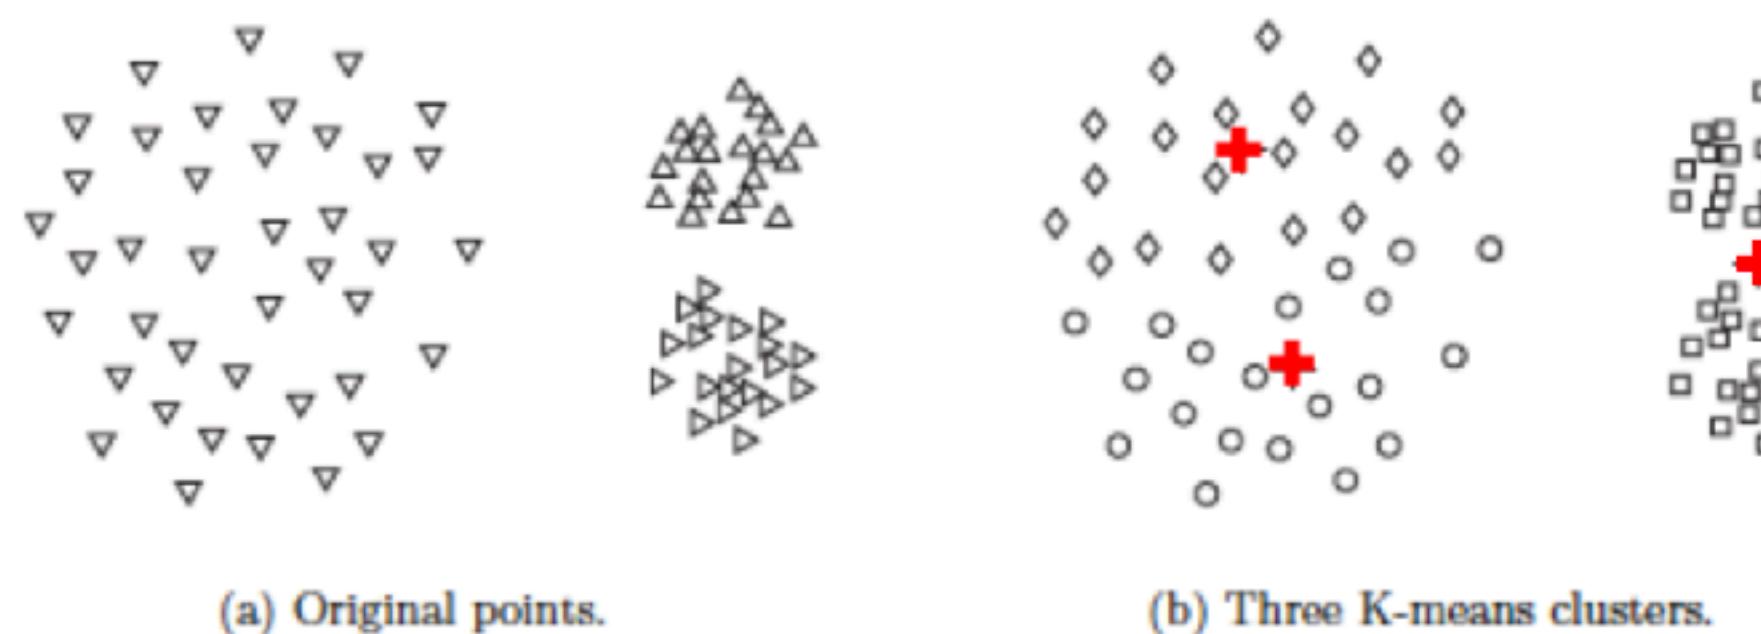


Figure 8.10 Clusters of different density

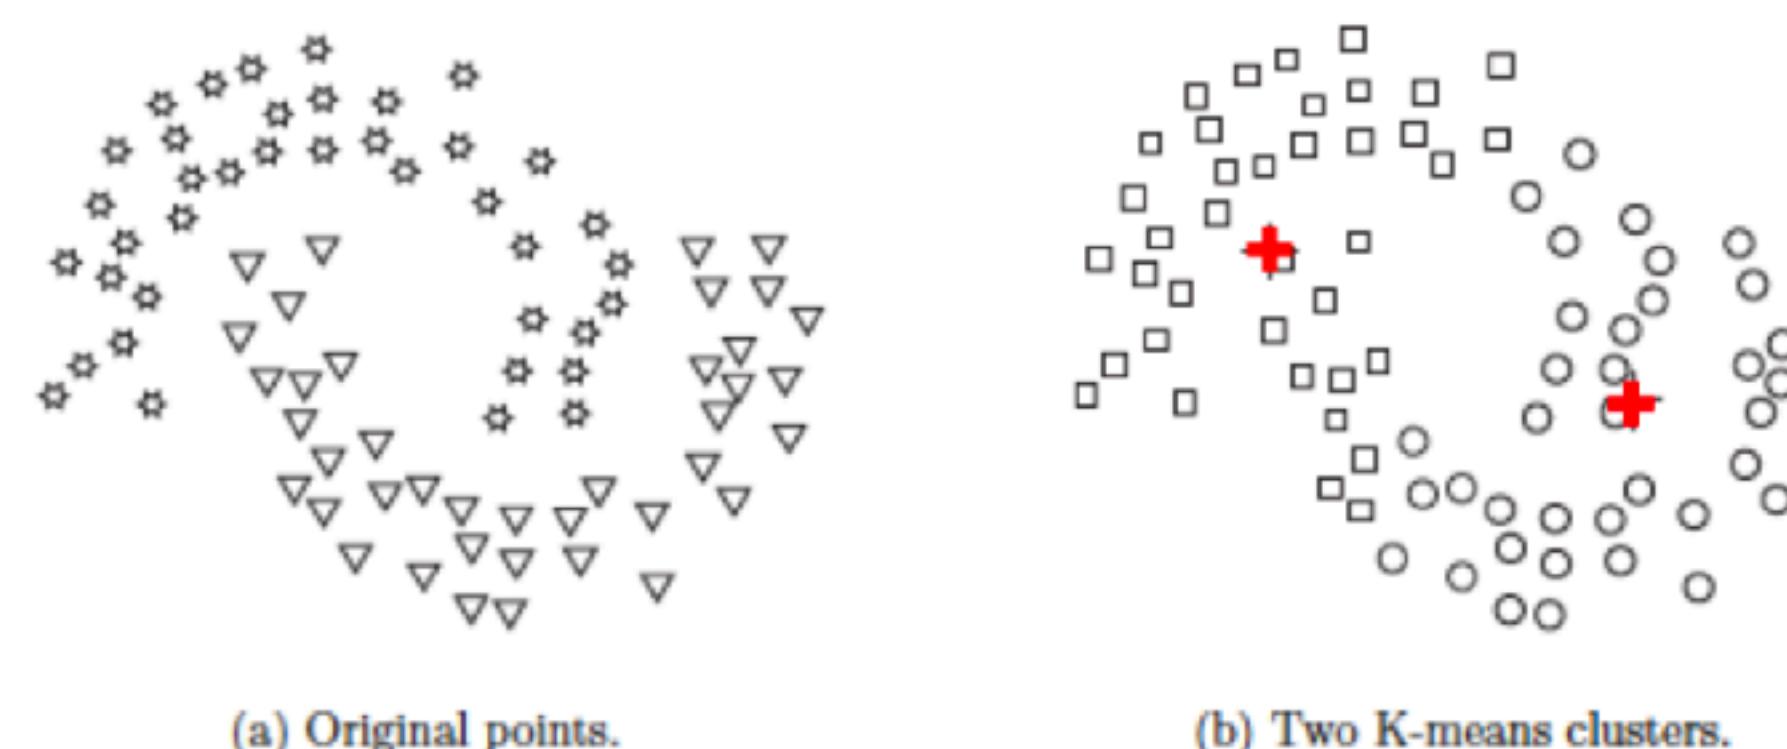


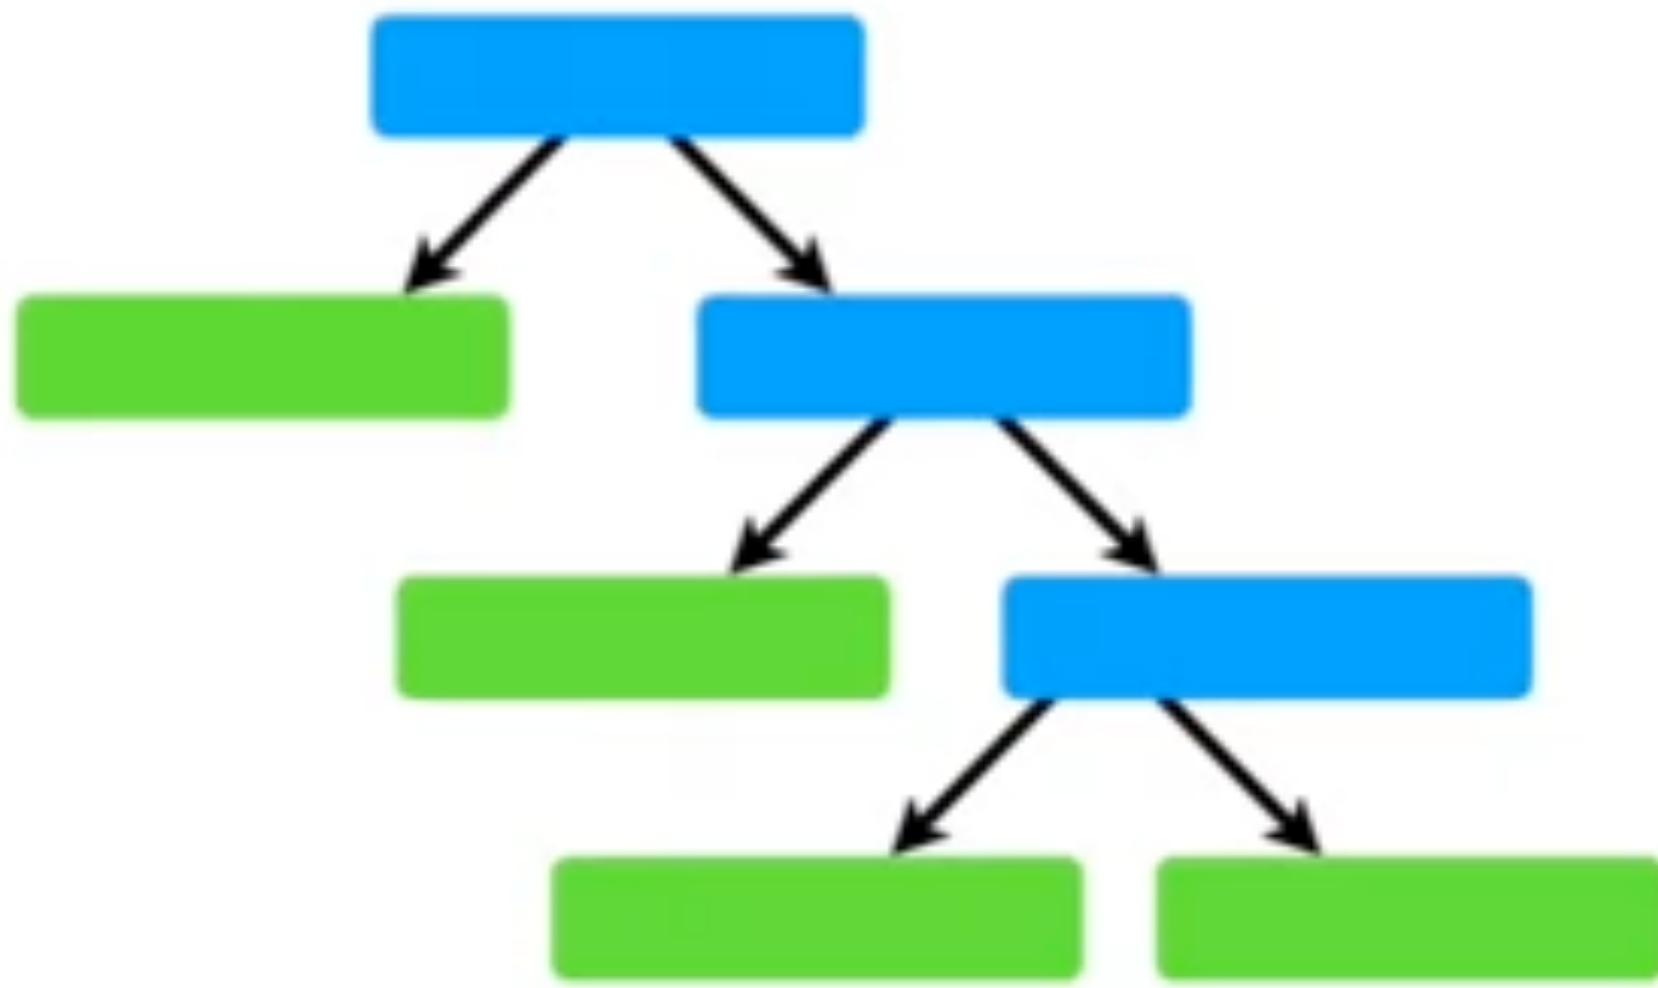
Figure 8.11. Non globular clusters

Today you learned about information theory

Entropy is the expected surprise

$$H(X) = \mathbb{E}(I(X))$$

..useful for:
Decision trees



Data Science setup



Clustering

